

# Distributed Computing

## A-15. Apache Spark

# References


- The paper and presentation by Zaharia et al. At USENIX NSDI 2012: *“Resilient Distributed Datasets: A Fault-Tolerant Abstraction for In-Memory Cluster Computing”*
- The Spark website: <https://spark.apache.org>

# MapReduce: Virtues & Shortcomings

- MapReduce has been a big improvement for “big data” on large clusters of unreliable machines
- However, it's less than perfect for important use cases
  - Multi-stage applications (e.g., iterative machine learning, graph processing)
  - Interactive ad-hoc queries
- Several specialized frameworks were designed to handle these cases


# Everything On Disk

- Each MapReduce job (i.e., something that has “one shuffle phase”) has to **read and write from disk**
- This is the solution to deal with unreliability: write everything on disk and replicate it
- However, this is **slow**




# In-Memory Processing

- What if things could be kept in RAM without the need of touching the disk every time?



- Huge speedup—10 to 100x
- The goal of Spark is to **get both this and fault tolerance**



# **Resilient Distributed Datasets (RDDs)**

# Resilient Distributed Datasets

- **Immutable, distributed** in-memory data structures
- **Partitioned** among the machines in the cluster
- Only built through **coarse-grained** operations that process the whole RDD
- Based on functional programming (map/filter/join...)
- Fault recovery performed using **lineage**
- A log of all the operations done to get there
- Recover lost partitions by **recomputing what's missing**
- No cost if nothing fails



# RDDs vs. Databases

- Databases & key-value stores handle **small updates** and **store everything on disk**
  - They're good for small modifications (transactions) that don't modify most of the state
- RDDs are efficient for **large operations**



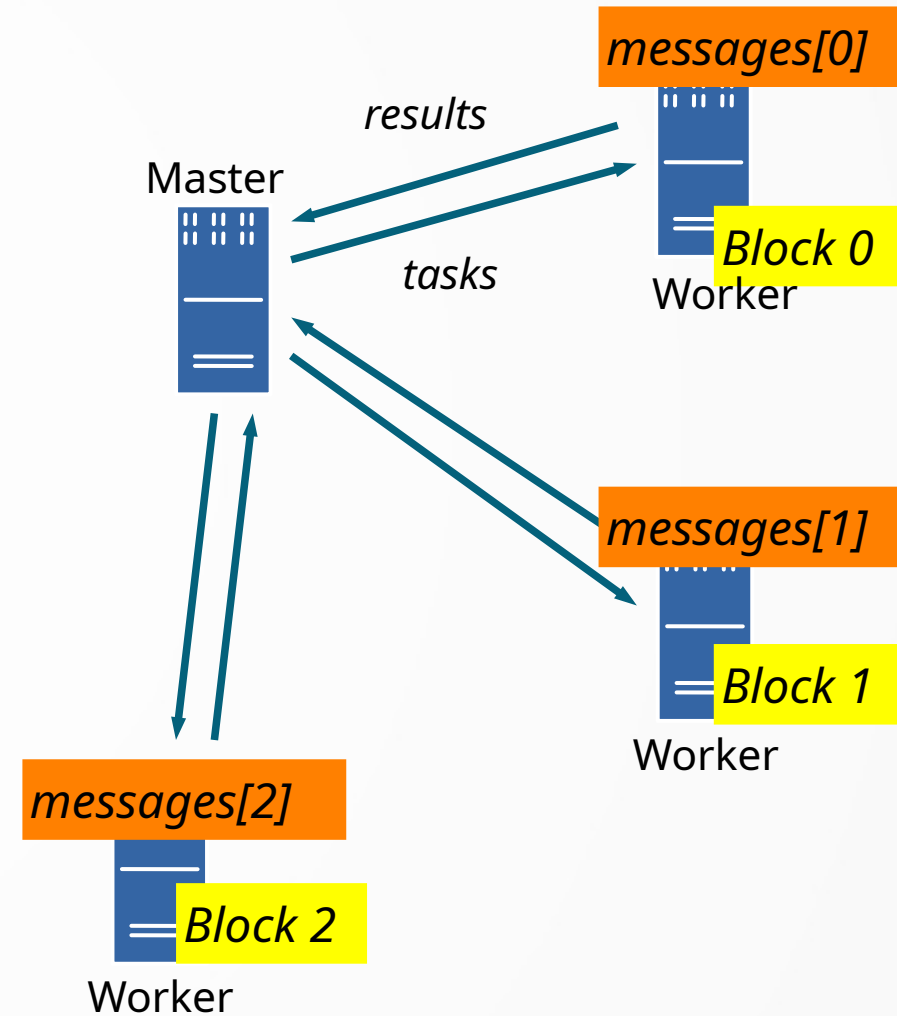
# Example: Log Mining

```
lines = spark.textFile("hdfs://...")

def is_error(line):
    return line.startswith('ERROR')
errors = lines.filter(is_error)

def get_message(line):
    return line.strip().split()[1]
messages = errors.map(get_message)
messages.persist()

messages.filter(lambda m: 'foo' in m).count()
messages.filter(lambda m: 'bar' in m).count()
```



# Fault Recovery

- RDDs tracks **lineage** (i.e. dependencies) for each block

```
messages=textFile(...) \  
    .filter(lambda x: 'error' in x) \  
    .map(lambda x: x.split()[1])
```

