

# Spark SQL

# SparkSQL

- Make data look like tables regardless how they really lay out
- SQL (standard!) based query can be directly against these tables
- Generate a specific execution plan for this query

# SparkSQL vs Spark

## Spark RDD API

```
rdd  
rdd1 = rdd.map(lambda x: x.split("\t"))  
rdd2 = rdd1.map(lambda x: (x[0], x[2]))
```

VS

## SQL on Spark

```
select user_id, url from access_log
```

- No data parsing
- Syntax is easier
- Code optimization
- No overhead

# Write Less Code: Compute an Average

## Using RDDs

```
data = sc.textFile(...).split("\t")
data.map(lambda x: (x[0], [int(x[1]), 1])) \
    .reduceByKey(lambda x, y: [x[0] + y[0], x[1] + y[1]]) \
    .map(lambda x: [x[0], x[1][0] / x[1][1]]) \
    .collect()
```

## Using SQL

```
SELECT name, avg(age)
FROM people
GROUP BY name
```

## Using DataFrames

```
sqlCtx.table("people") \
    .groupBy("name") \
    .agg("name", avg("age")) \
    .collect()
```

# SparkSQL: Spark as a database language? - Application scenario

- No real-time queries (high latency)!
- No row level updates!
- Not designed for online transaction processing!
- Best use: batch jobs over large sets of append-only data
  - Log processing
  - Data/Text mining
  - Business Intelligence
  - ...


# A step back ...What is Hive?



- A big data management system storing structured data on Hadoop file system
- It projects tabular schemas over folders and files in HDFS
- Enables the contents of folders in HDFS to be queried as tables, using SQL-like query semantics



# Spark 1.0





# Spark 2.0

Spark  
ML

Structured  
Streaming

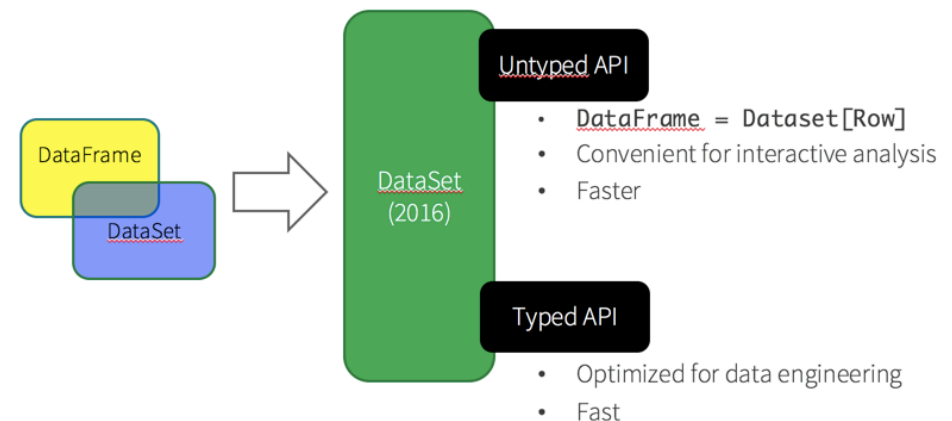
GraphFrames

Spark SQL

Spark RDD

# SparkSQL

- programming abstractions called DataFrame that can act as a distributed SQL query engine
- The input data can be queried by using
  - ad-hoc methods
  - an SQL-like language

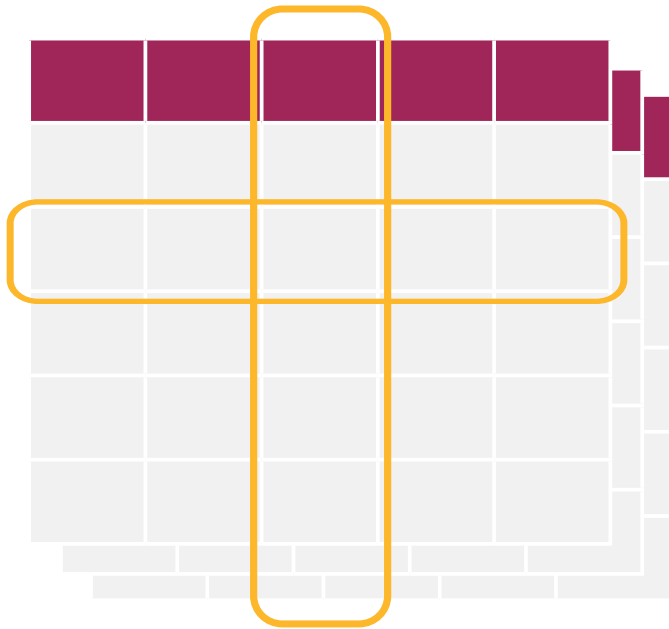


# Structured API

- The interfaces provide more information about the structure of both the data and the computation being performed
- Spark uses this extra information to perform extra optimizations based on an “SQL-like” optimizer called **Catalyst**
- ➔ • Compute the best execution plan before executing the code  
=> Programs are usually faster than standard RDD-based programs

# DataFrames

# DataFrames = RDD+Schema



DataFrame = RDD + schema

# DataFrame

**noun** – [dey-tuh-freym]

1. A distributed collection of rows organized into named columns.
2. An abstraction for selecting, filtering, aggregating and plotting structured data (*cf. R, Pandas*).
3. Archaic: Previously SchemaRDD (*cf. Spark < 1.3*).

# DataFrames = RDD+Schema

```
In: from pyspark.sql import SparkSession
    spark_session = SparkSession.builder\
        .enableHiveSupport()\
        .appName("spark sql")\
        .master("local")\
        .getOrCreate()
```

```
In: geoip_rdd = spark_session\
    .sparkContext\
    .textFile("..."/geoip")
```

```
In: geoip_rdd.take(3)
```

```
Out: [u'194.120.126.123, NL, Netherlands',
      u'94.126.119.173, FR, France',
      u'193.46.74.166, RU, Russian Federation']
```

# DataFrames = RDD+Schema

ip **STRING**,  
code **STRING**,  
country **STRING**

```
In: from pyspark.sql.types import *  
    schema = StructType().add("ip", StringType())\  
                             .add("code", StringType())\  
                             .add("country", StringType())
```

```
In: geoip_df = spark_session\  
    .createDataFrame(geoip_rdd1, schema)
```

```
In: geoip_df
```

```
Out: DataFrame[ip: string, code: string, country: string]
```



# DataFrames = RDD+Schema

In: `geoip_df.show(3)`

```
+-----+-----+-----+
|          ip|code|          country|
+-----+-----+-----+
|194.120.126.123|  NL|      Netherlands|
| 94.126.119.173|  FR|          France|
| 193.46.74.166|  RU|Russian Federation|
+-----+-----+-----+
```

only showing top 3 rows

# DataFrames = RDD+Schema

```
In: geoip_df.rdd.take(3)
```

```
Out: [Row(ip=u'194.120.126.123', code=u'NL',  
        country=u'Netherlands'),  
      Row(ip=u'94.126.119.173', code=u'FR',  
        country=u'France'),  
      Row(ip=u'193.46.74.166', code=u'RU',  
        country=u'Russian Federation')]
```

# DataFrames = RDD+Schema

```
In: geoip_df.printSchema()
```

```
root
```

```
 |-- ip: string (nullable = true)  
 |-- code: string (nullable = true)  
 |-- country: string (nullable = true)
```

```
In: geoip_df
```

```
Out: DataFrame[ip: string, code: string, country: string]
```

## show

```
In: geoip_df.show(3)
```

```
+-----+-----+-----+
|          ip|code|          country|
+-----+-----+-----+
|194.120.126.123|  NL|      Netherlands|
| 94.126.119.173|  FR|          France|
| 193.46.74.166|  RU|Russian Federation|
+-----+-----+-----+
only showing top 3 rows
```

# Projection

$\pi$

**select**

```
In: geoip_df.select("country", "ip")\  
      .show(3)
```

```
+-----+-----+  
|          country|          ip|  
+-----+-----+  
|    Netherlands|194.120.126.123|  
|          France| 94.126.119.173|  
|Russian Federation| 193.46.74.166|  
+-----+-----+
```

only showing top 3 rows

# where

*Filtering  
(selection)*

```
In: geoip_df\  
    .select("country", "ip")\  
    .where("country = 'Russian Federation'")\  
    .show(3)
```

$\sigma$

country	ip
Russian Federation	193.46.74.166
Russian Federation	46.235.67.202
Russian Federation	193.161.193.64

only showing top 3 rows

```
In: step1 = geoip_df.select("country", "ip")
```

```
In: step2 = step1.where("country = 'Russian Federation'")
```

```
In: step3 = step2.show(3)
```

```
+-----+-----+
|          country|          ip|
+-----+-----+
|Russian Federation| 193.46.74.166|
|Russian Federation| 46.235.67.202|
|Russian Federation|193.161.193.64|
+-----+-----+
only showing top 3 rows
```

# Spark execution model: transformations, actions, lazy evaluation

In: `type(step1)`

Out: `pyspark.sql.dataframe.DataFrame`

Transformation: DataFrame -> DataFrame

In: `type(step2)`

Out: `pyspark.sql.dataframe.DataFrame`

Action: Dataframe -> None

In: `type(step3)`

Out: `NoneType`



```
In: %%time
step1 = geoip_df.select("country","ip")
```

Out: CPU times: user 4 ms, sys: 0 ns, total: 4 ms  
Wall time: 28.8 ms

```
In: %%time
step2 = step1.where("country = 'Russian Federation'")
```

Out: CPU times: user 4 ms, sys: 0 ns, total: 4 ms  
Wall time: 13.8 ms

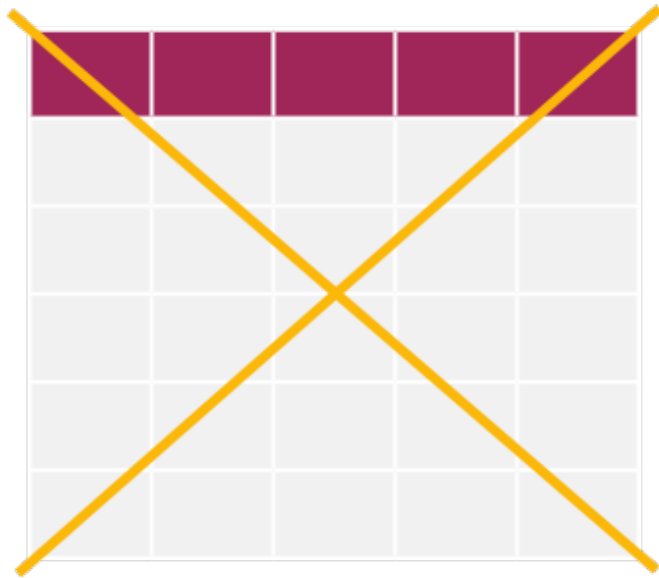
```
In: %%time
step2.show(3)
```

Out: +-----+-----+  
| country| ip|  
+-----+-----+  
Russian Federation	193.46.74.166
Russian Federation	46.235.67.202
Russian Federation	193.161.193.64
+-----+-----+  
only showing top 3 rows

CPU times: user 0 ns, sys: 0 ns, total: 0 ns  
Wall time: 181 ms


DataFrames vs SQL  
Which Spark-API to use?

# DataFrames and SQL



Table

DataFrame



View

```
In: geoip_df.createTempView("geoip")
```

```
spark_session.sql
```

```
In: spark_session.sql("""
    select country from geoip
    """)
```

"registering"  
the DataFrame so that  
it can be used in the from

```
Out: DataFrame[country: string]
```

# Which Spark API to use?

## RDD API

In [9]:

```
geoip_df.rdd\  
  .map(lambda x: Row(ip=x.ip, country=x.country))\  
  .filter(lambda x: x.country == "Russian Federation")\  
  .take(3)
```

## DataFrame API

In [10]:

```
geoip_df.select("ip", "country")\  
  .where("country='Russian Federation'")\  
  .filter(lambda x: x.country == "Russian Federation")\  
  .show(3)
```

## SQL

In [11]:

```
geoip_df.createOrReplaceTempView("geoip")
```

In [12]:

```
spark_session.sql("""  
  select ip,  
         country  
  from geoip  
  where country='Russian Federation'  
""").show(3)
```

# Which Spark API to use?

RDD

vs

DataFrame & SQL



# Which Spark API to use?

By DataFrame API      vs      By SQL command

```
geoip_df\  
  .selec("ip", "country")\  
  .where("country='Russia'")\  
  .show(3)
```

Error will be found at  
compilation

```
spark_session.sql("""  
  selec ip,  
          country  
  from geoip  
  where country='Russia'  
""").show(3)
```

Error will be found at  
query call

# Creating Data Frames

# Creating DataFrames

- DataFrames can be created
  - from CSV (with options, e.g. `inferschema`, `header`)
  - from text file (standard «multiline» JSON)
  - from an RDD
    - Programmatically specifying the schema
    - Inferring the schema through reflection



# Reading and Writing DataFrames

Built-In



External



Unified interface to reading/writing data in a variety of formats:

```
df = sqlContext.read \
```

# Write - Formats

```
In: geoip_df.write.save("geoip_csv",  
                        format='csv')
```

```
In: spark_session\  
    .sparkContext\  
    .textFile("geoip_csv")\  
    .take(3)
```

```
Out: [u'194.120.126.123,NL,Netherlands',  
      u'94.126.119.173,FR,France',  
      u'193.46.74.166,RU,Russian Federation']
```

# Write - Formats

```
In: geoip_df.write.save("geoip_json",  
                        format='json')
```

```
In: spark_session\  
    .sparkContext\  
    .textFile("geoip_json")\  
    .take(3)
```

```
Out: [u'{"ip": "194.120.126.123", "code": "NL", "country": "Net  
herlands"}',  
      u'{"ip": "94.126.119.173", "code": "FR", "country": "Fran  
ce"}',  
      u'{"ip": "193.46.74.166", "code": "RU", "country": "Russi  
an Federation"}']
```

# Read

```
In: geoip_from_table = spark_session\  
    .read.table("web.geoip")
```

```
In: geoip_from_table.show(3)
```

ip	code	country
194.120.126.123	NL	Netherlands
94.126.119.173	FR	France
193.46.74.166	RU	Russian Federation

only showing top 3 rows

# Read

```
In: geoip_from_json = spark_session\  
    .read.json("geoip_json")
```

```
In: geoip_from_json.show(3)
```

code	country	ip
NL	Netherlands	194.120.126.123
FR	France	94.126.119.173
RU	Russian Federation	193.46.74.166

only showing top 3 rows

# Read

```
In: geoip_from_csv = spark_session\  
    .read.csv("geoip_csv")
```

```
In: geoip_from_csv.show(3)
```

```
+-----+-----+-----+  
|          _c0|_c1|          _c2|  
+-----+-----+-----+  
|   217.8.92.38| RU|Russian Federation|  
|185.102.10.199| RU|Russian Federation|  
|   217.73.57.80| RU|Russian Federation|  
+-----+-----+-----+
```

only showing top 3 rows

# Read – Adding column names

```
In: schema = StructType().add("ip",      StringType())\  
                                .add("code",  StringType())\  
                                .add("country", StringType())
```


```
In: geoip_from_csv = spark_session\  
    .read.csv("geoip_csv")
```

```
In: geoip_from_csv.show(3)
```

```
+-----+-----+-----+  
|          ip|code|          country|  
+-----+-----+-----+  
|  217.8.92.38|  RU|Russian Federation|  
|185.102.10.199|  RU|Russian Federation|  
|  217.73.57.80|  RU|Russian Federation|  
+-----+-----+-----+  
only showing top 3 rows
```

# JDBC

## Import and export data to DBMSs



## JDBC - Java DataBase Connectivity

```
In: connection_string="jdbc:mysql://localhost:3306/demo?"\  
    "user=demo&"\  
    "password=demo"
```



```
In: geoip_from_jdbc = spark_session\  
    .read.jdbc(connection_string, "geoip")
```

```
In: geoip_from_jdbc.show(3)
```

ip	code	country
194.120.126.123	NL	Netherlands
94.126.119.173	FR	France
193.46.74.166	RU	Russian Federation

only showing top 3 rows

```
In: geoip_df.write.jdbc(connection_string, "geoip")
```

```
In: geoip_from_jdbc = spark_session\  
    .read.jdbc(connection_string, "geoip")
```

```
In: geoip_from_jdbc.show(3)
```

ip	code	country
194.120.126.123	NL	Netherlands
94.126.119.173	FR	France
193.46.74.166	RU	Russian Federation

only showing top 3 rows

# Data Frames Operations

# Projection & Filtering

```
spark_session.read.table("web.access_log")\
    .show(3)
```

http_code		ip	response_length	time	url	user_agent
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...

# Projection - SQL

```
spark_session.sql("""  
    select ip, url  
    from web.access_log  
""").show(3)
```

	ip	url
0	109.106.133.8	/id53821
1	46.31.82.254	/id33929
2	193.124.254.46	/id35754

# Projection - Dataframes

```
access_log_df.select("ip", "url")\n    .show(3)
```

	ip	url
0	109.106.133.8	/id53821
1	46.31.82.254	/id33929
2	193.124.254.46	/id35754

# Selection - SQL

```
spark_session.sql("""
    select *
    from web.access_log
    where http_code <> '200'
""").show(3)
```

http_code		ip	response_length	time	url	user_agent
0	404	91.206.117.71	0	12/Dec /2015:01:32:04 +0400	/favicon.ico	Mozilla/5.0 (X11; Linux x86_64; rv:31.0) Gecko...
1	404	23.39.172.114	0	12/Dec /2015:01:32:05 +0400	/favicon.ico	Mozilla/5.0 (Windows NT 6.3; WOW64) AppleWebKi...
2	404	176.120.130.254	0	12/Dec /2015:01:32:06 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-at; H...

# Selection - Dataframes

Where

Filter

```
access_log_df.where("http_code <> '200'") show(3)
```

http_code		ip	response_length	time	url	user_agent
0	404	91.206.117.71	0	12/Dec /2015:01:32:04 +0400	/favicon.ico	Mozilla/5.0 (X11; Linux x86_64; rv:31.0) Gecko...
1	404	23.39.172.114	0	12/Dec /2015:01:32:05 +0400	/favicon.ico	Mozilla/5.0 (Windows NT 6.3; WOW64) AppleWebKi...
2	404	176.120.130.254	0	12/Dec /2015:01:32:06 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-at; H...



# Selection – SQL (Composite Conditions)

```
spark_session.sql("""
    select *
    from web.access_log
    where http_code<>'200'
    and user_agent like '%Android%'
""").show(3)
```

http_code		ip	response_length	time	url	user_agent
0	404	176.120.130.254	0	12/Dec /2015:01:32:06 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-at; H...
1	404	93.188.131.176	0	12/Dec /2015:01:32:07 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-de; L...
2	404	87.245.244.151	0	12/Dec /2015:01:32:08 +0400	/favicon.ico	Mozilla/5.0 (Linux; Android 5.1.1; D6603 Build...

# Selection – Dataframes (Composite Conditions)

```
access_log_df.where(  
    (access_log_df.http_code <> '200') &  
    (access_log_df.user_agent.like('%Android%')))\n    • show(3)
```

http_code		ip	response_length	time	url	user_agent
0	404	176.120.130.254	0	12/Dec /2015:01:32:06 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-at; H...
1	404	93.188.131.176	0	12/Dec /2015:01:32:07 +0400	/favicon.ico	Mozilla/5.0 (Linux; U; Android 4.2.2; de-de; L...
2	404	87.245.244.151	0	12/Dec /2015:01:32:08 +0400	/favicon.ico	Mozilla/5.0 (Linux; Android 5.1.1; D6603 Build...

# Join

a dataframe  
access\_log

```
access_log = spark_session.read.table("web.access_log")
```

```
access_log.limit(3).show(3)
```

http_code		ip	response_length	time	url	user_agent
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...

# Join

another dataframe  
geoip

```
geoip = spark_session.read.table("web.geoip")
```

```
geoip.show(3)
```

http_code		ip	code	country
0	200	194.120.126.123	NL	Netherlands
1	200	94.126.119.173	FR	France
2	200	193.46.74.166	RU	Russian Federation

# Join

```
spark_session.sql("""
select *
from web.access_log l
join web.geoip g
on l.ip = g.ip
""").show(3)
```

http_code		ip	response_length	time	url	user_agent	code	country
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	RU	Russian Federation
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	GB	United Kingdom
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	RU	Russian Federation

```
access_log.join(geoip_df, on="ip")\
    . show(3)
```

http_code		ip	response_length	time	url	user_agent	code	country
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	RU	Russian Federation
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	GB	United Kingdom
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	RU	Russian Federation

```
access_log.join(geoip_df,
                 on = access_log.ip == geoip_df.ip)\
    . show (3)
```

http_code		ip	response_length	time	url	user_agent	code	country
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	RU	Russian Federation
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	GB	United Kingdom
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	RU	Russian Federation

# Join Types

```
access_log.join(geoip_df,  
                on = "ip",  
                how = "inner")\
```

```
access_log.join(geoip_df,  
                on = "ip",  
                how = "left")\
```

```
access_log.join(geoip_df,  
                on = "ip",  
                how = "left_semi")\
```

```
access_log.join(geoip_df,  
                on = "ip",  
                how = "left_anti")\
```

```
access_log.crossJoin(geoip_df)\  
                .count()
```



# Join Types – Left join recall

```
SELECT * FROM web.geo_ip LEFT OUTER JOIN  
web.access_log ON web.geo_ip.ip = web.accesslog.ip
```

geoip

ip

130.251.61.19

code

1

country

Italy

null null

access-log


ip

# Functions


```
spark_session.sql("""
    select user_agent,
           length (user_agent)
    from web.access_log
    limit 3
    """).show()
```

	user_agent	length (user_agent)
0	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	120
1	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	88
2	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	153


# Functions



mapping



generating



aggregating

# Mapping (aka scalar) functions

```
access_log.select("url",  
                  f.concat(f.lit("http://vk.com"),  
                           access_log_df.url)  
                  ).show(5)
```

	url	concat (http://vk.com, url)
0	/id53821	http://vk.com/id53821
1	/id33929	http://vk.com/id33929
2	/id35754	http://vk.com/id35754
3	/id78231	http://vk.com/id78231
4	/id39395	http://vk.com/id39395

# Aggregation

```
access_log = spark_session.read.table("web.access_log")
```

```
access_log.show(3)
```

http_code		ip	response_length	time	url	user_agent
0	200	109.106.133.8	21546	12/Dec /2015:01:31:46 +0400	/id53821	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...
1	200	46.31.82.254	8777	12/Dec /2015:01:31:47 +0400	/id33929	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...
2	200	193.124.254.46	8731	12/Dec /2015:01:31:48 +0400	/id35754	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...

# Aggregation

```
spark_session.sql("""  
    select url,  
           count(ip)  
    from web.access_log  
    group by url  
""").show(4)
```

	url	count(ip)
0	/id53821	3
1	/id33929	2
2	/id35754	1
3	/id11744	1

# Aggregation

```
import pyspark.sql.functions as f
```

```
access_log.groupBy("url")\  
    .agg(f.count("ip"))\  
    .show(3)
```

	url	count(DISTINCT ip)
0	/id37020	4
1	/id47695	2
2	/id77559	1

# Aggregation

```
access_log.groupBy("url")\  
    .agg({"ip": "count"})\  
    .show(3)
```

	url	count(DISTINCT ip)
0	/id37020	4
1	/id47695	2
2	/id77559	1



# Aggregation

```
access_log.groupBy(f.length("url"))\  
  .agg(f.count("*"))\  
  • show(4)
```

	length(url)	count(1)
0	31	5
1	34	4
2	28	4
3	27	7

# Aggregation – a single group

```
access_log.groupBy()\n    .agg(f.count("*"))\n    .show(4)
```

	count(ip)
0	89206

```
access_log.\n    .agg(f.count("*"))\n    .show(4)
```

	count(ip)
0	89206

```
SELECT COUNT(*)\nFROM web.access_log
```

# Aggregates

[1, 1, 2, 2, 2, 3, 3, 3, 3]

collect\_list = [1, 1, 2, 2, 2, 3, 3, 3, 3]

count = 9

sum = 20

**math**

min max avg var ...

**distinct**

collect\_set = [1, 2, 3]

countDistinct = 3

sumDistinct = 6

# Combining aggregation with join

```
access_log.join(geoip_df, on = "ip",)\n    .groupBy("country")\n    .agg(f.countDistinct("ip").alias("cnt"))\n    .orderBy(f.col("cnt").desc())\n    .show(3)
```

	country	cnt
0	Russian Federation	4556
1	France	1474
2	Germany	1287

Select country, count(distinct ip) as cnt  
From web.access\_log natural join web.geo\_ip  
Group by country  
Order by cnt desc

# Mapping functions

Remark non pure-relational model: list includes a list of values

```
access_log.select("user_agent")\
    .select("user_agent",
            f.split("user_agent", " ")
            .alias("list"))\
    • show(5)
```

1 2 3 4

[1, 2, 3, 4]

	user_agent	list
0	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	[Mozilla/5.0 (Macintosh; Intel...
1	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	[Mozilla/5.0 (Windows NT 5.1...
2	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	[Mozilla/4.0 (compatible; MSIE..
3	Mozilla/5.0 (Linux; Android 4.4.2; nb-no; SAMS...	[Mozilla/5.0 (Linux; Android...
4	Mozilla/5.0 (Linux; Android 4.4.2; nb-no; SAMS...	[Mozilla/5.0 (Linux; Android...

# Mapping functions

```
access_log.select("user_agent")\
    .select("user_agent",
            f.split("user_agent", " ")
            .alias("list")
            )\
    .select("user_agent",
            f.col("list")[0],
            f.col("list")[1],
            f.col("list")[2]
            )\
    .show(4)
```

	user_agent	list[0]	list[1]	list[2]
0	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	Mozilla/5.0	(Macintosh	Intel
1	Mozilla/5.0 (Windows NT 5.1; U; de; rv:1.9.1.6...	Mozilla/5.0	(Windows	NT
2	Mozilla/4.0 (compatible; MSIE 7.0; Windows NT...	Mozilla/4.0	(compatible	MSIE
3	Mozilla/5.0 (Linux; Android 4.4.2; nb-no; SAMS...	Mozilla/5.0	(Linux	Android

# Generating functions

```
access_log.select("user_agent")\
    .select("user_agent",
            f.split("user_agent", " ")
            .alias("list")
            )\
    .select("user_agent",
            f.explode("list"),
            )\
    .show(4)
```

t 1  
t 2  
t 3  
t 4  
t [1, 2, 3, 4]

	user_agent	col
0	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	Mozilla/5.0
1	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	(Macintosh;
2	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	Intel
3	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	Mac
4	Mozilla/5.0 (Macintosh; Intel Mac OS X 10_9_4)...	OS

# Generating functions

```
access_log.select("user_agent")\  
    .select("user_agent",  
            f.split("user_agent", " ")  
              .alias("list")  
            )\  
    .select("user_agent",  
            f.explode("list"),  
            )\  
    .where(f.col("col") == "Android")\  
    .show(4)
```

	user_agent	col
0	Mozilla/5.0 (Linux; Android 4.4.2; nb-no; SAMS...	Android
1	Mozilla/5.0 (Linux; Android 4.4.2; nb-no; SAMS...	Android
2	Mozilla/5.0 (Linux; Android 4.4.1; Caesar Buil...	Android



```
access_log.select(f.split("user_agent", " ").alias("words"))\
    .select(f.explode("words").alias("word"))\
    .groupBy("word")\
    .agg(f.count("*").alias("count"))\
    .orderBy(f.col("count").desc())\
    .show(10)
```

	words	count
0	Mozilla/5.0	75565
1	like	63791
2	Gecko)	58926
3	(KHTML,	58551
4	NT	50439
5	AppleWebKit/537.21	48648
6	Safari/537.36	48648
7	(Windows	37942
8	CLR	32140
9	.NET	31648

# more operations ...

- `when()`
- Partitions and Windows - OLAP SQL extensions
  - `over()`, `partitionBy()`

# Credits & References

- Yandex “Big Data Analysis” course
- Riccardo Torlone, Università Roma Tre, Roma, slides for the “Big Data” course
- Paolo Garza, Politecnico di Torino, slides for the “ Big data ” course
- Reference:  
Chapter 9 – Spark SQL

