

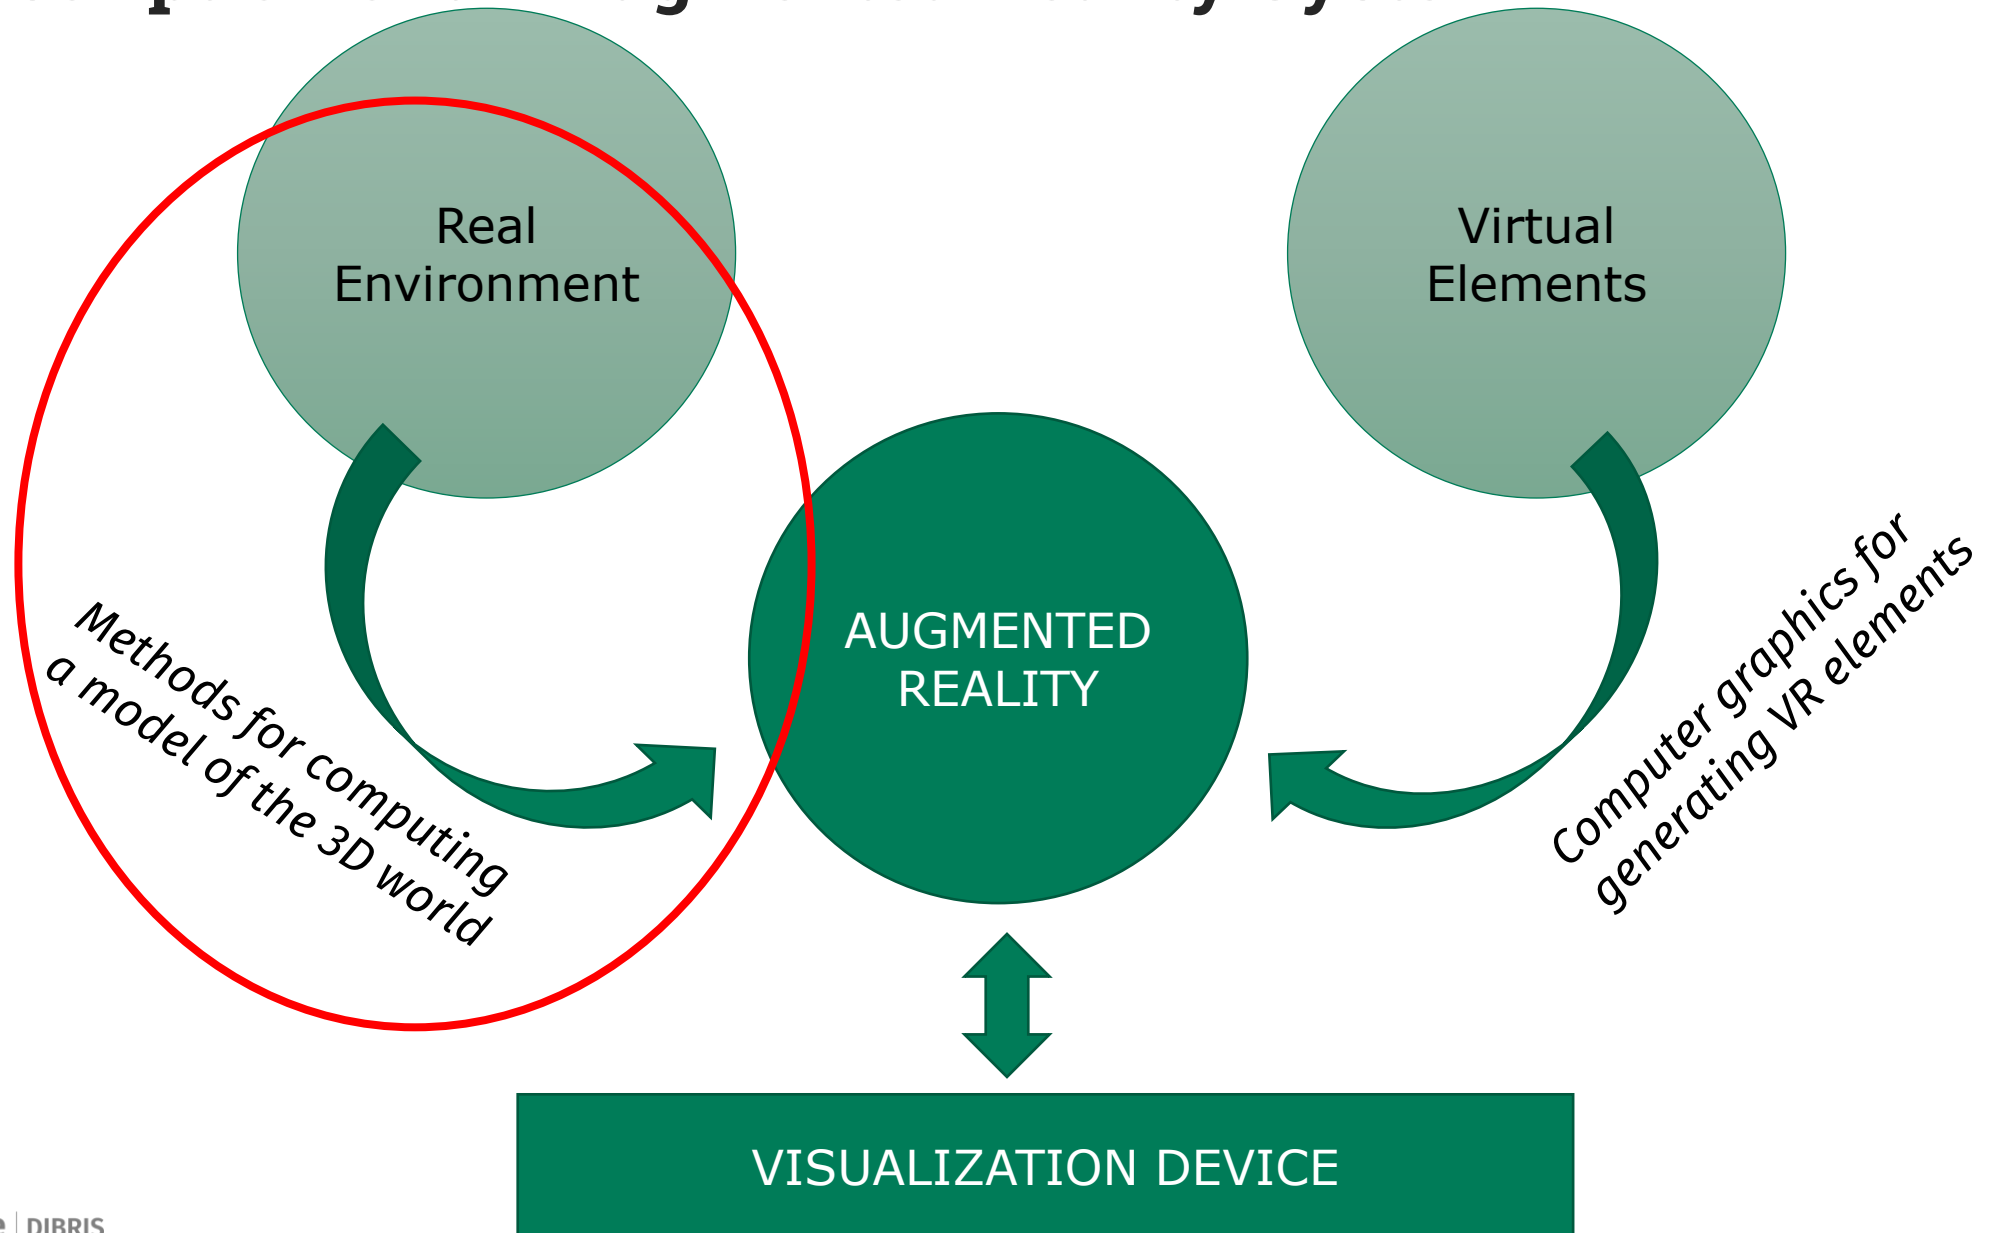
Augmented Reality

Lecture 8 – image formation and edge detection

Manuela Chessa – manuela.chessa@unige.it

Fabio Solari – fabio.solari@unige.it

Description of an Augmented Reality System




CG, AR and 3D CV



- Computer Graphics (CG): creating an image from scratch by using a 3D model.
- Computer Vision (CV): understanding the “content” of an image, usually for estimating a “3D model” of the depicted scene.
- Augmented Reality (AR): using CG and CV to blend real and virtual contents in a coherent way.

CG, AR and 3D CV

Computer Graphics: the **unknown** is the **2D image**, i.e. to create the 2D image(s) from 3D model and camera position





Computer Vision: the **unknown** is the real camera position in the world, i.e. to estimate the **3D camera pose** from 2D image(s), then the **3D scene**




CG, AR and 3D CV

Computer Vision: we infer 3D models



Computer Graphics: we know 3D models







w: world
c: camera
CG: computer graphics

CG, AR and 3D CV

User tracking

Head pose

Hand pose





To change the view of VR world



3D Computer Vision for AR


The *unknown* is the *real camera position* in the world (i.e., *pose estimation problem*). There are several approaches:

1. Approaches where **3D models are available**:
 - classical pose estimation method (this is an *inverse problem* and uses images)




3D Computer Vision for AR

2. Approaches where **3D models are not available** (*by using images*):
 - 3D models can be estimated on-line (by using *RGB cameras*) thanks to Simultaneous Localization and Mapping (SLAM) techniques



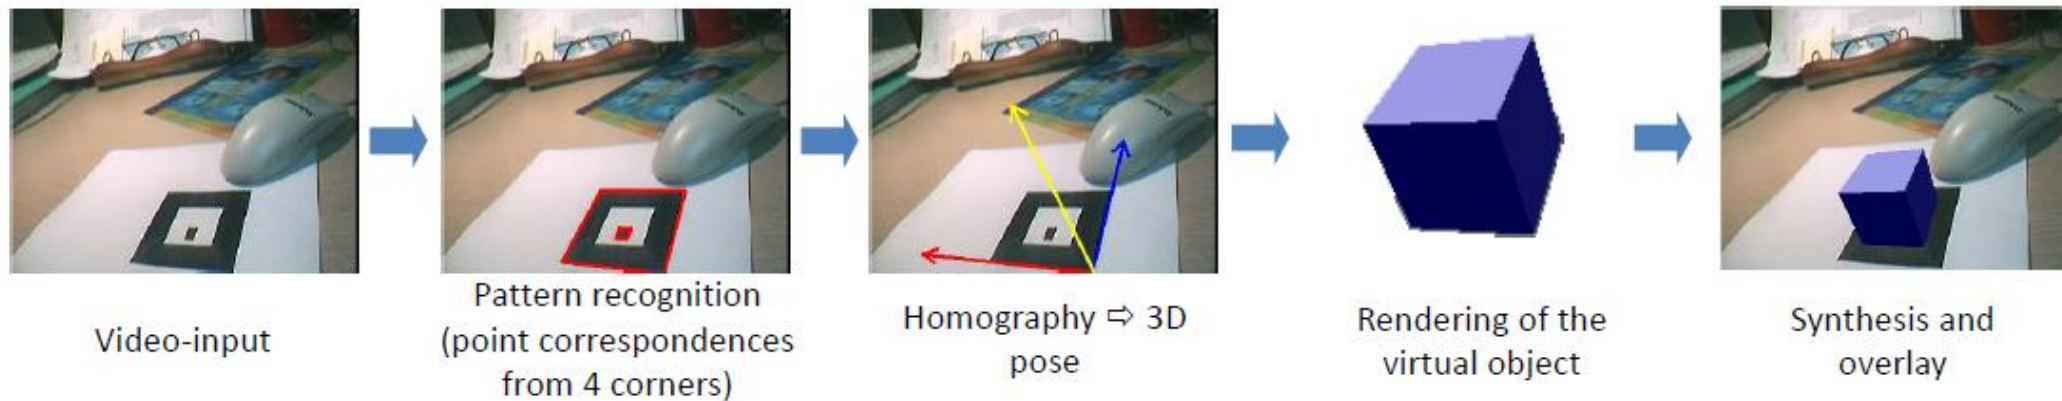
3D Computer Vision for AR

3. Approaches where **3D models are not available** (*by using point cloud*):
- when 3D data can be directly measured (by *RGB-D cameras*), registration can be done directly in the 3D space, e.g. by Iterative Closest Point (ICP) technique




3D Computer Vision for AR

4. Approaches where **2D models are available** (*markers*): the pose estimation problem can be simplified when the **scene is planar** (*homography*)



3D Computer Vision for AR

- From a practical point of view, the development of **actual AR** applications rises the question of the **features extraction and matching** and of the **3D reconstruction**: we consider **calibration**, **registration** and **tracking**.



Summary

- Image processing:
 - image formation and edge detection
- Computer Vision:
 - image segmentation, keypoints (corners), stereopsis, disparity computation
- 3D Computer Vision for AR:
 - camera calibration, pose estimation, epipolar geometry, RGBD camera
 - camera tracking, homography, SLAM, SPAAM






**Università
di Genova**

DIBRIS DIPARTIMENTO
DI INFORMATICA, BIOINGEGNERIA,
ROBOTICA E INGEGNERIA DEI SISTEMI

Image formation and edge detection

Introduction

- Vision is deceptively easy:
 - vision is immediate for us
 - we perceive the visual world as external to ourselves, but it is a **reconstruction within our brain**



- Vision is computationally demanding

Introduction

What do you see?



perceived as
a 3D scene

but really just a planar
surface (screen) under
nonuniform lighting
(projector).

- It is something mathematically impossible: **recovering 3D from a single image is a mathematically inverse ill-posed problem.**
- How?
- You used **assumptions** based on prior knowledge/experience about the way the world works

Introduction

The goal of computer vision: **to bridge the gap** between “meaning” and pixels



What we see

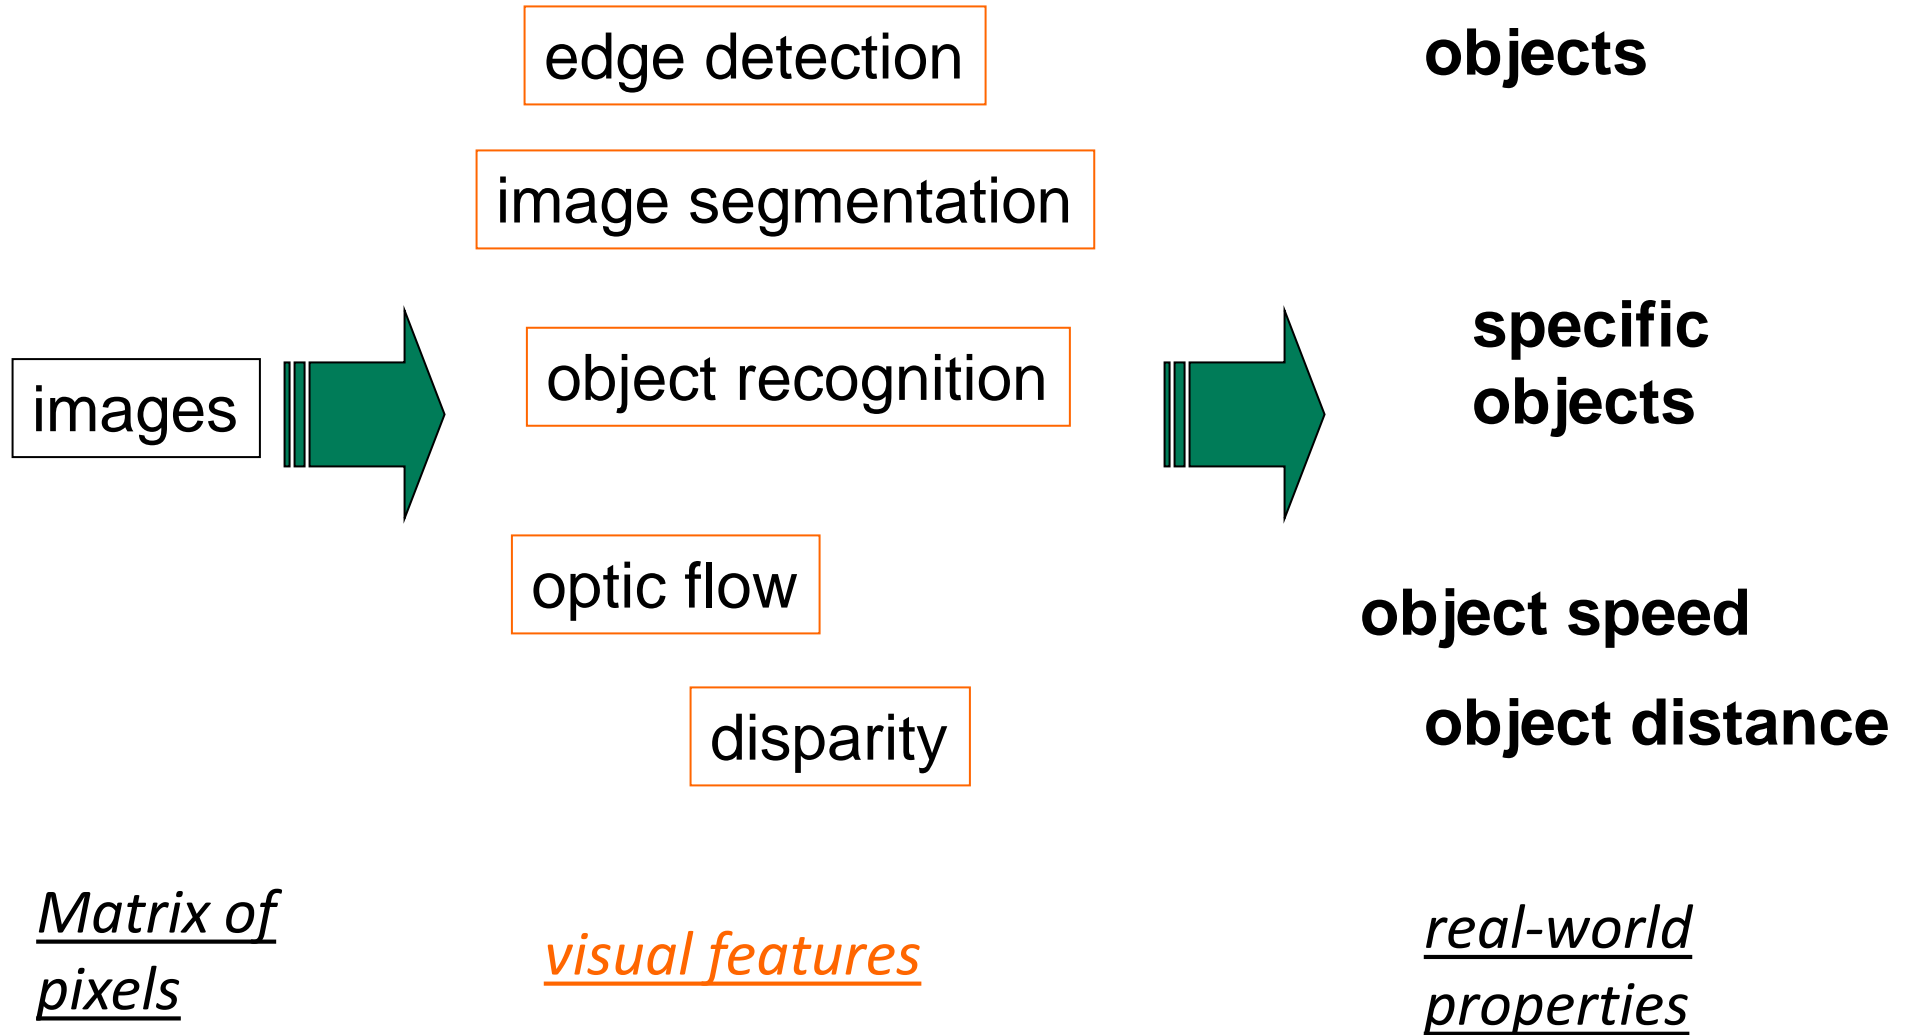
0	3	2	5	4	7	6	9	8
3	0	1	2	3	4	5	6	7
2	1	0	3	2	5	4	7	6
5	2	3	0	1	2	3	4	5
4	3	2	1	0	3	2	5	4
7	4	5	2	3	0	1	2	3
6	5	4	3	2	1	0	3	2
9	6	7	4	5	2	3	0	1
8	7	6	5	4	3	2	1	0

What a computer sees

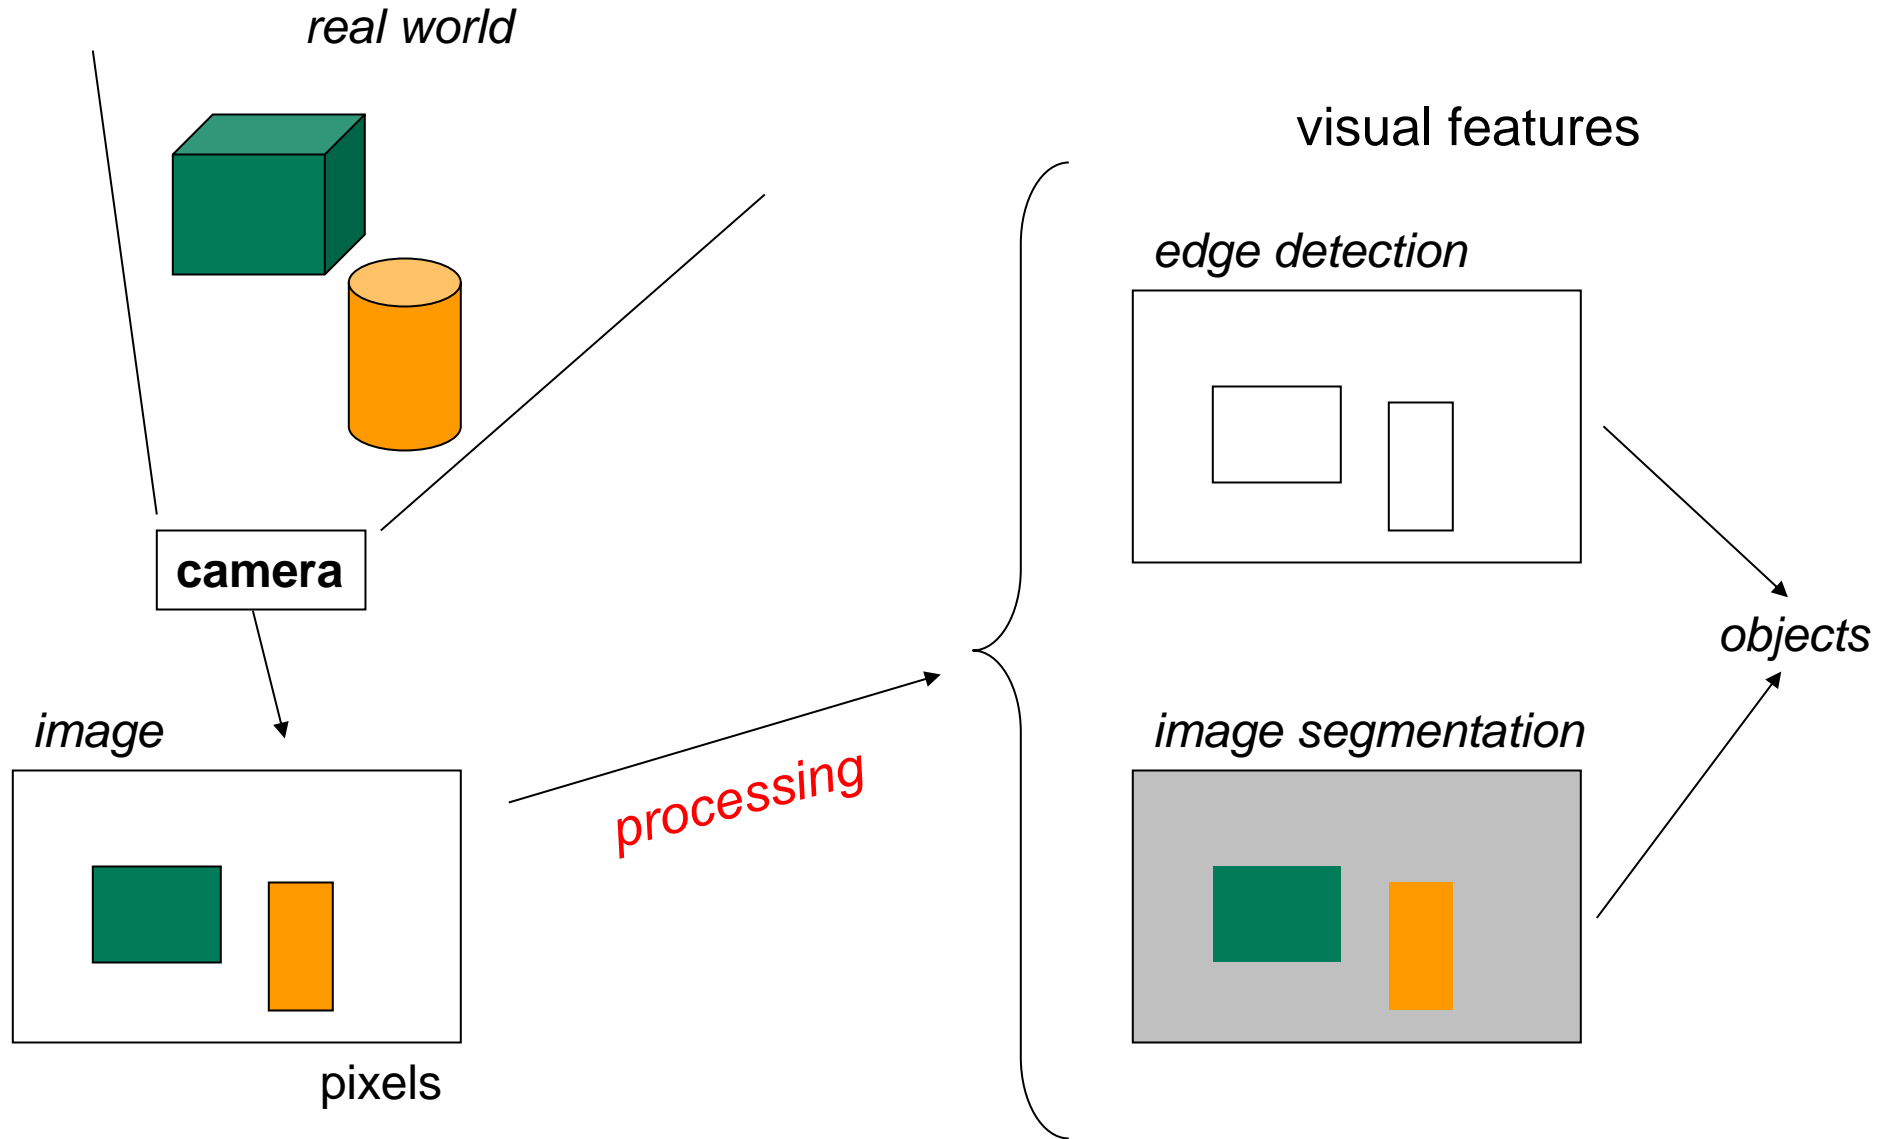
Introduction

- To interact with the real-world, we must tackle the **problem of inferring 3-D information of a scene from a set of 2-D images**.
- In general, this problem falls into the category of so-called *inverse problems*, which are prone to be *ill-conditioned* and difficult to solve in their full generality unless *additional assumptions* are imposed.
- Before we address how to reconstruct 3-D geometry from 2-D images, ***we first need to understand how 2-D images are generated and processed.***

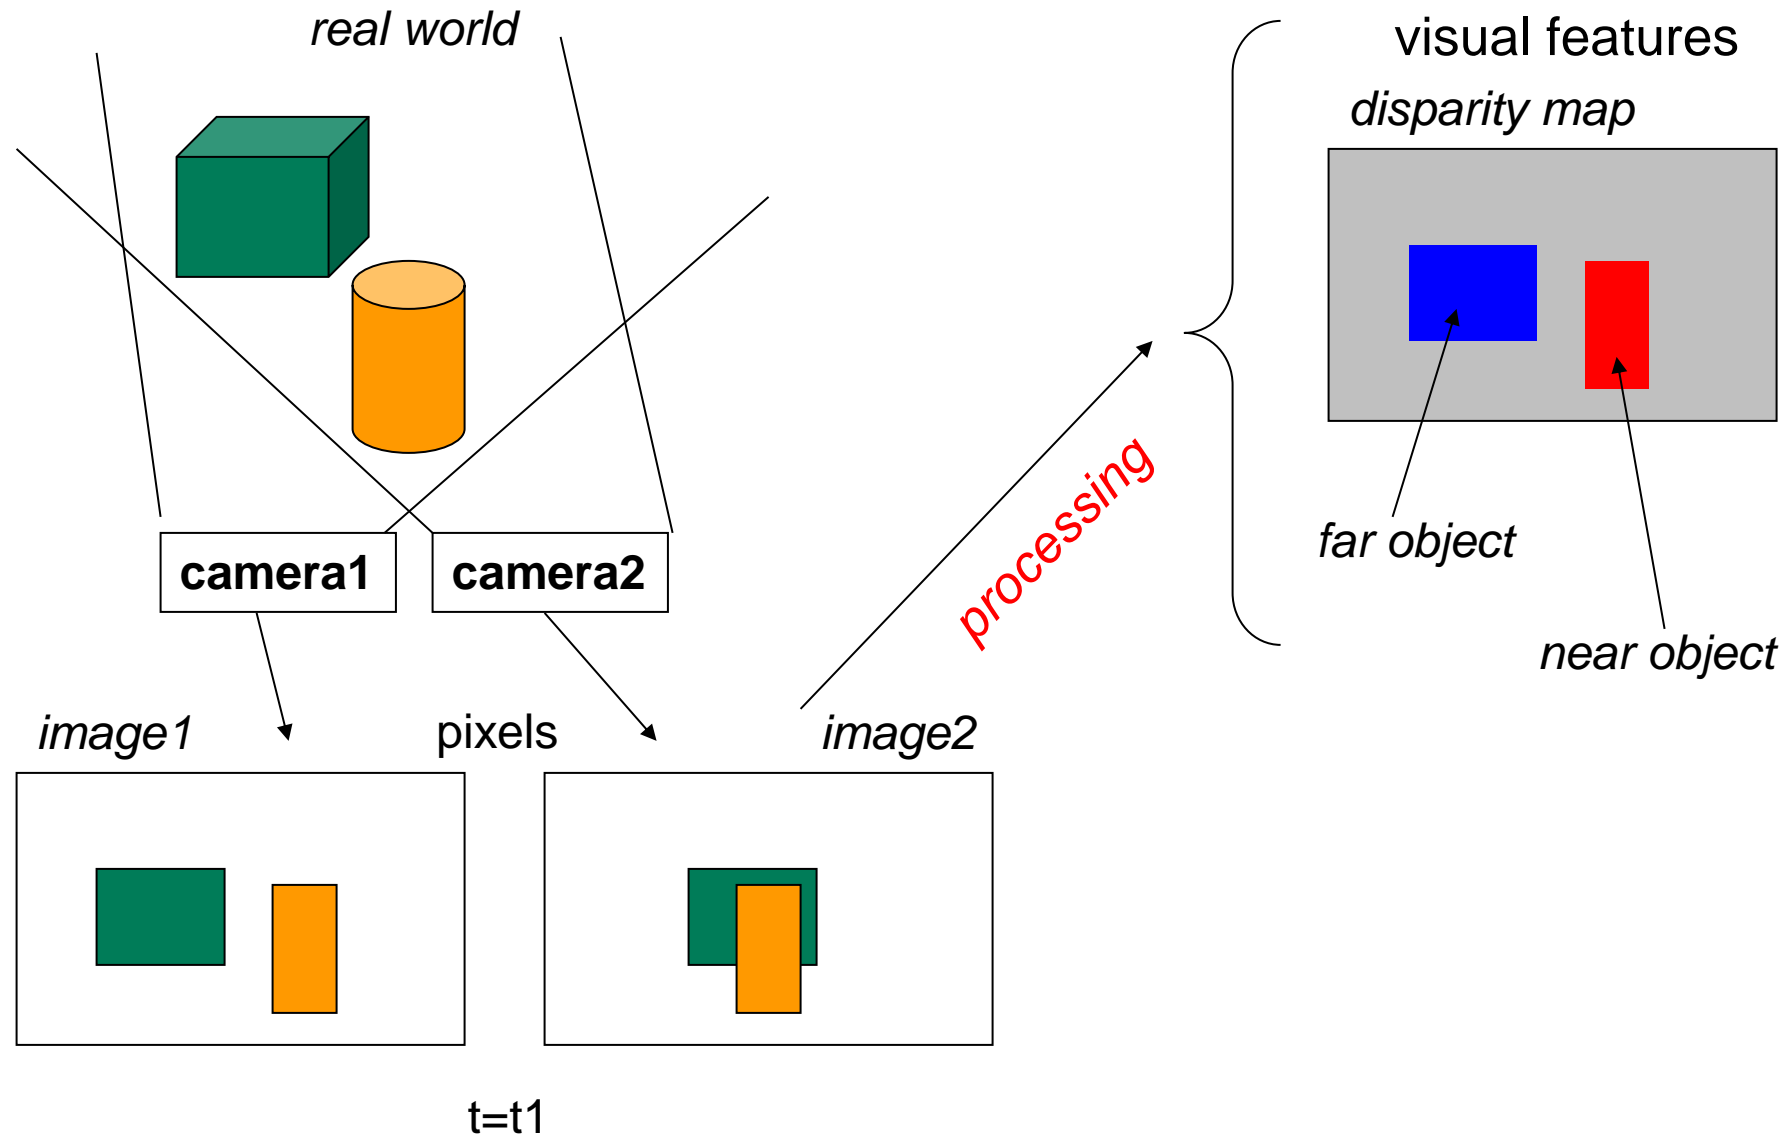
Computer Vision



Computer Vision



Computer Vision



Computer Vision

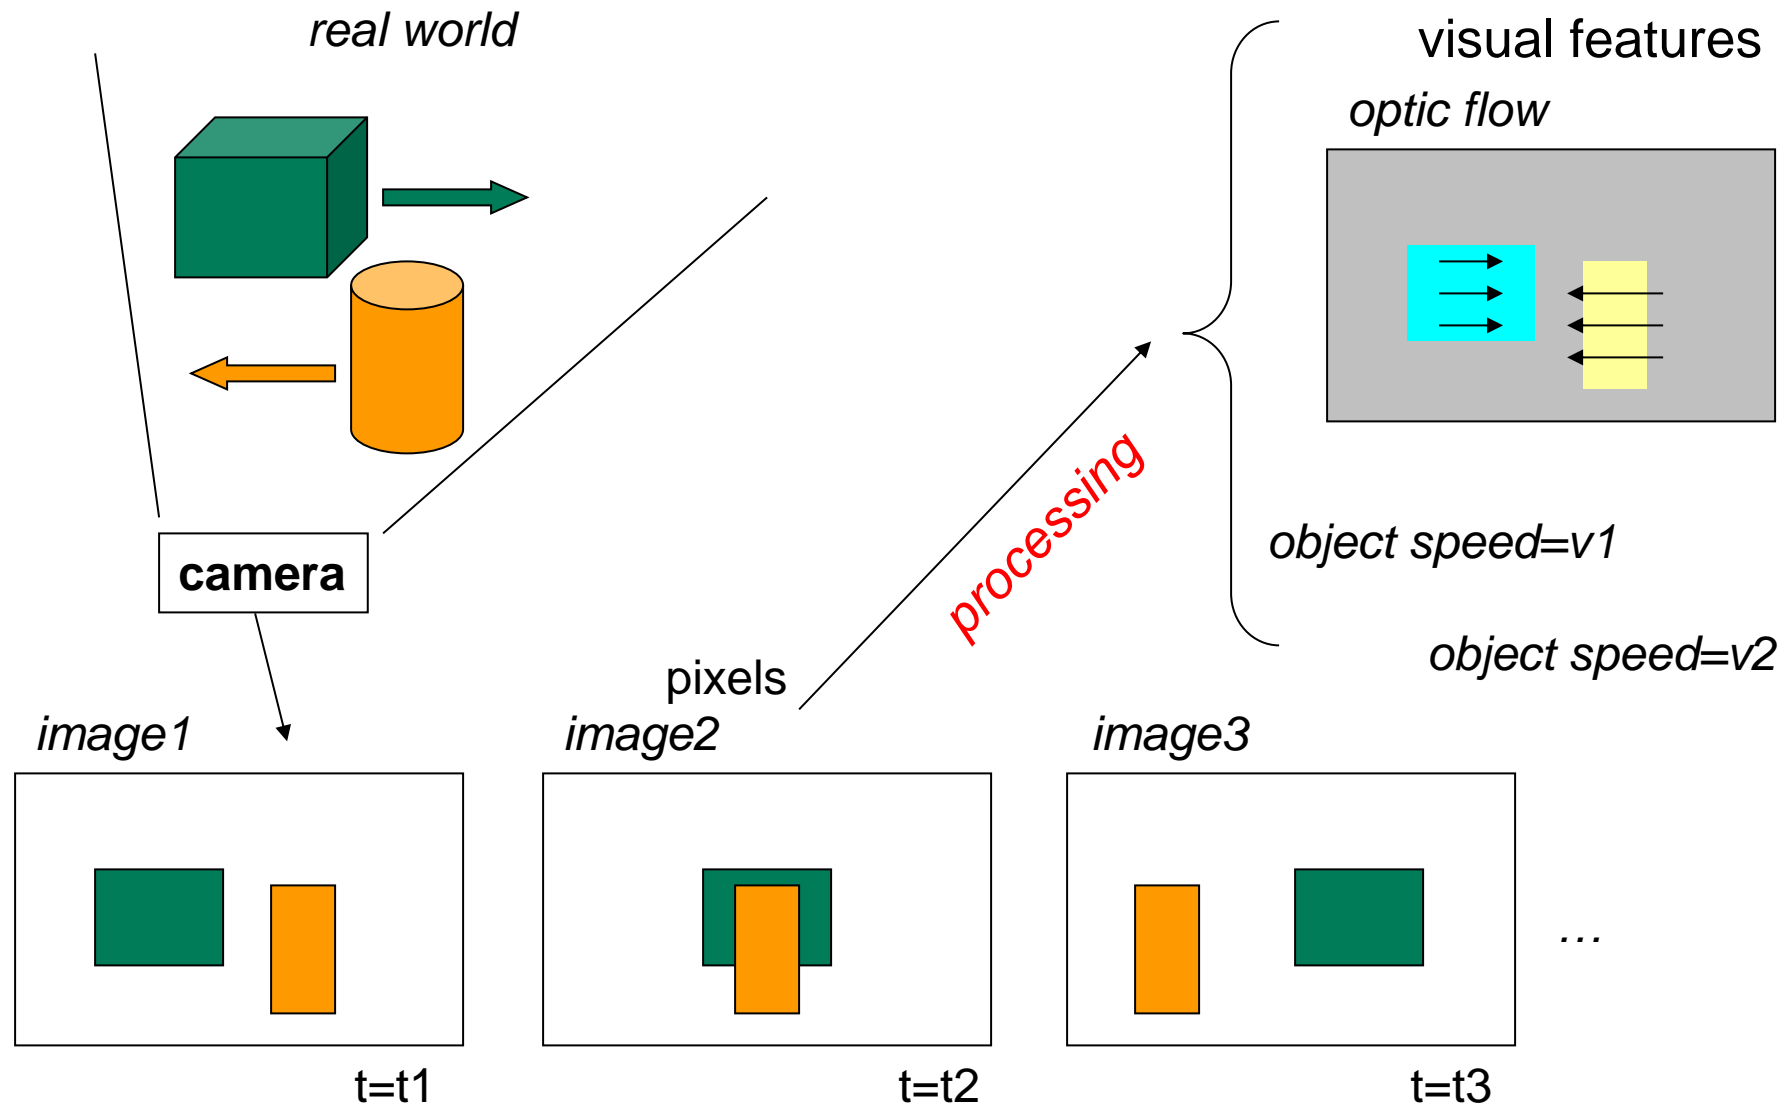


Image formation

Photometry overview




Image formation

Photometry overview

Visual signal $s(x,y)$:

$$s: \mathbb{R}^2 \rightarrow \mathbb{R}$$

- $s(x,y)$ is a **two-dimensional function**: x and y are spatial coordinates.
- The **amplitude** of s is called light **intensity** or gray level at the point (x, y) .

$$s(x,y) = il(x,y) \cdot r(x,y)$$

$s(x,y)$: **intensity** at the point (x,y)

$il(x,y)$: **illumination** at the point (x,y)
(*the amount of source illumination incident on the objects*)

$r(x,y)$: **reflectance** at the point (x,y)
(*the amount of illumination reflected/transmitted by the object*)

Where $0 < il(x,y) < \infty$ and $0 < r(x,y) < 1$

Image formation

Photometry overview

- Illumination

Lumen: a unit of light flow or luminous flux

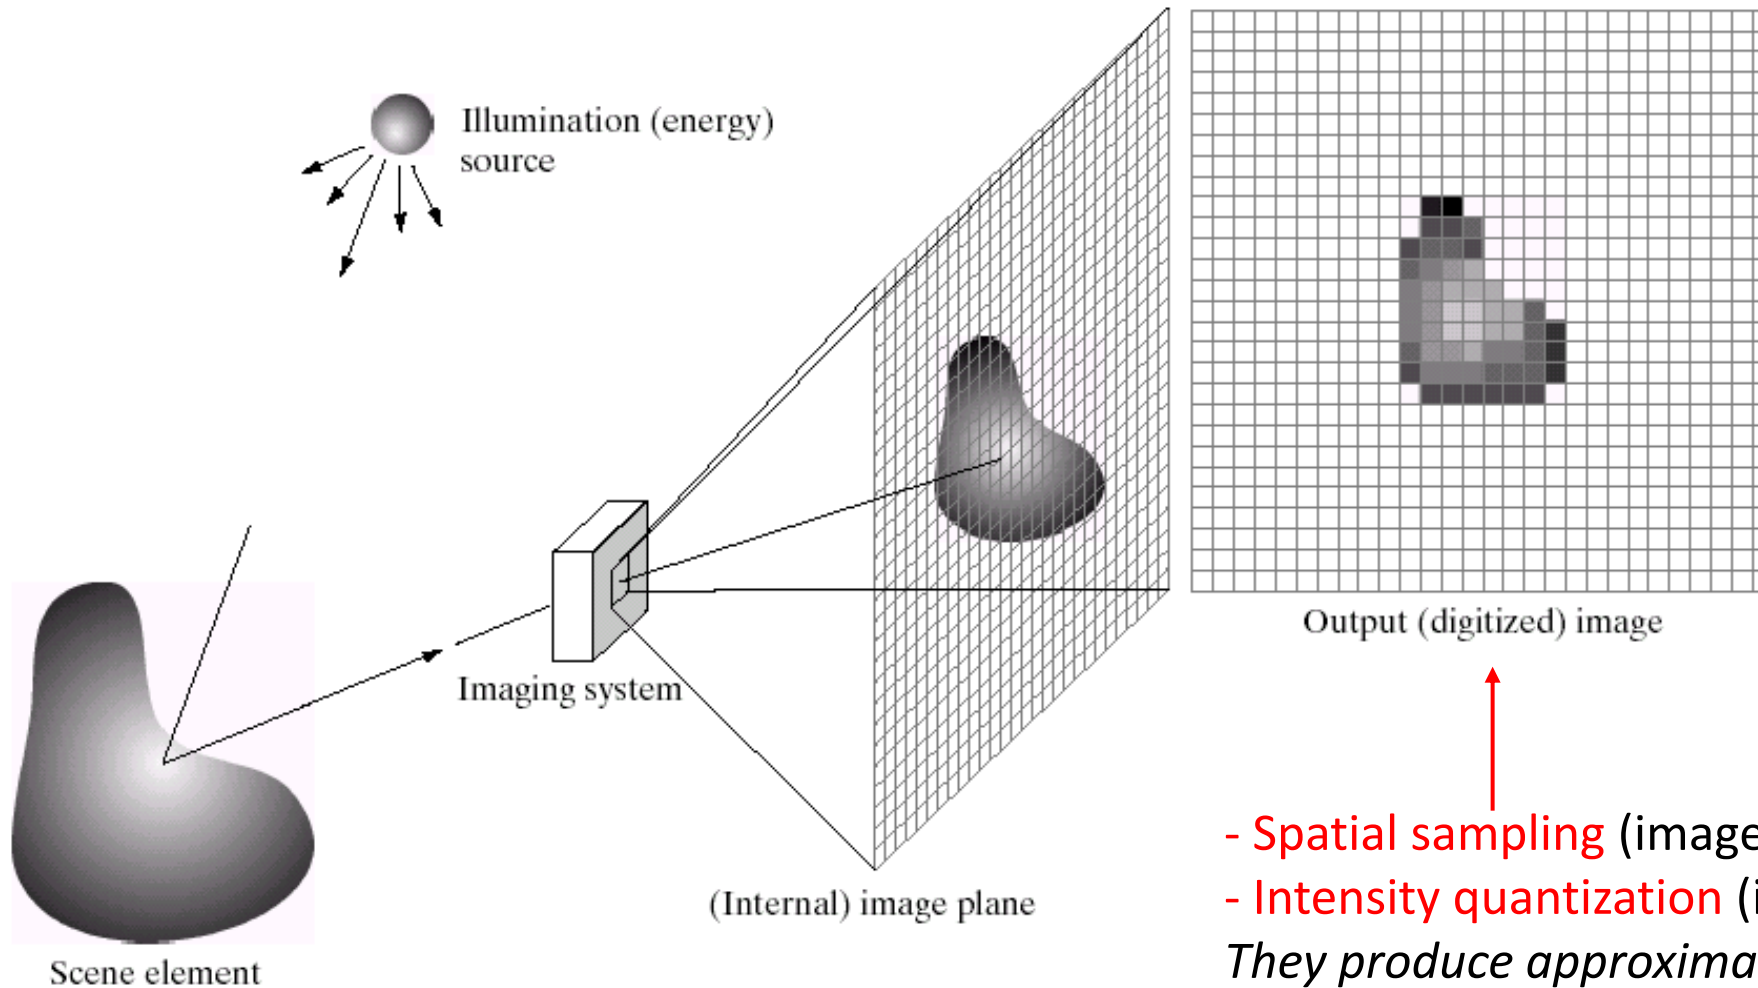
Lumen per square meter (lm/m^2) — The metric unit of measure for illuminance of a surface

- On a clear day, the sun may produce in excess of $90,000 \text{ lm/m}^2$ of illumination on the surface of the Earth
- On a cloudy day, the sun may produce less than $10,000 \text{ lm/m}^2$ of illumination on the surface of the Earth
- On a clear evening, the moon yields about 0.1 lm/m^2 of illumination
- The typical illumination level in a commercial office is about 1000 lm/m^2

- Reflectance

- 0.01 for black velvet
- 0.65 for stainless steel
- 0.80 for flat-white wall paint
- 0.90 for silver-plated metal
- 0.93 for snow


Image formation



- **Spatial sampling** (image resolution)
 - **Intensity quantization** (image gray levels)
- They produce approximation (errors)*

Image formation

- Digital image formation is the first step in any digital image processing application.
- The *digital image formation system* (camera) consists basically of the *optical system*, the *sensor* and the *digitizer*.



The effect of the recording process is the addition of a noise contribution.
The recorded image i is called noisy image.

Optical system

- The *optical system* can be modeled as a ***linear shift invariant system*** having a two-dimensional impulse response $h(x,y)$.
- The **input-output relation** of the optical system is described by a **2D convolution** (both signals s and f represent optical intensities):


$$f(x, y) = \iint s(\xi, \eta) h(x - \xi, y - \eta) d\xi d\eta$$

Optical system

- The two-dimensional impulse response $h(x,y)$ is also called PSF (***point spread function***), and its Fourier transform is called ***transfer function***.
 - Linear motion blur: it is due to the relative motion, during exposure, between the camera and the object being photographed, $H(\mathbf{w}) \sim \exp(jT\mathbf{v}\mathbf{w})\text{sinc}(T\mathbf{v}\mathbf{w})$.
 - Out-of-focus blur: $H(\mathbf{w}) \sim (1/D|\mathbf{w}|)^{3/2} \cos(D|\mathbf{w}|)$.
 - Atmospheric turbulence blur: $H(\mathbf{w}) \sim \exp(-s^2 |\mathbf{w}|^2)$.
 - No blur: $h(x, y) = \delta(x, y)$.



Motion blur →




Camera model


- **Images are two-dimensional patterns of brightness values.**
They are formed by the **projection** of **3D objects** on the camera **image plane**.
- Basic abstraction is the **pinhole camera**:
 - Lenses required to ensure image is not too dark.
 - Pinhole camera model works in practice.

Pinhole camera


Abstract camera model



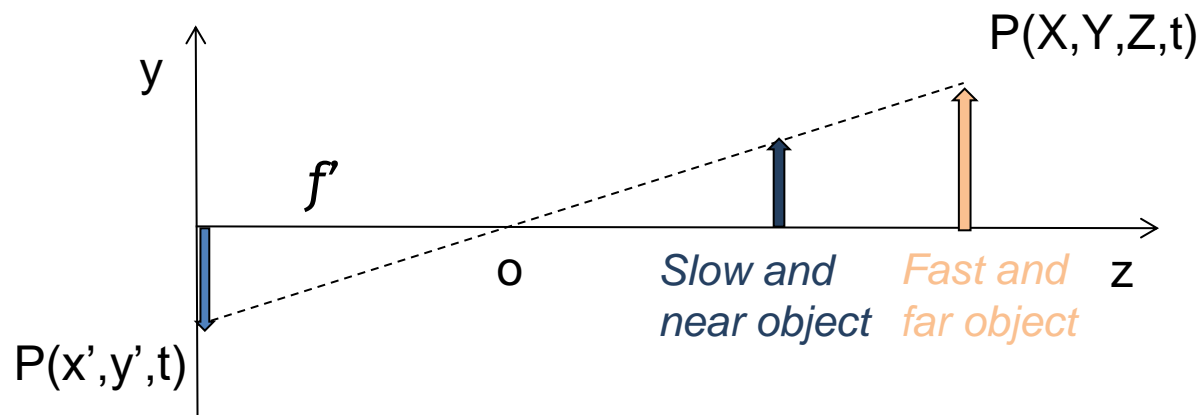

Animal eye



Pinhole Perspective Equation




Pinhole camera



Inverse problems, which are prone to be ill-conditioned


$$\begin{cases} x' = f' \frac{X}{Z} \\ y' = f' \frac{Y}{Z} \end{cases}$$

Cameras with lenses



- A lens focuses parallel rays onto a single focal point
- Gather more light, while keeping focus; **make pinhole perspective projection practical**

Cameras with lenses




Thin lens equation

$$\frac{1}{f} = \frac{1}{u} + \frac{1}{v}$$

- Any **object** point satisfying this equation is **in focus**


Cameras with lenses

- Depth of field: distance between image planes *where blur is tolerable*




Cameras with lenses

- Depth of field: a *smaller aperture* increases the range in which the object is approximately in focus




Cameras with lenses

- In VR/AR, the depth of field is a post-processing effect *to be implemented*




Without depth of field



With depth of field (*Unity*)




- In SAR, the depth of field is an effect *to be compensated*



Huang, B., & Ling, H. (2021). Deprocams: Simultaneous relighting, compensation and shape reconstruction for projector-camera systems. *IEEE Transactions on Visualization and Computer Graphics*, 27(5), 2725-2735.

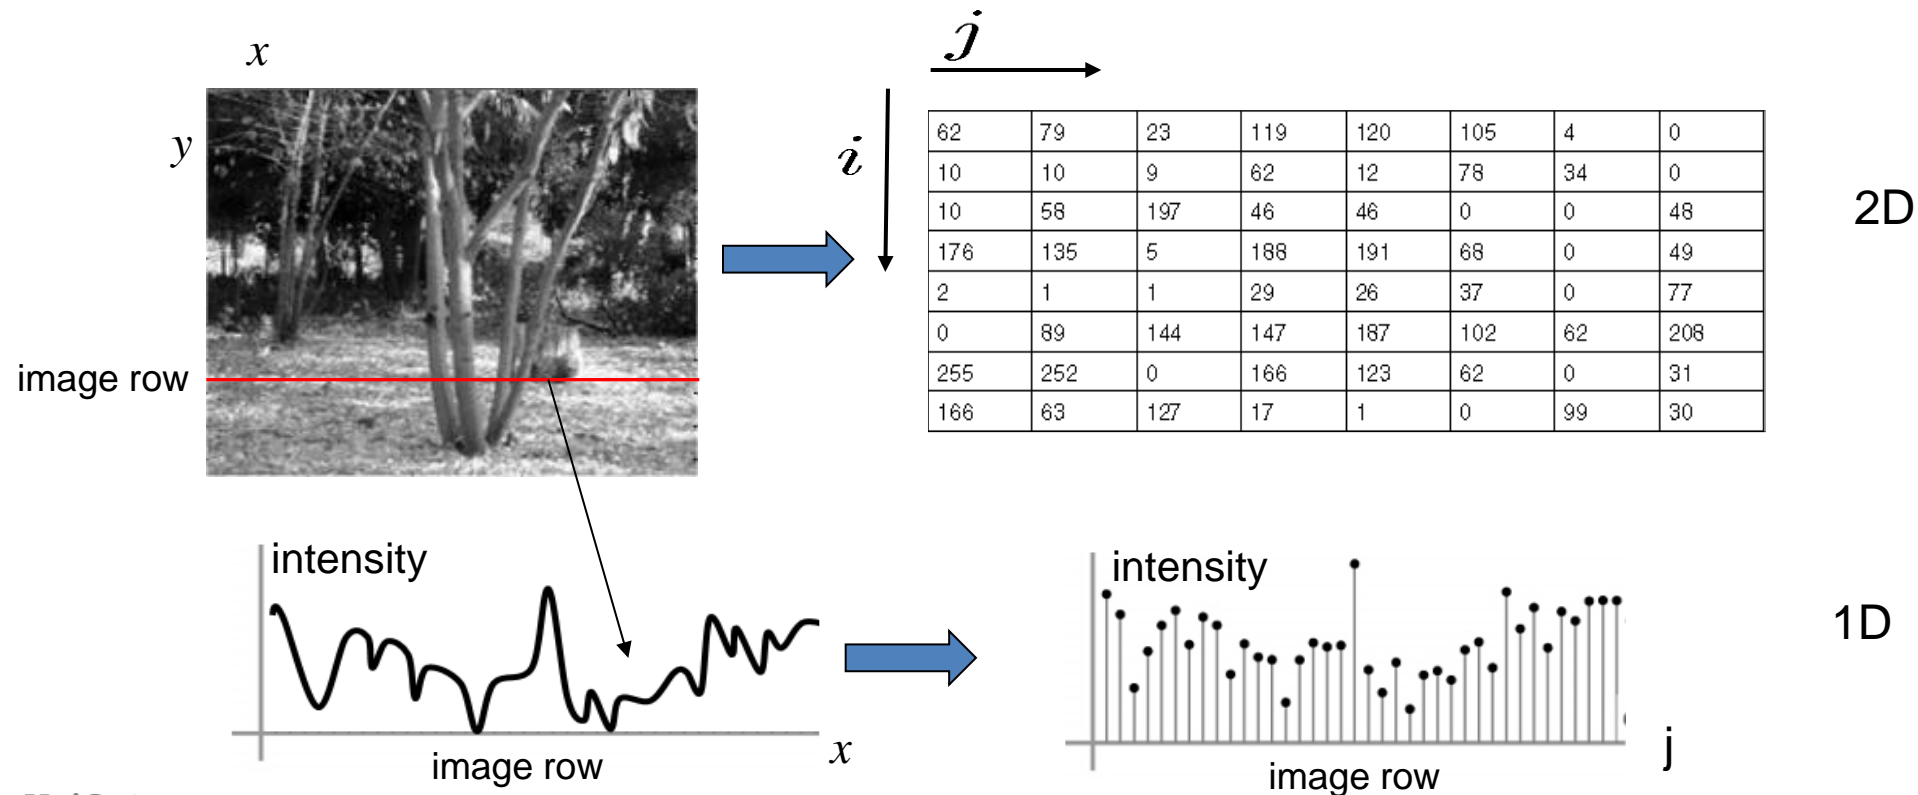
Cameras with lenses

- Field of view (FOV): Angular measure of portion of 3D space seen by the camera
- As **f gets smaller**, image becomes more *wide angle*
 - more world points project onto the finite image plane d
- As **f gets larger**, image becomes *more telescopic*
 - smaller part of the world projects onto the finite image d plane



Digital images representation

- In Computer Vision we operate on **digital (discrete)** images:
 - **To sample** the 2D space on a regular grid (**spatial resolution**)
 - **To quantize** the amplitude of each sample, e.g. round to nearest integer, (**gray levels**)
- Image is represented as a *matrix of integer values (intensity)*.



Digital images representation

Image **sampling**



200x200



100x100



50x50



25x25 (pixels)

**Spatial
resolution**
(pixels)

Image **quantization**



8 bits



5 bits



4 bits



3 bits



2 bits




1 bit
(2 levels)

**Intensity
resolution**
(gray levels)

Digital images representation

Human retina: **three different receptors**



Each pixel is specified
by three values (RGB)

We perceive

Monochrome image




$$R[x,y] = G[x,y] = B[x,y]$$

Color image (RGB)




Digital images representation

- A **truecolor image** is an image in which each pixel is specified by three values (RGB), one each for the **red**, **blue**, and **green** components of the pixel's color.



- An **indexed image** consists of an array and a *colormap matrix* (LUT). The pixel values in the array are direct indices into a colormap. The colormap matrix is an m-by-3 array of values in the range [0,1]. Each row of map specifies the red, green, and blue components of a single color.



Digital image formation system

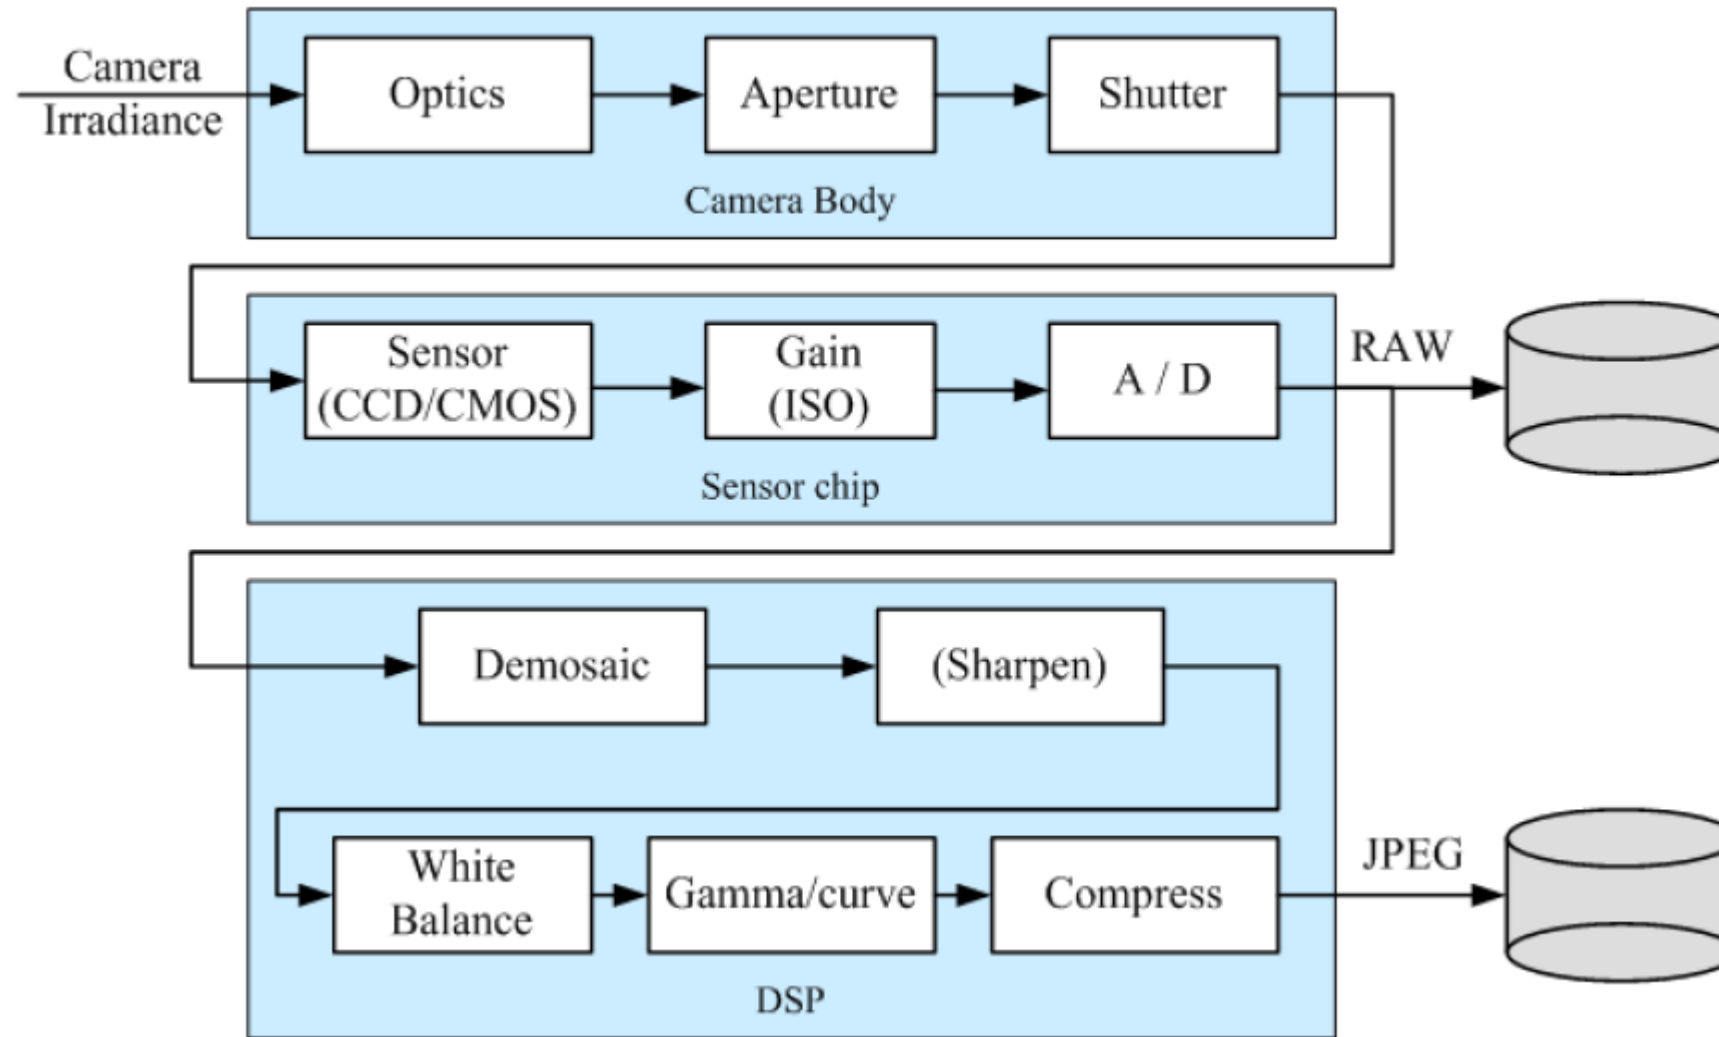



Image features: edges

- Edges: **discontinuities in intensity**
 - Boundaries of material properties
 - Boundaries of objects
 - Boundaries of lighting




- Edge detection:
- **Goal:** **Identify** visual changes (**discontinuities**) in an image.
- **Why?** Intuitively, semantic information is encoded in edges (e.g. to recognize objects and to *recover geometry and viewpoint*).



Image features: edges

- Discontinuities in signal can be **detected** by computing the **derivative of the signal**.



- In particular, we compute the image gradient:

$$\nabla I(x, y) = \left[\frac{\partial I}{\partial x}, \frac{\partial I}{\partial y} \right]^T = [I(x, y)_x, I(x, y)_y]^T = [I_x, I_y]^T$$

Plotted as a **vector field**, the *gradient* vector at each pixel *points “uphill”*.

The gradient indicates the direction of steepest ascent.







Image features: image gradient

The gradient of an image:

$$\nabla f = \left[\frac{\partial f}{\partial x}, \frac{\partial f}{\partial y} \right]$$

The gradient points in the direction of most rapid change in intensity



The gradient direction (*orientation of edge normal*) is given by:

$$\theta = \tan^{-1} \left(\frac{\partial f}{\partial y} / \frac{\partial f}{\partial x} \right)$$

The *edge strength* is given by the gradient magnitude

$$\|\nabla f\| = \sqrt{\left(\frac{\partial f}{\partial x} \right)^2 + \left(\frac{\partial f}{\partial y} \right)^2}$$

Image features: numerical derivatives

- We can compute the gradient vector at each pixel by convolving image with horizontal and vertical derivative filters.

- By considering the *finite forward difference* and $h=1$ for images, we have:

$$\frac{\partial f(x, y)}{\partial x} \approx \frac{f(x+1, y) - f(x, y)}{1}$$

the partial derivative with respect to columns as


$$img(i, j+1) - img(i, j)$$

- It can be considered as a convolution with the kernel $[1, -1]$

1	-1
---	----


Image features: numerical derivatives

$I(x, y)$






$$I_x = \frac{\partial I(x, y)}{\partial x}$$

1	-1
---	----



$$I_y = \frac{\partial I(x, y)}{\partial y}$$

1
-1



Simple Edge Detection Using Gradients

- Issue: the noise
 - **smooth before differentiation**
 - two convolutions: to smooth, then to differentiate
 - actually, we can use the derivative of the kernel $\rightarrow \frac{\partial}{\partial x}(f * g) = \frac{\partial f}{\partial x} * g$

For instance, *Prewitt's operator*:

$$\longrightarrow \begin{bmatrix} 1 & 1 \\ 1 & 1 \\ 1 & 1 \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & -1 \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 0 & -1 \\ 1 & 0 & -1 \\ 1 & 0 & -1 \end{bmatrix}$$

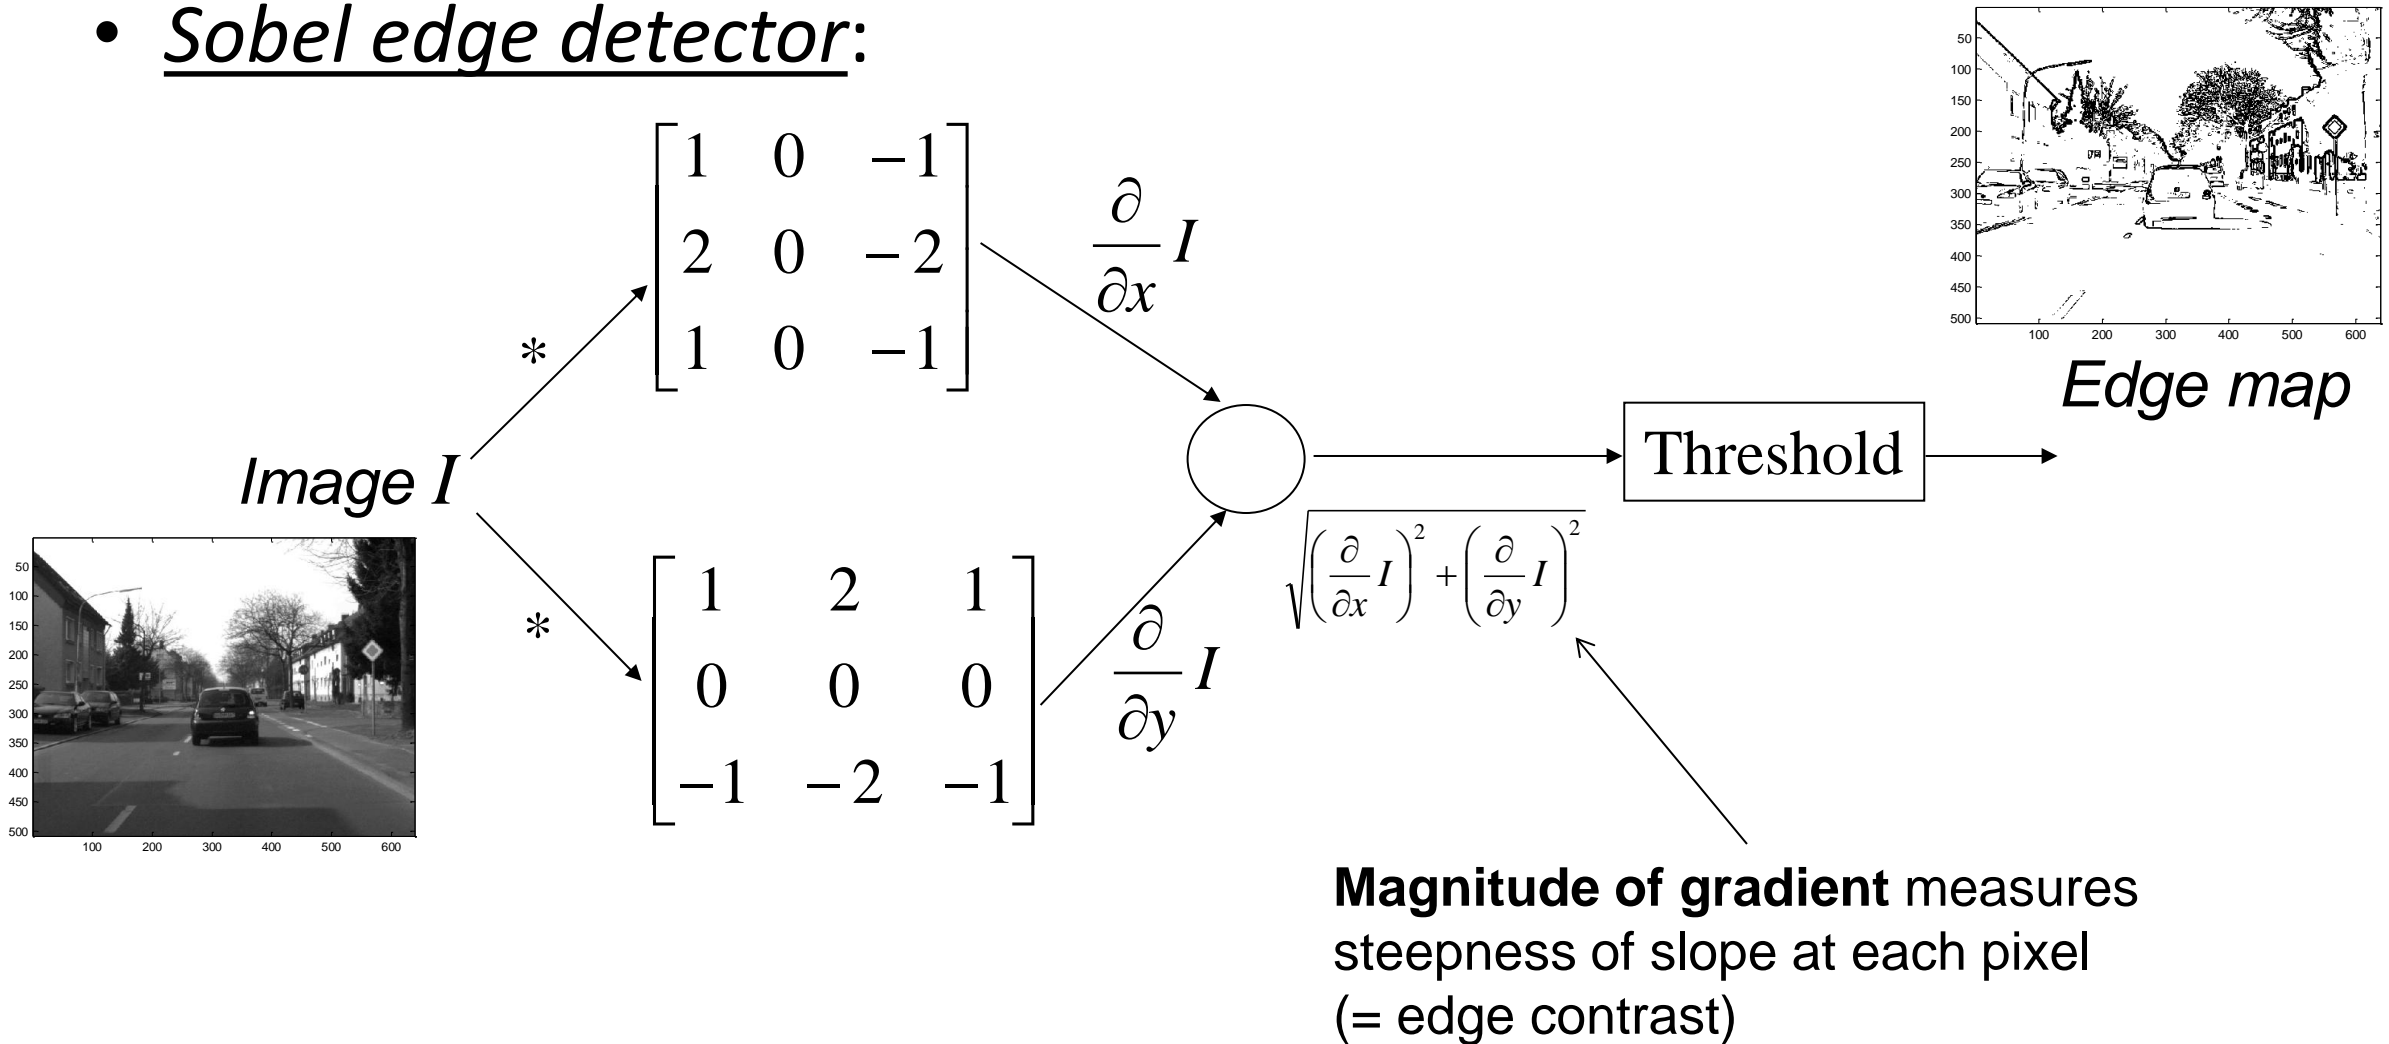
Smooth

Differentiate

$$\longrightarrow \begin{bmatrix} 1 & 1 & 1 \\ 1 & 1 & 1 \end{bmatrix} \longrightarrow \begin{bmatrix} 1 \\ -1 \end{bmatrix} \longrightarrow \begin{bmatrix} 1 & 1 & 1 \\ 0 & 0 & 0 \\ -1 & -1 & -1 \end{bmatrix}$$

Simple Edge Detection Using Gradients


- Sobel edge detector:




Sobel edge detector: Matlab implementation

```
dx=[1 0 -1; 2 0 -2; 1 0 -1];%mask
dy=[1 2 1; 0 0 0; -1 -2 -1];
tmprgb=imread('left_#290.bmp','bmp');
tmp=rgb2gray(tmprgb);
I=double(tmp);
figure,imagesc(I),colormap gray
Ix=conv2(I,dx,'same');%numerical derivatives
Iy=conv2(I,dy,'same');
figure,imagesc(Ix),colormap gray
figure,imagesc(Iy),colormap gray
```




```
M=sqrt(Ix.^2 + Iy.^2);%magnitude
figure,imagesc(M),colormap gray, colorbar, title('Magnitude map')
I_edge=M>100; %threshold
figure,imagesc(~I_edge),colormap gray,title('Edge map')%binary image
```



$$\frac{\partial}{\partial x} I$$



$$\frac{\partial}{\partial y} I$$



Sobel edge detector: Python implementation

```
import numpy as np

from scipy import signal

from PIL import Image

from matplotlib import pyplot as plt

im = Image.open('left_#290.bmp')

img = im.convert("L")

plt.figure()

plt.imshow(img,cmap='gray')

kh = np.asarray([[1, 0, -1], [2, 0, -2], [1, 0, -1]])

kv = np.asarray([[1, 2, 1], [0, 0, 0], [-1, -2, -1]])

gx = signal.convolve2d(img, kh, mode="same",boundary="symm", fillvalue=0)

gy = signal.convolve2d(img, kv, mode="same",boundary="symm", fillvalue=0)

plt.figure()

plt.imshow(gx,cmap='gray')

plt.figure()

plt.imshow(gy,cmap='gray')
```

```
g = np.sqrt(gx * gx + gy * gy)

g *= 255.0 / np.max(g)

plt.figure()

plt.imshow(g, cmap='gray')

plt.figure()

plt.imshow(255-g>210, cmap='gray')

plt.show()
```

Then ... to use libraries ...

```
sobelmap = edge(I, 'sobel');

import cv2
sobelx = cv2.Sobel(img,cv2.CV_64F,1,0,ksize=3)
```

ARCore SDK for Unity

