


Convolutional Neural Network

Deep Learning




A refresh

Deep Neural Network Recap




The role of a neuron



$$y = f(\mathbf{x}) = \sigma\left(\sum_i w_i x_i + b\right)$$

The role of a neuron



- Each neuron is connected to all the others
- Correlations between input are not taken into account
- As the size of the input and the depth of the architecture increase, the number of parameters increases dramatically

$$y = f(\mathbf{x}) = \sigma\left(\sum_i w_i x_i + b\right)$$

Convolutional Neural Networks


- A specialized kind of neural network for processing data with a known grid-like topology
- Examples:

- Time-series



1D grid

- Images




2D grid

NNs don't scale to images!


- Let us consider a fully connected network with a single unit
 - Tiny color images of size $32 \times 32 \times 3$
 - Size of the input layer: $32 \times 32 \times 3 = 3072$
 - Size of the weights: 3072
 - Small color images of size $200 \times 200 \times 3$
 - Size of input layer and weights: $200 \times 200 \times 3 = 120000$

From shallow to deep models

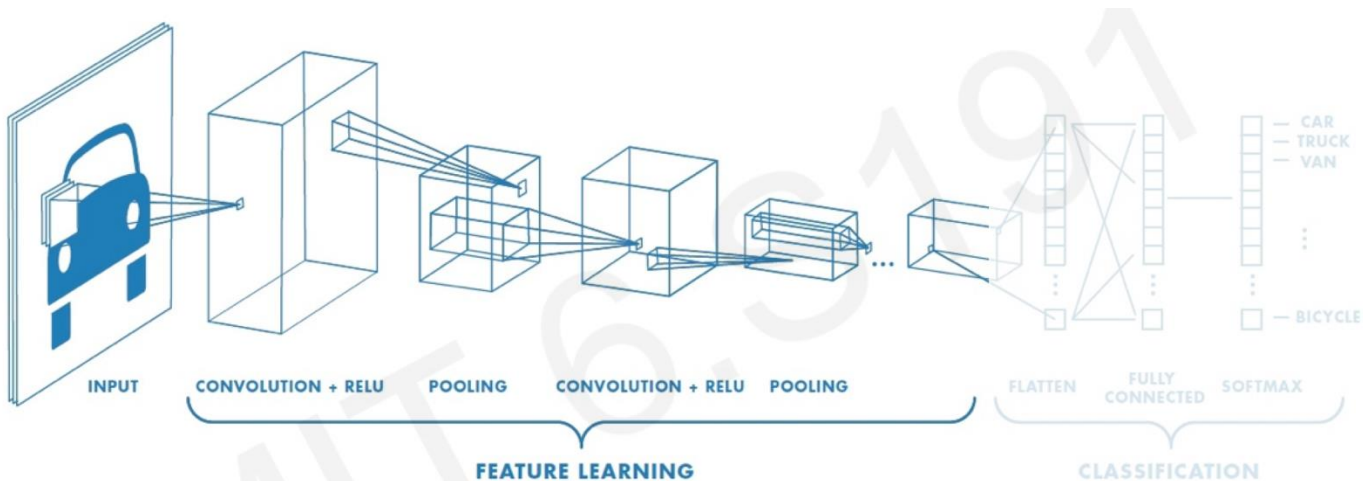
Shallow models



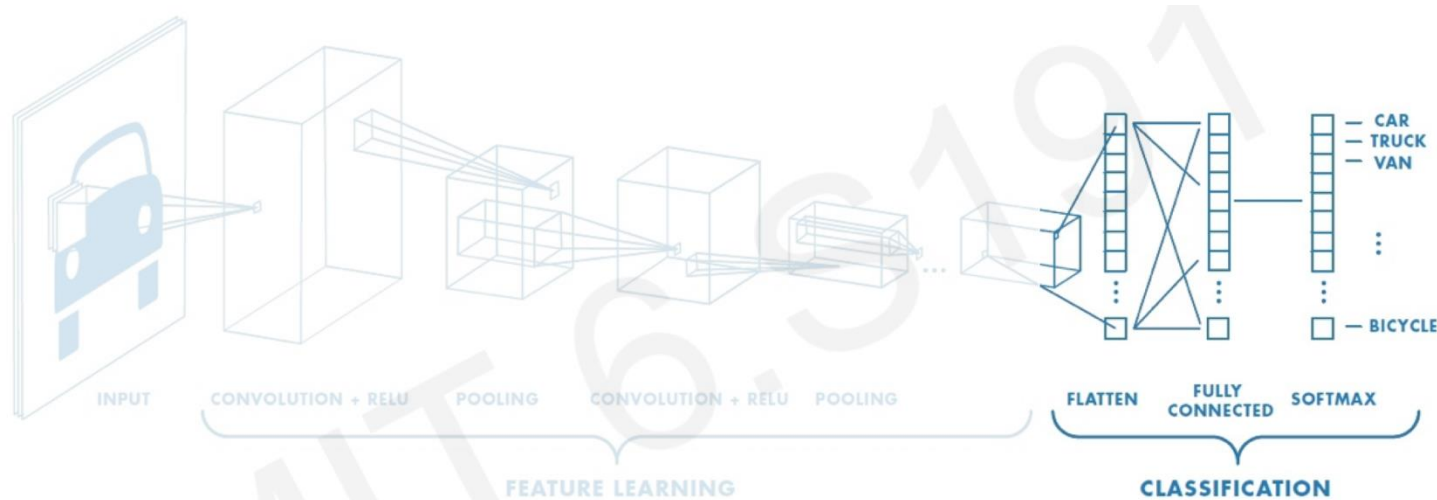
Deep models



A typical CNN



A typical CNN



$$\text{softmax}(y_i) = \frac{e^{y_i}}{\sum_j e^{y_j}}$$



Interlude: convolution

The convolution/cross-correlation operation

For 2D input arrays:

$$\begin{aligned} s(i, j) &= (K * I)(i, j) = \sum_m \sum_n I(m, n) K(i - m, j - n) = \\ &= \sum_m \sum_n I(i - m, j - n) K(m, n) \end{aligned}$$

$$s(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

Cross-correlation (with an example)

$$s(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$

| | | | |
|-----------------|-----------------|-----------------|-----------------|
| X ₁₁ | X ₁₂ | X ₁₃ | X ₁₄ |
| X ₂₁ | X ₂₂ | X ₂₃ | X ₂₄ |
| X ₃₁ | X ₃₂ | X ₃₃ | X ₃₄ |
| X ₄₁ | X ₄₂ | X ₄₃ | X ₄₄ |

*

| | |
|-----------------|-----------------|
| W ₁₁ | W ₁₂ |
| W ₂₁ | W ₂₂ |

=


| | | |
|-----------------|-----------------|-----------------|
| Y ₁₁ | Y ₁₂ | Y ₁₃ |
| Y ₂₁ | Y ₂₂ | Y ₂₃ |
| Y ₃₁ | Y ₃₂ | Y ₃₃ |

$$\begin{aligned} Y_{11} &= X_{11} W_{11} + X_{12} W_{12} + X_{21} W_{21} + X_{22} W_{22} \\ Y_{12} &= X_{12} W_{11} + X_{13} W_{12} + X_{22} W_{21} + X_{23} W_{22} \\ Y_{13} &= X_{13} W_{11} + X_{14} W_{12} + X_{23} W_{21} + X_{24} W_{22} \\ &\dots\dots \end{aligned}$$

2D convolution


A "feature detector" (kernel) slides over the inputs to generate a feature map

$$s(i, j) = (K * I)(i, j) = \sum_m \sum_n I(i + m, j + n) K(m, n)$$



Input tensor
of size 10x10

Kernel of
size 3x3




Output tensor of size
8x8




Convolutional layer

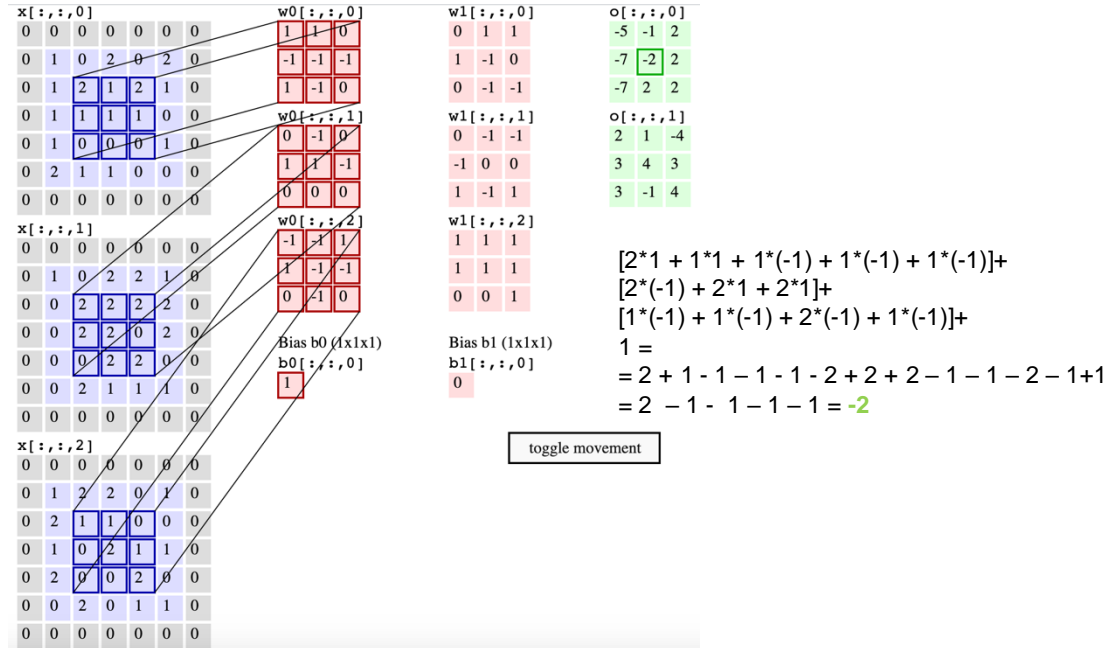
A sketch of a dense layer



A sketch of a convolutional layer



An example from <https://cs231n.github.io/convolutional-networks/>




Another animation to clearly understand

$W \times H \times 3$




$K \times K \times 3 \times 4$
($C=4$)



$W' \times H' \times 4$
($C=4$)



A sketch of a convolutional layer




Also called the **detector** stage, it provides a set of **linear activations**



As the kernel slides on the image, it is able to capture the **same property in different image regions** → THERE IS PARAMETER SHARING

Multiple feature detectors can be used to capture different image properties → Their number is called **channels**



From dense to sparse interaction



From dense to sparse interaction




From dense to sparse interaction



With the application of the kernel, each input elements interact with only a subset of other input elements → THERE IS SPARSE INTERACTION

Output feature size of conv layers

- Three parameters control the size of the output of a layer
 - **Channles**, the number of filters (kernels) of the layer
 - **Stride**, the step used to slide the filter on the input
 - When stride > 1 we are down-sampling the input data
 - **Tiling** refers to the special case where stride = kernel span
 - **Padding** to enlarge the input and allow for kernels application in each one of the (original) point



Output features size of conv layers

Size BEFORE convolution

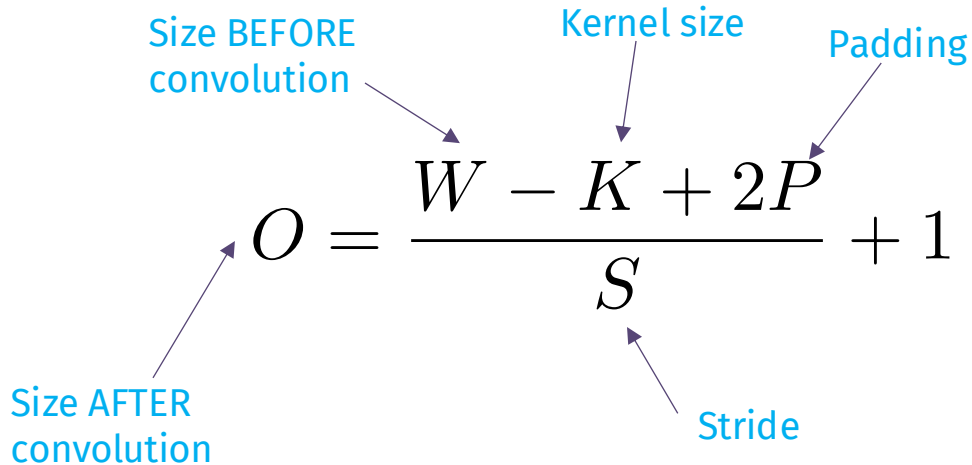
Kernel size

Padding

$$O = \frac{W - K + 2P}{S} + 1$$


Size AFTER convolution

Stride

The diagram shows the formula for the output feature size of a convolutional layer. The formula is $O = \frac{W - K + 2P}{S} + 1$. Annotations with arrows point to each variable: 'Size BEFORE convolution' points to W , 'Kernel size' points to K , 'Padding' points to P , 'Stride' points to S , and 'Size AFTER convolution' points to O .

Output features size of conv layers

Examples




No padding, stride 1

$$O = \frac{W - K + 2P}{S} + 1$$

Output features size of conv layers

Examples




No padding, stride 2

$$O = \frac{W - K + 2P}{S} + 1$$

Diagram illustrating the formula for output size O given input width W , kernel size K , padding P , and stride S . The values are: $W=5$, $K=3$, $P=0$, and $S=2$.

Output features size of conv layers

Examples



Padding 2, stride 1


$$O = \frac{W^5 - K^4 + 2P^2}{S_1} + 1$$

6



To sum up: output features size of conv layers

- Input size: $W_1 \times H_1 \times D_1$
- Parameters:
 - Number of kernels N
 - Kernel size K
 - Stride S
 - Padding P
- Output size: $W_2 \times H_2 \times D_2$
 - where
 - $W_2 = (W_1 - K + 2P)/S + 1$
 - $H_2 = (H_1 - K + 2P)/S + 1$
 - $D_2 = N$
 - Number of weights per filter: $K \times K \times D_1$
 - Number of parameters in total:
 - $K \times K \times D_1 \times N$ weights
 - N biases



A sketch of a convolutional layer



A typical choice: ReLU



A sketch of a convolutional layer



A way to further **reduce the dimensionality** of the representation while providing **invariance to small shifts** of the inputs

Common choices: **average** or **max pooling**

Pooling with an example

| | | | | | |
|---|---|---|---|---|---|
| 2 | 1 | 7 | 1 | 2 | 5 |
| 5 | 0 | 3 | 4 | 1 | 2 |
| 1 | 7 | 8 | 3 | 3 | 0 |
| 0 | 3 | 2 | 0 | 1 | 1 |
| 3 | 6 | 5 | 3 | 0 | 3 |
| 3 | 6 | 0 | 2 | 1 | 0 |

Max
pooling

| | |
|---|---|
| 8 | 5 |
| 6 | 3 |

Average
pooling

| | |
|-----|-----|
| 3.8 | 2.3 |
| 3 | 1.2 |

Pooling can help with local invariance although some information is lost

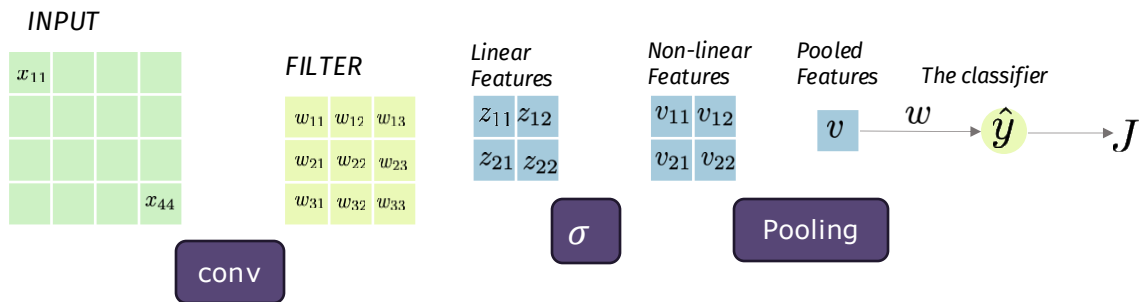
No parameter to be estimated here!

To sum up: output features size of pooling layer


- Input size: $W_1 \times H_1 \times D_1$
- Parameters:
 - Window size H
 - Stride S
- Output size: $W_2 \times H_2 \times D_2$
where
 - $W_2 = (W_1 - H)/S + 1$
 - $H_2 = (H_1 - H)/S + 1$
 - $D_2 = D_1$
- Number of weights per filter: $K \times K \times D_1$
- Number of parameters in total:
 - $K \times K \times D_1 \times N$ weights
 - N biases

Backpropagation in CNNs

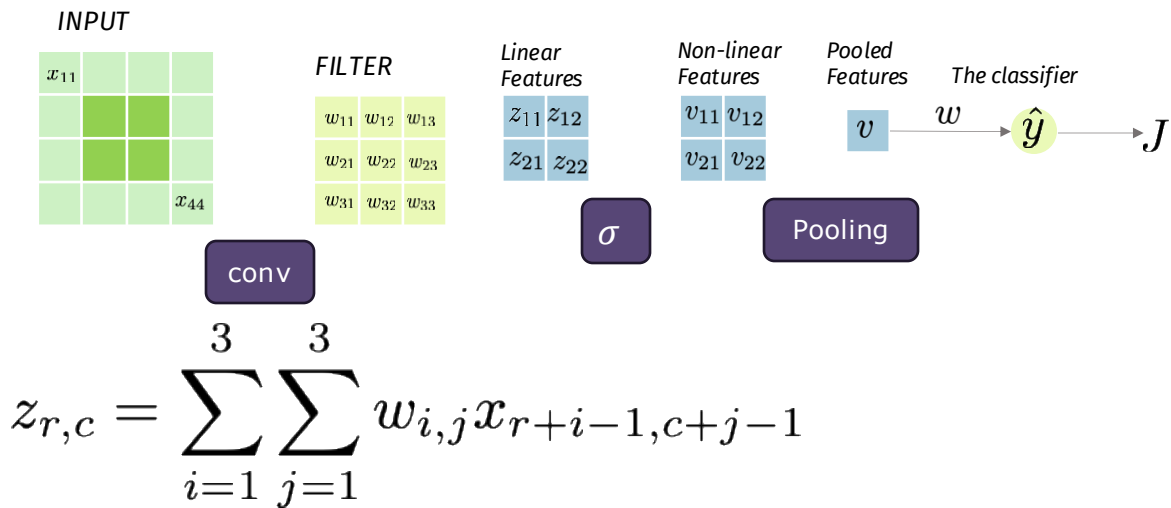
Backpropagation in CNNs (intuition)



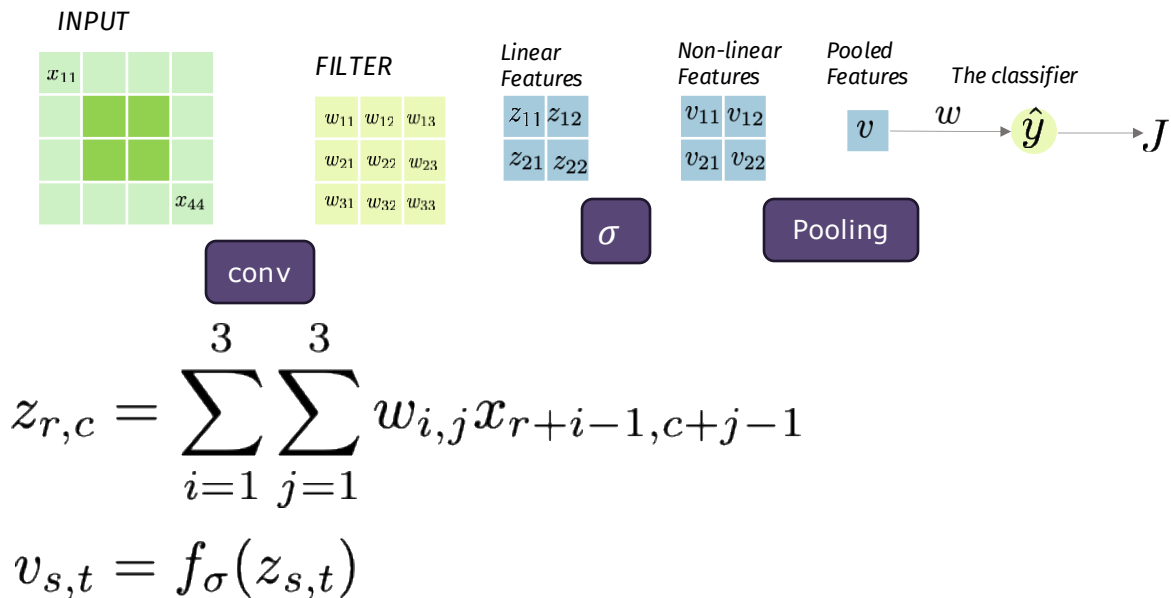
Backpropagation in CNNs (intuition)



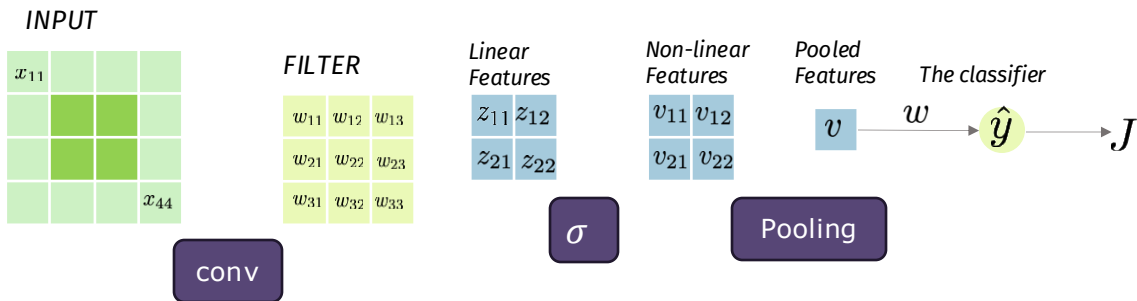
Backpropagation in CNNs (intuition)



Backpropagation in CNNs (intuition)



Backpropagation in CNNs (intuition)

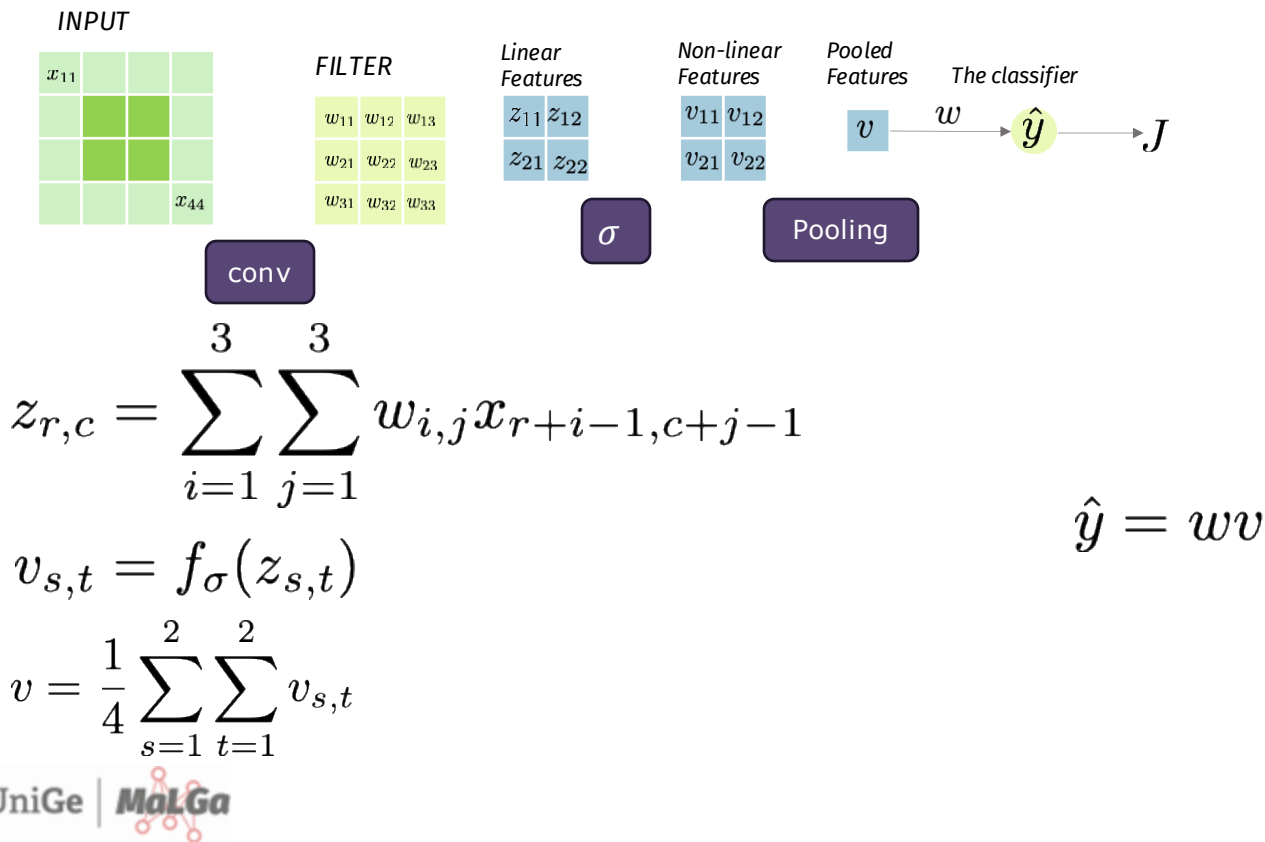


$$z_{r,c} = \sum_{i=1}^3 \sum_{j=1}^3 w_{i,j} x_{r+i-1, c+j-1}$$


$$v_{s,t} = f_{\sigma}(z_{s,t})$$

$$v = \frac{1}{4} \sum_{s=1}^2 \sum_{t=1}^2 v_{s,t}$$

Backpropagation in CNNs (intuition)



Backpropagation in CNNs (intuition)



$$z_{r,c} = \sum_{i=1}^3 \sum_{j=1}^3 w_{i,j} x_{r+i-1, c+j-1}$$

$$v_{s,t} = f_{\sigma}(z_{s,t})$$

$$v = \frac{1}{4} \sum_{s=1}^2 \sum_{t=1}^2 v_{s,t}$$

$$\hat{y} = wv$$

$$J(\mathbf{w}) = \frac{1}{N} \sum_{k=1}^N (\hat{y}_k - y_k)^2$$

Backpropagation

$$J(\mathbf{w}) = \frac{1}{N} \sum_{k=1}^N (\hat{y}_k - y_k)^2$$

$$J_k(\mathbf{w}) = (\hat{y}_k - y_k)^2$$

$$\nabla(J_k(\mathbf{w})) = \begin{bmatrix} \frac{\partial J_k(\mathbf{w})}{\partial w_{1,1}} \\ \dots \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{3,3}} \end{bmatrix}$$

Backpropagation

$$J(\mathbf{w}) = \frac{1}{N} \sum_{k=1}^N (\hat{y}_k - y_k)^2$$

$$J_k(\mathbf{w}) = (\hat{y}_k - y_k)^2$$

$$\nabla(J_k(\mathbf{w})) =$$

$$\begin{bmatrix} \frac{\partial J_k(\mathbf{w})}{\partial w} \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{1,1}} \\ \dots \\ \dots \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{3,3}} \end{bmatrix}$$

$$\frac{\partial J_k(\mathbf{w})}{\partial w} = \frac{\partial J_k(\mathbf{w})}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial v}$$

Backpropagation

$$J(\mathbf{w}) = \frac{1}{N} \sum_{k=1}^N (\hat{y}_k - y_k)^2$$

$$J_k(\mathbf{w}) = (\hat{y}_k - y_k)^2$$

$$\nabla(J_k(\mathbf{w})) =$$

$$\begin{bmatrix} \frac{\partial J_k(\mathbf{w})}{\partial w_{1,1}} \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{1,2}} \\ \dots \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{3,3}} \end{bmatrix}$$

$$\frac{\partial J_k(\mathbf{w})}{\partial w_{r,c}} = \frac{\partial J_k(\mathbf{w})}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial v} \left[\sum_{s=1}^2 \sum_{t=1}^2 \frac{\partial v}{\partial v_{s,t}} \frac{\partial v_{s,t}}{\partial f_{\sigma}(z_{s,t})} \frac{\partial f_{\sigma}(z_{s,t})}{\partial z_{s,t}} \frac{\partial z_{s,t}}{\partial w_{r,c}} \right]$$

Backpropagation

$$J(\mathbf{w}) = \frac{1}{N} \sum_{k=1}^N (\hat{y}_k - y_k)^2$$

$$J_k(\mathbf{w}) = (\hat{y}_k - y_k)^2$$

$$\nabla(J_k(\mathbf{w})) =$$

$$\begin{bmatrix} \frac{\partial J_k(\mathbf{w})}{\partial w_{1,1}} \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{1,1}} \\ \dots \\ \dots \\ \frac{\partial J_k(\mathbf{w})}{\partial w_{3,3}} \end{bmatrix}$$


$$\frac{\partial J_k(\mathbf{w})}{\partial w_{r,c}} = \frac{\partial J_k(\mathbf{w})}{\partial \hat{y}_k} \frac{\partial \hat{y}_k}{\partial v} \left[\sum_{s=1}^2 \sum_{t=1}^2 \frac{\partial v}{\partial v_{s,t}} \frac{\partial v_{s,t}}{\partial f_{\sigma}(z_{s,t})} \frac{\partial f_{\sigma}(z_{s,t})}{\partial z_{s,t}} \frac{\partial z_{s,t}}{\partial w_{r,c}} \right]$$



To further discuss...

A nice visualization

https://adamharley.com/nn_vis/cnn/3d.html



Intepretable models or interpretable data?



Gradient-weighted Class Activation Mapping (Grad-CAM), uses the gradients of any target concept flowing into the final convolutional layer to produce a coarse localization map highlighting the important regions in the image for predicting the concept

From <https://towardsdatascience.com/understand-your-algorithm-with-grad-cam-d3b62fce353>

Neuroscientific basis for convolutional networks




- Some of the design principles of Neural Networks have been drawn from neuroscience
- We now briefly discuss some of the connections between CNNs and a simplified version of the brain functions
- We refer to the primary visual cortex (V1 area), the first one in the brain performing some significantly advanced processing of visual input

V1 area & CNNs

- V1 is arranged in a spatial map
- V1 contains simple cells , that an to some extent be characterized by a linear function (as for the detection step in CNNs)
- V1 also contains complex cells, that show some level of invariance to some changes in the visual input
- It is generally believed that the same basic principles apply to other areas in the visual stream, repeatedly


Again on V1 cells

- Experiments showed that most V1 cells have weights that can be described by Gabor functions




What about CNN weights?

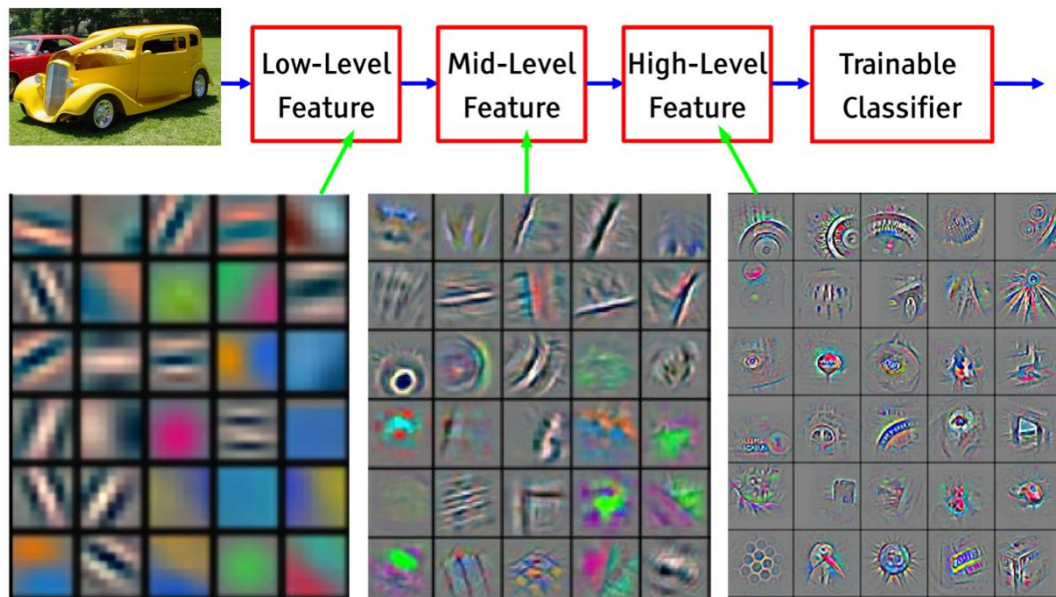
- At the very first layers the weights learnt by a CNN on natural images are very similar to Gabor filters



On the properties of weights learnt by convolutional layers




A hierarchical representation



Feature visualization of convolutional net trained on ImageNet, from [Zeiler & Fergus 2013]

Transfer learning



- It refers to the possibility of exploiting knowledge in terms of pre-trained models that can be used on different data and tasks (with some constraint)
- Fine-tuning is a well-assessed procedure in which the weights are somehow adapted to the new problem/data starting from the pre-trained model
- This may imply a domain shift (also known as covariate shift) problem, due to the fact that the data distribution may change as you change the problem/data

CNN training

- Very data hungry and computationally intensive
- One of the trick for coping with data lack us data augmentation
 - The idea is to generate more data by applying some transformation to the image

Data augmentation



https://m2dsupsdclass.github.io/lectures-labs/slides/04_conv_nets/index.html#82

CNN training

- Very data hungry and computationally intensive
- One of the trick for coping with data lack us data augmentation
 - The idea is to generate more data by applying some transformation to the image
- An alternative is to use synthetic data, but the model may be affected by domain shift issue (and thus it would need a specific domain adaptation strategy)

UniGe

