

15 - Computational geodesics

In this lecture

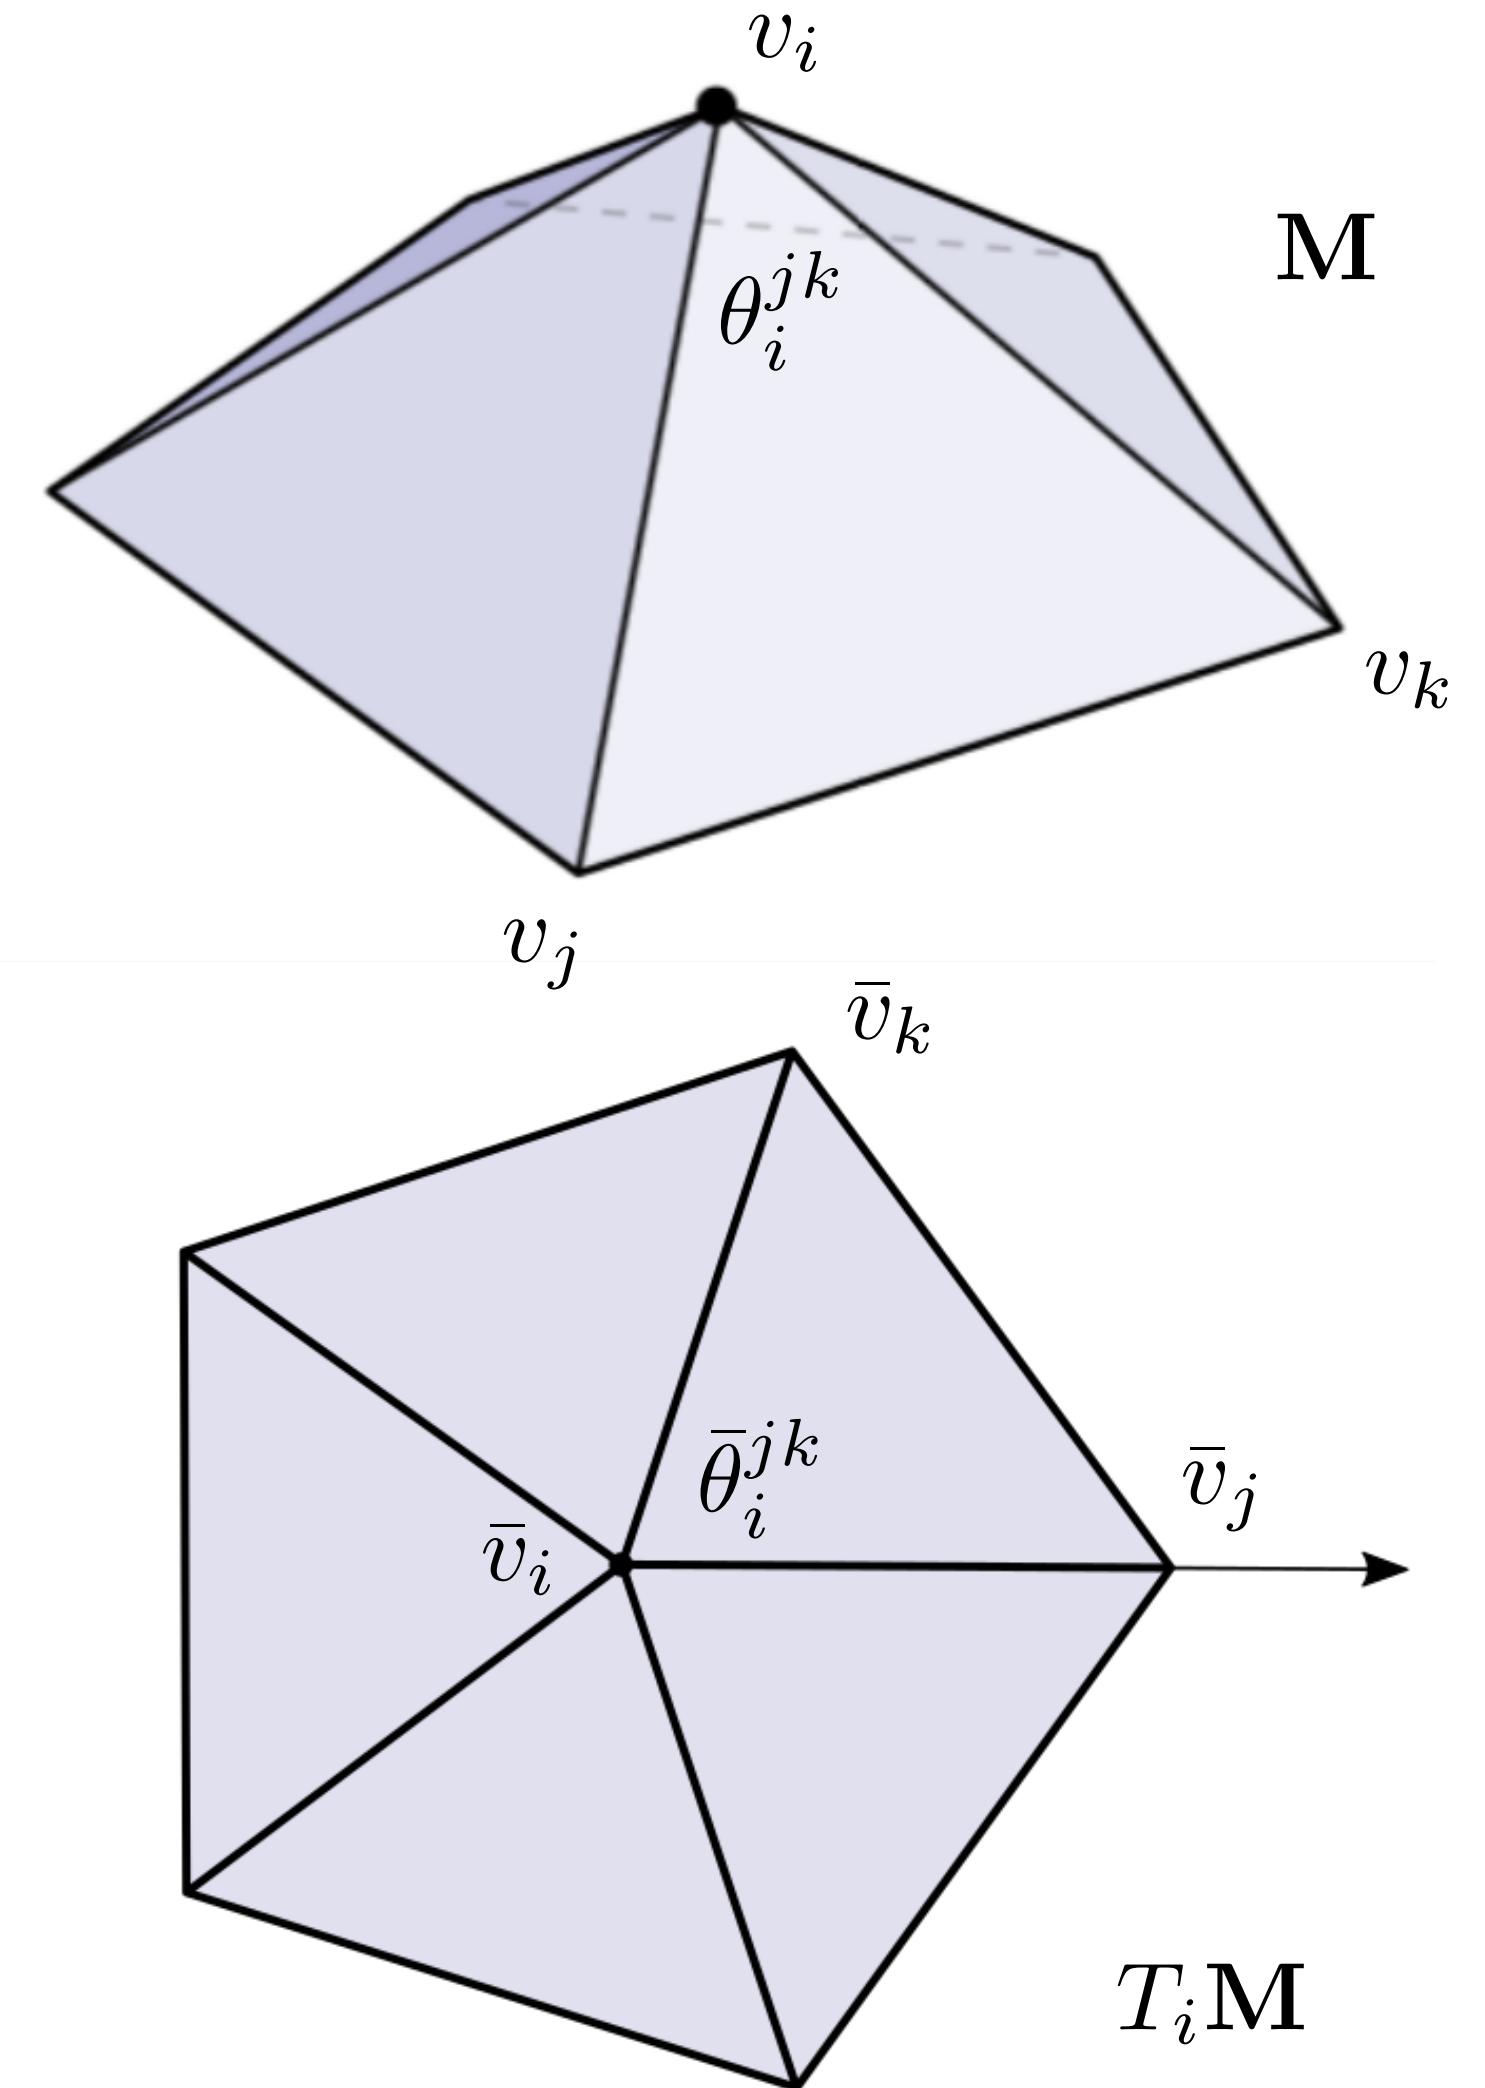
- Algorithms for computational geodesy:
 - on meshes
 - how to trace geodesics
 - how to find shortest paths
 - how to estimate the distance field

Geodesic Tracing

- We assume that the starting point is a vertex v_i (points on edges or inside triangles are easier to handle)
- Tangent direction t is given on the tangent plane $T_i\mathbf{M}$ at v_i
- Basic steps:
 1. project t on \mathbf{M} and find the triangle f crossed by such projection
 2. propagate to the triangle adjacent to f by the crossed edge
 3. repeat step 2 until reaching distance $|t|$

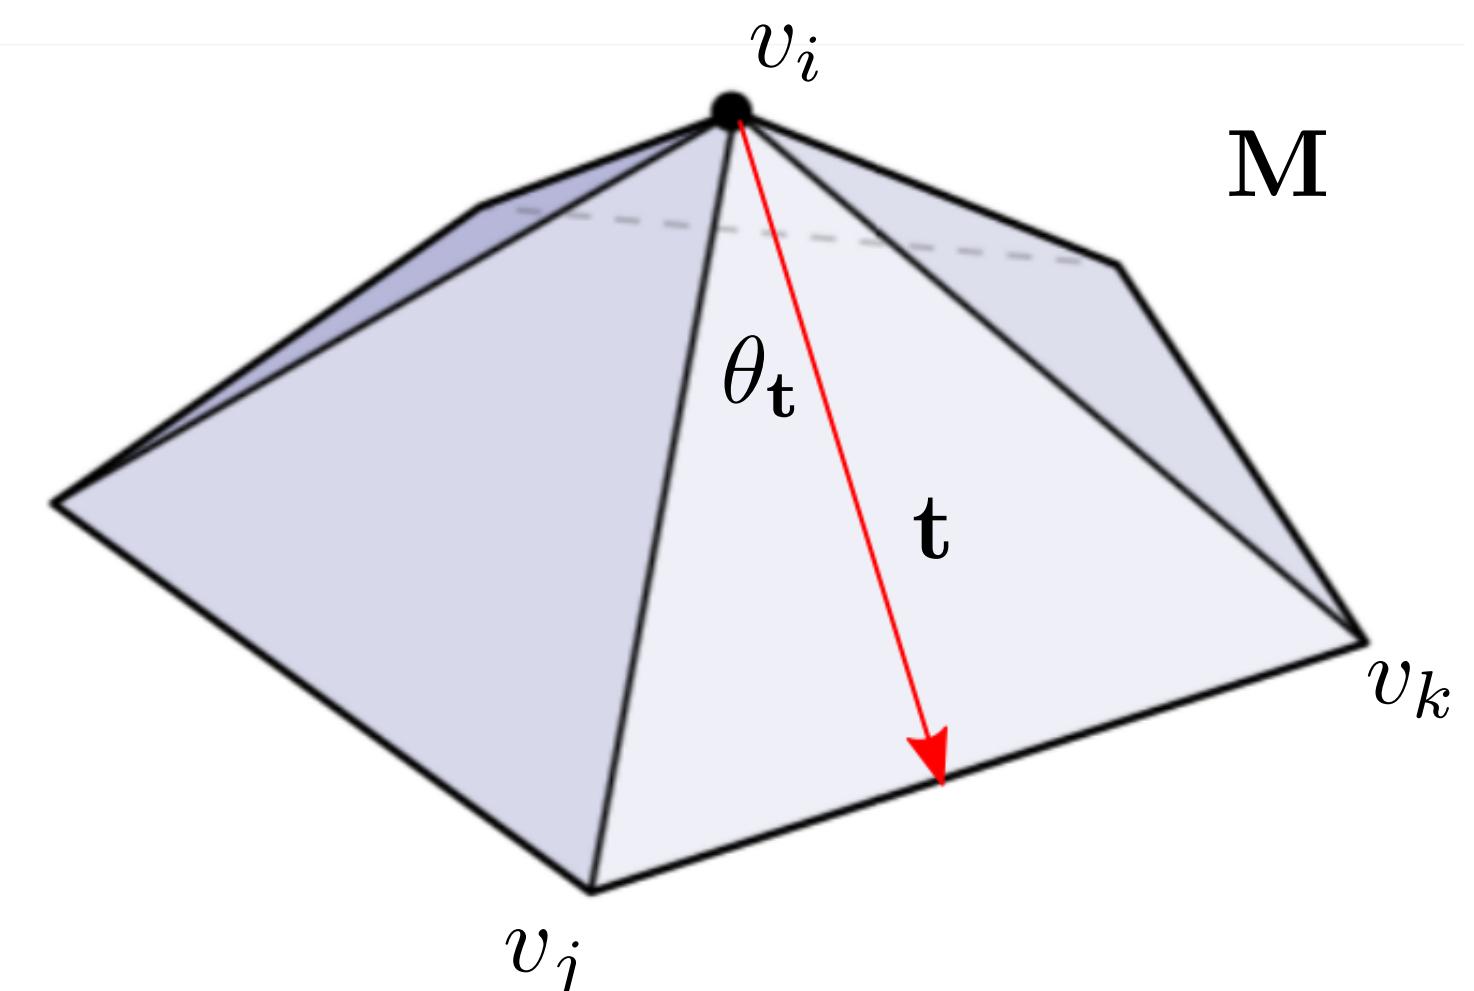
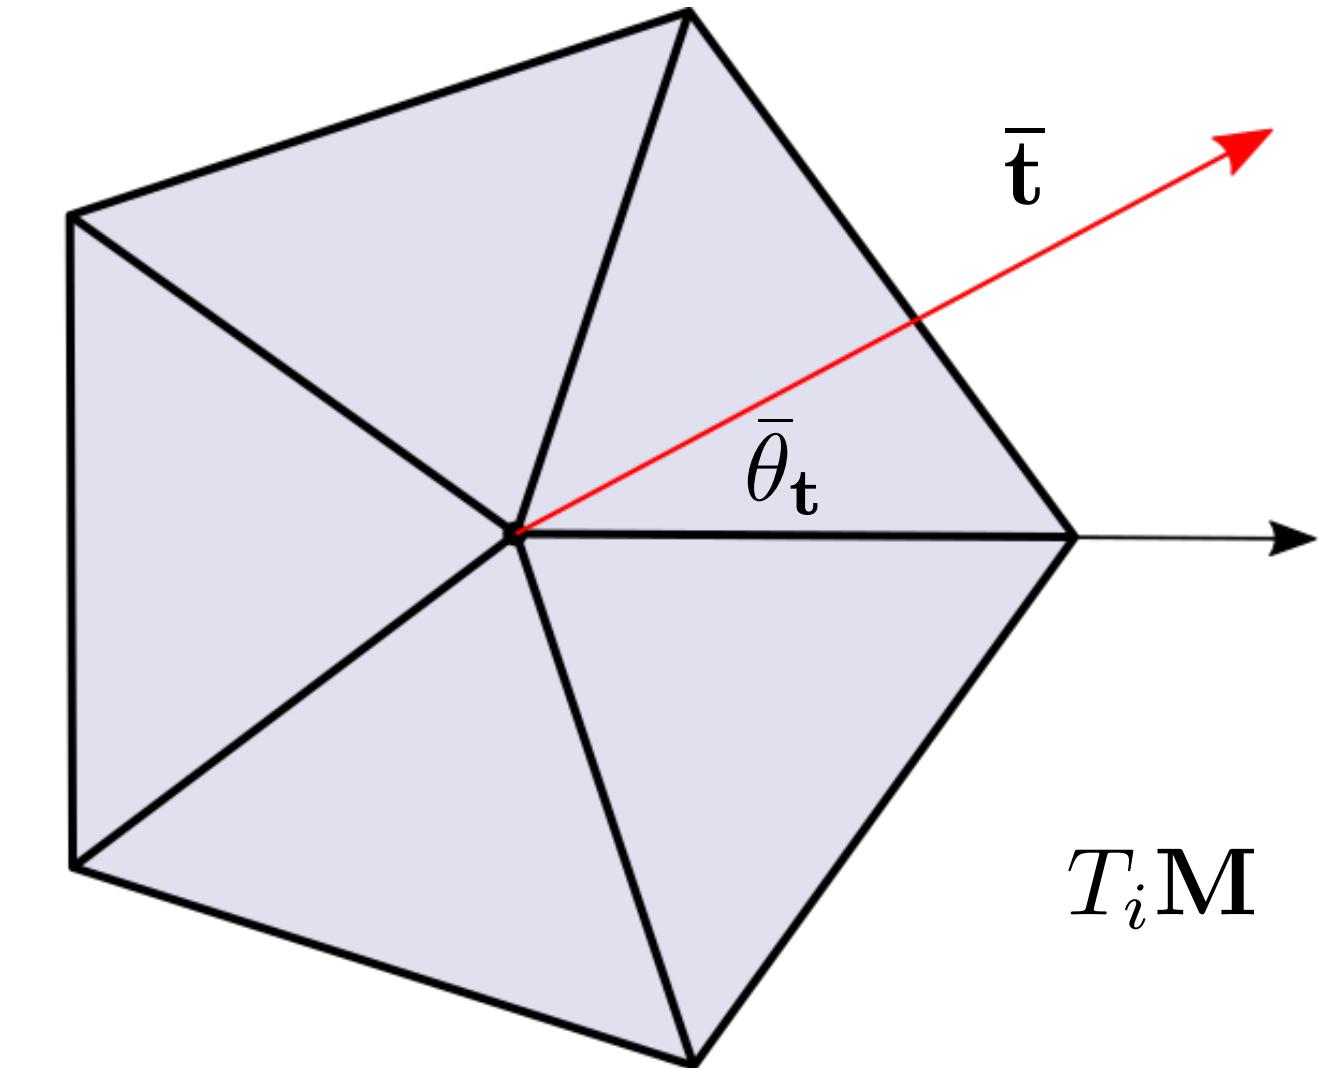
Geodesic Tracing

- Mapping the 1-ring to the tangent plane:
 - Compute the total angle Θ_i about vertex v_i
 - Rescale all incident angles by $2\pi/\Theta_i$
 - Map one edge to the x axis and distribute the others by rescaled angles



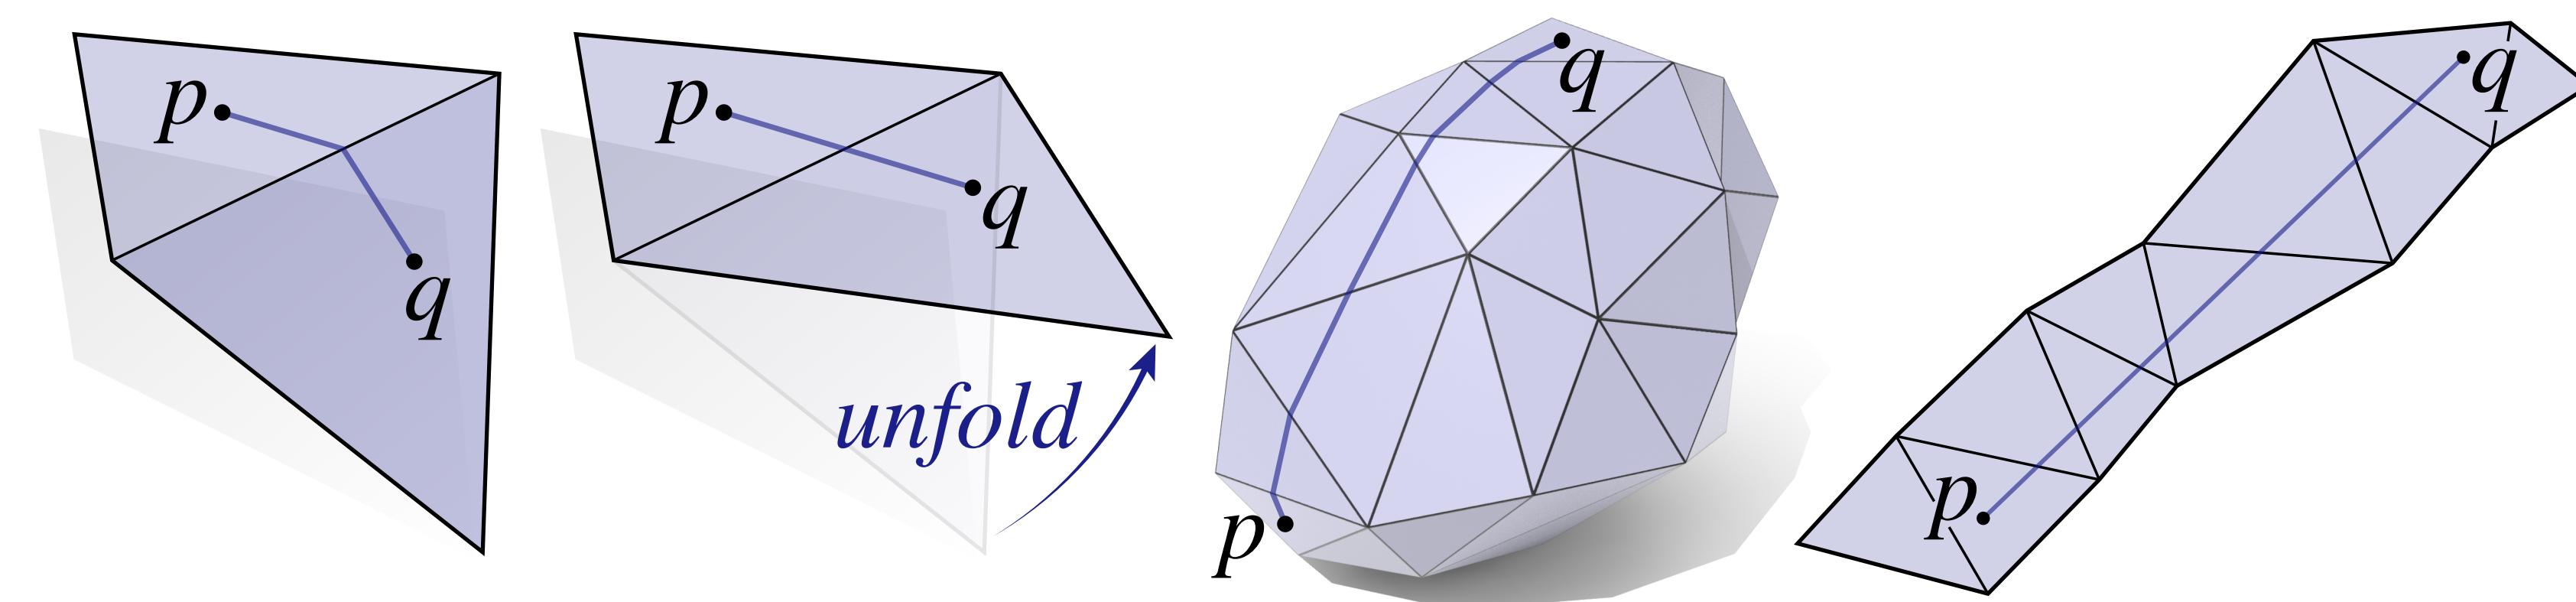
Geodesic Tracing

- Mapping the tangent direction to the mesh:
 - Find the triangle in tangent plane containing direction \bar{t}
 - Rescale angle by $\Theta_i/2\pi$
 - Map angle to the corresponding triangle of M



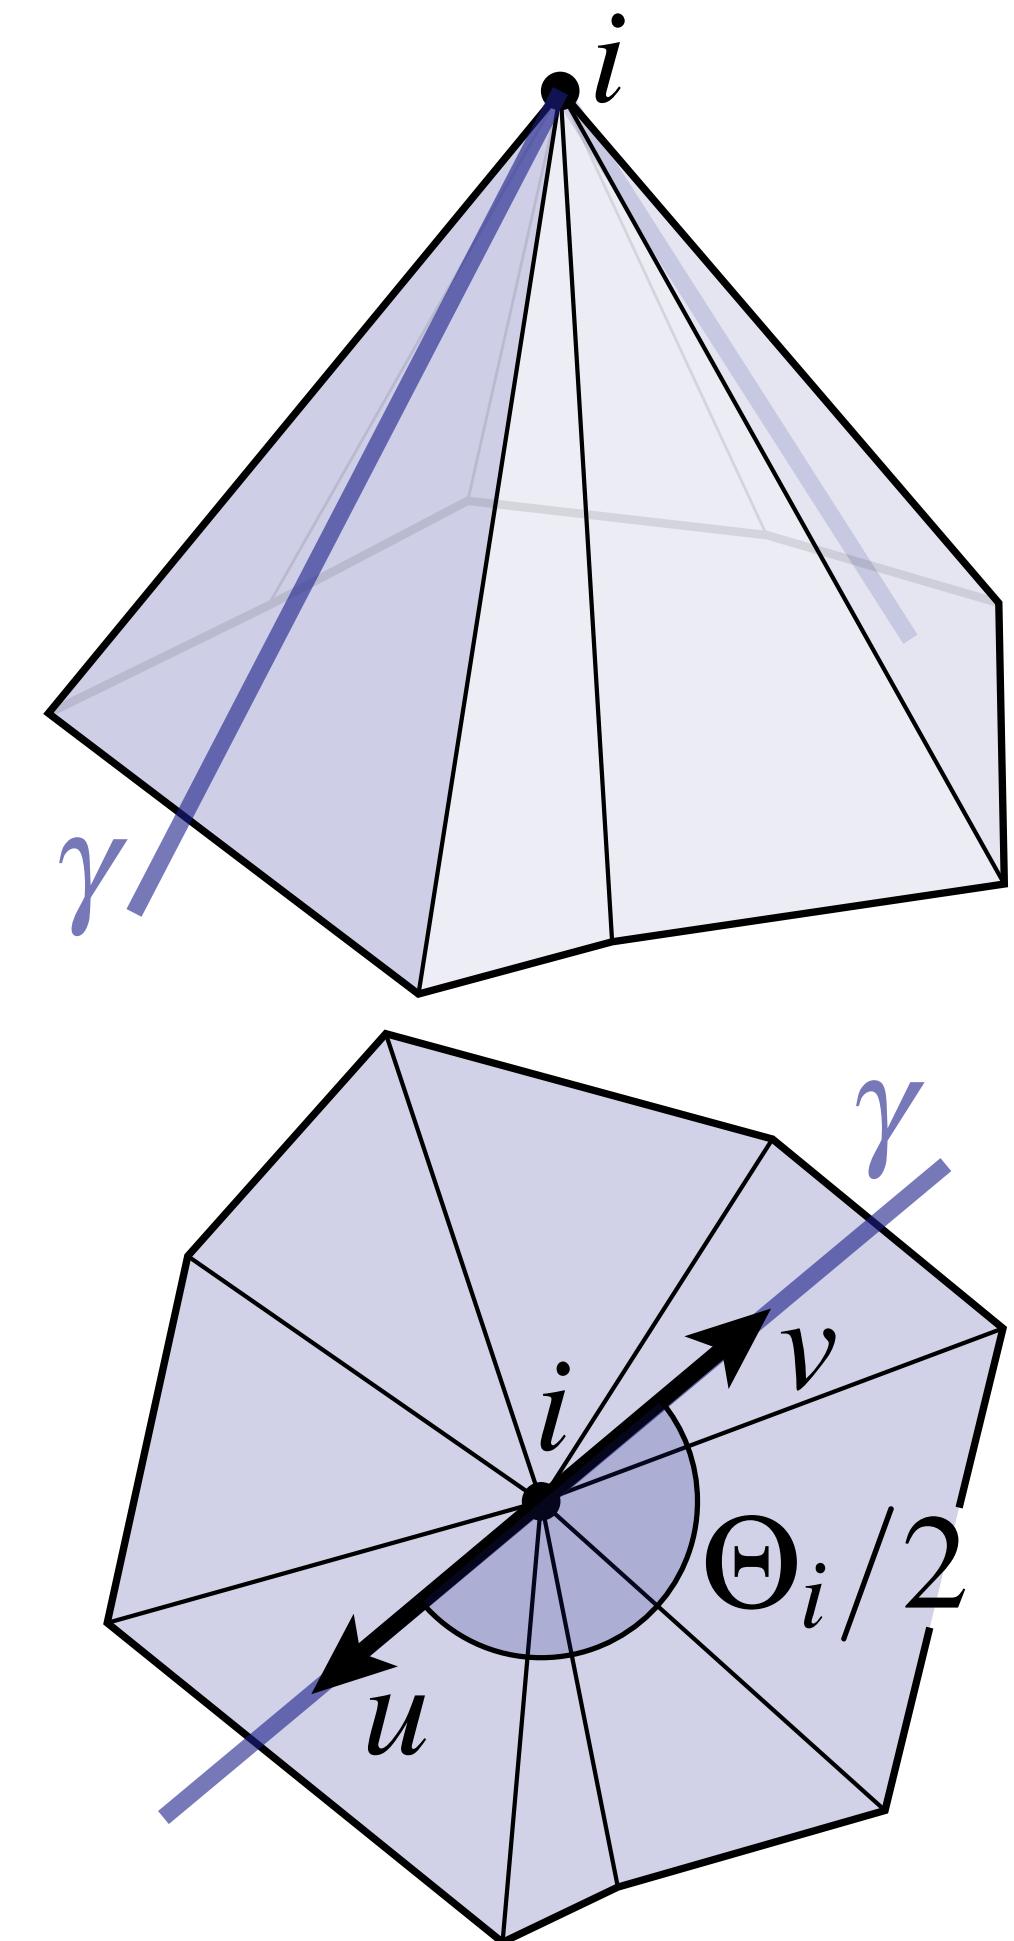
Geodesic Tracing

- Propagating the direction:
 - Find the outgoing edge of the direction in the current triangle
 - Unfold the neighboring triangle on the same plane and propagate a straight line through it
 - Repeated unfolding gives a triangle strip



Geodesic Tracing

- What to do if the geodesic crosses a vertex?
 - *straightest geodesic*: extend with line that halves the total angle [Polthier & Schmies 1998]
 - straightest geodesics do not distribute uniformly at saddle vertices
 - better uniform Geodesic Tracing exploits the fact that a geodesic is a normal curve [Biermann et al. 2002]



Single Source Geodesic Distance

- Three classes of algorithms:
 - Graph-based: restrict possible routes to paths in a graph; compute distance function on the graph
 - PDE-based: resolve a global PDE whose result approximates the distance function
 - Polyhedral waveform propagation: propagate a front through the mesh while building a data structure that supports computing the exact polyhedral distance from any point

Single Source Geodesic Distance

- Graph-based:
 - Pro: fast, scalable; easy pre-computation; trade-off between accuracy and complexity
 - Con: approximated; fields just at vertices (nodes of graph)
- PDE-based:
 - Pro: fast after pre-computation; approximates distance on the smooth surface
 - Con: approximated; field just at vertices; pre-computation does not scale well
- Polyhedral waveform propagation:
 - Pro: exact on polyhedra; field at all points; supports SSGP; no pre-computation
 - Con: slow

Graph-based techniques

Shortest path on graphs

- $G=(V,E)$ graph; V set of nodes, E set of arcs
- $w(e)$ weight associated to any edge (for us: geodesic length)
- Weight of a path $p = \langle v_0, v_1, \dots, v_k \rangle$ is the sum of weights on the arcs

$$w(p) = \sum_{i=1}^k w(v_{i-1}, v_i)$$

- Shortest path between u and v :

$$\bar{p}_{uv} = \operatorname{argmin}_p \{w(p) : p \text{ path connecting } u \text{ to } v\}$$

- Length of shortest path: $\delta(u, v) = w(\bar{p}_{uv})$
- Warning: the shortest path is not necessarily unique!

Single source shortest paths

- Assume all weights are non-negative
- Let $s \in V$, it is well defined the *tree of shortest paths* rooted at s

$$G_s = (V_s, E_s)$$

- where
 - V_s is the set of vertices that can be reached from s
 - for each $v \in V_s$ the (only) simple path from s to v in G_s is a shortest path in G

Optimal substructure of shortest paths

- For each shortest path $p = \langle v_0, v_1, \dots, v_k \rangle$ and for each pair v_i, v_j with $i < j$ the sub-path $p_{ij} = \langle v_i, v_{i+1}, \dots, v_j \rangle$ is a shortest path between v_i and v_j
- Let s be a source and u, v be two nodes connected with an arc in G , then we have

$$\delta(s, v) \leq \delta(s, u) + w(u, v)$$

- We can find shortest distances by expanding known shortest paths to neighbor nodes
- Since one node can be reached from different neighbors, we select the shortest among possible paths - all candidate paths must be visited

Graph relaxation

- For each node v let us define:
 - $d[v]$ candidate length of shortest path, initialized at $+\infty$
 - $\pi(v)$ node preceding v in a candidate path, initialized empty
- Relaxation of an arc (u, v) :

Relax (u, v, w)

if $d[v] > d[u] + w(u, v)$ then

$d[v] = d[u] + w(u, v)$

$\pi(v) = (\pi(u), (u, v))$

Graph relaxation

- If we repeatedly relax all arcs of E :
 - The value of $d[v]$ is always greater or equal than $\delta(s, v)$
 - If $d[v]$ reaches the value $\delta(s, v)$ then it becomes stable
- If $$ is a shortest path between s and v and we have $d[u] = \delta(s, u)$ then after relaxing edge (u, v) we have $d[v] = \delta(s, v)$

Dijkstra expansion

- Input: graph $G=(V,E)$, weight w , source s
- Output: graph $G=(V,E)$ where for each v we have $d[v] = \delta(s,v)$ and $\pi(v)$ points to his predecessor on the shortest path from s
- Shortest paths can be extracted next in optimal time by backtracking from destination to source following $\pi(v)$

Dijkstra expansion

```
Dijkstra( $G, w, s$ )
forall  $v$  do  $d[v] = +\infty$ ;  $\pi(v) = \text{NIL}$ 
 $d[s] = 0$ 
//  $S$ : vertices with a known shortest path
 $S = \emptyset$ 
//  $Q$ : priority queue, using  $d$  as key
forall  $v$  in  $V$  do Insert( $Q, v$ )
while  $Q \neq \emptyset$  do
     $u = \text{Min}(Q)$ 
     $S = S \cup \{u\}$ 
    forall  $v$  adjacent to  $u$  do
        Relax( $u, v, w$ )
        if relaxed cost is smaller than previous cost then
            Insert( $Q, v$ )
```

Dijkstra expansion

- Correctness: by absurd, assuming there is a node u that is added to S while $d[u] > \delta(s, u)$... Exercise
- Complexity: $O((|V|+|E|) \log |V|)$... Exercise
 - for a planar graph, $O(|V| \log |V|)$
 - the $\log |V|$ factor is due to the priority queue

SLF-LLL heuristics

- The priority queue of Dijkstra warrants optimal worst-case time, but it is expensive to maintain in practice
- For planar graphs and graphs used in geodesic algorithms, an algorithm using a double ended queue achieves better practical performances, while being not optimal in the worst case:
 - Small Label First (SLF): a new node is added either to the front or to the back of the queue, depending on its relative cost with respect to the first node in the queue
 - Large Label Last (LLL): nodes from the front of the queue that have a cost higher than the average are moved to the back of the queue

SLF-LLL heuristics

```
SLF-LLL( $G, w, s$ )
forall  $v$  do  $d[v] = +\infty$ ;  $\pi(v) = \text{NIL}$ ;  $d[s] = 0$ ;
add $_$ front( $Q, s$ );
while  $Q \neq \emptyset$  do
    while  $d[\text{front}(Q)] > \text{average cost in } Q$  do
        move  $\text{front}(Q)$  to the back of  $Q$ 
     $u = \text{pop}_\text{front}(Q)$ ;
    forall  $v$  adjacent to  $u$  do
        Relax( $u, v, w$ );
        if relaxed cost is smaller than previous cost then
            if  $d[v] < \text{average cost in } Q$  then add $_$ front( $Q, v$ )
            else add $_$ back( $Q, v$ );
```

Which graph to use?

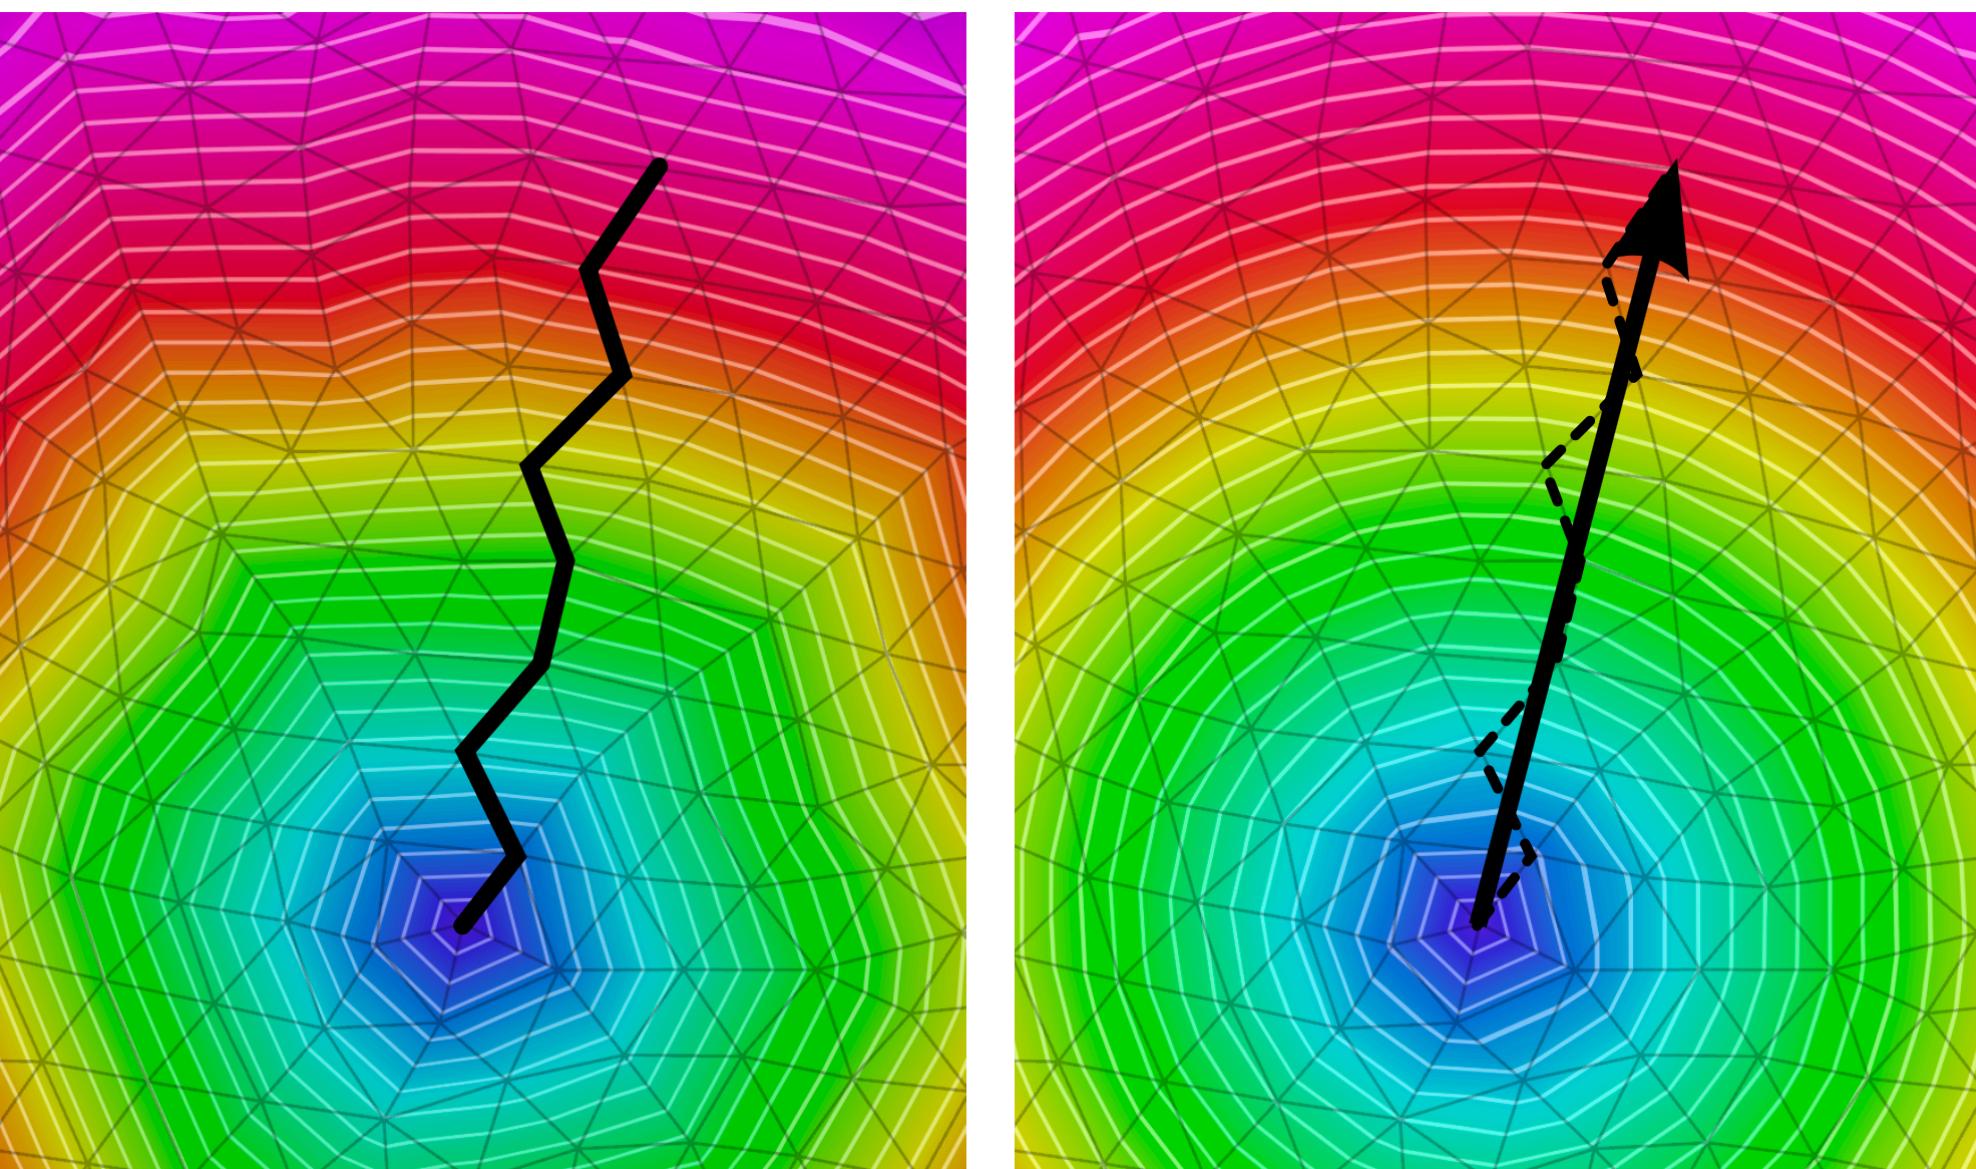
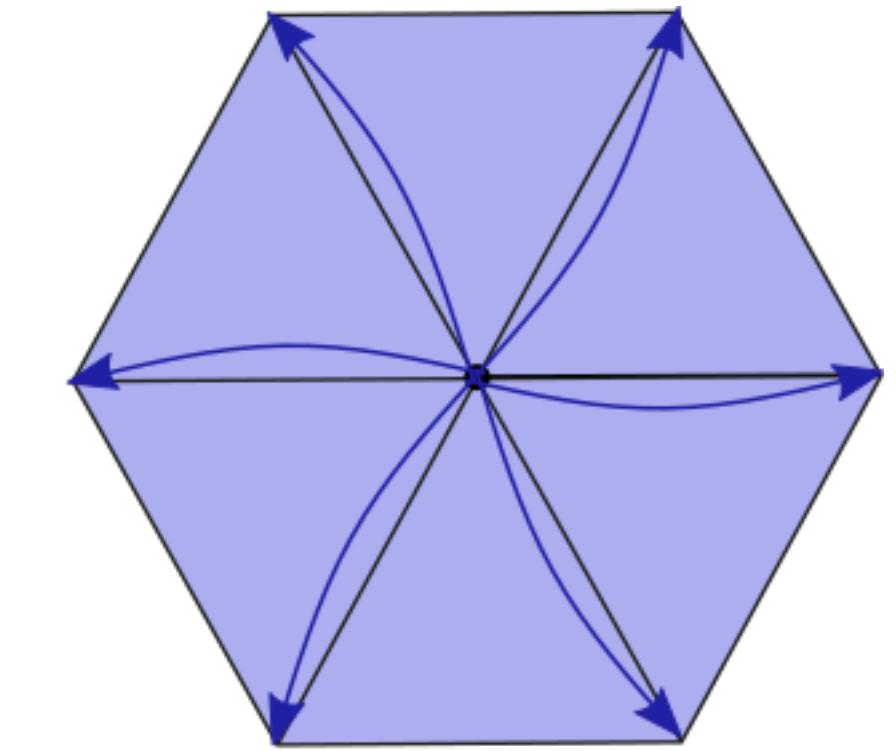
- The graph restricts the possible sources to its nodes, and the possible paths to graph paths
- A new source can be added by modifying the graph on-the-fly
- The distance field is defined only at the nodes of the graph
- The more nodes and arcs, the more accurate the result
- Trade-off: find a graph that achieves a good accuracy with a small number of nodes/arcs

Which graph to use?

- In most cases, defining the distance field just at the vertices of a mesh is sufficient
- The distance field is interpolated inside triangles
 - linear interpolation is not accurate!
- A graph having just the vertices at nodes is ok
- Edges are more important: wiggly paths are highly inaccurate

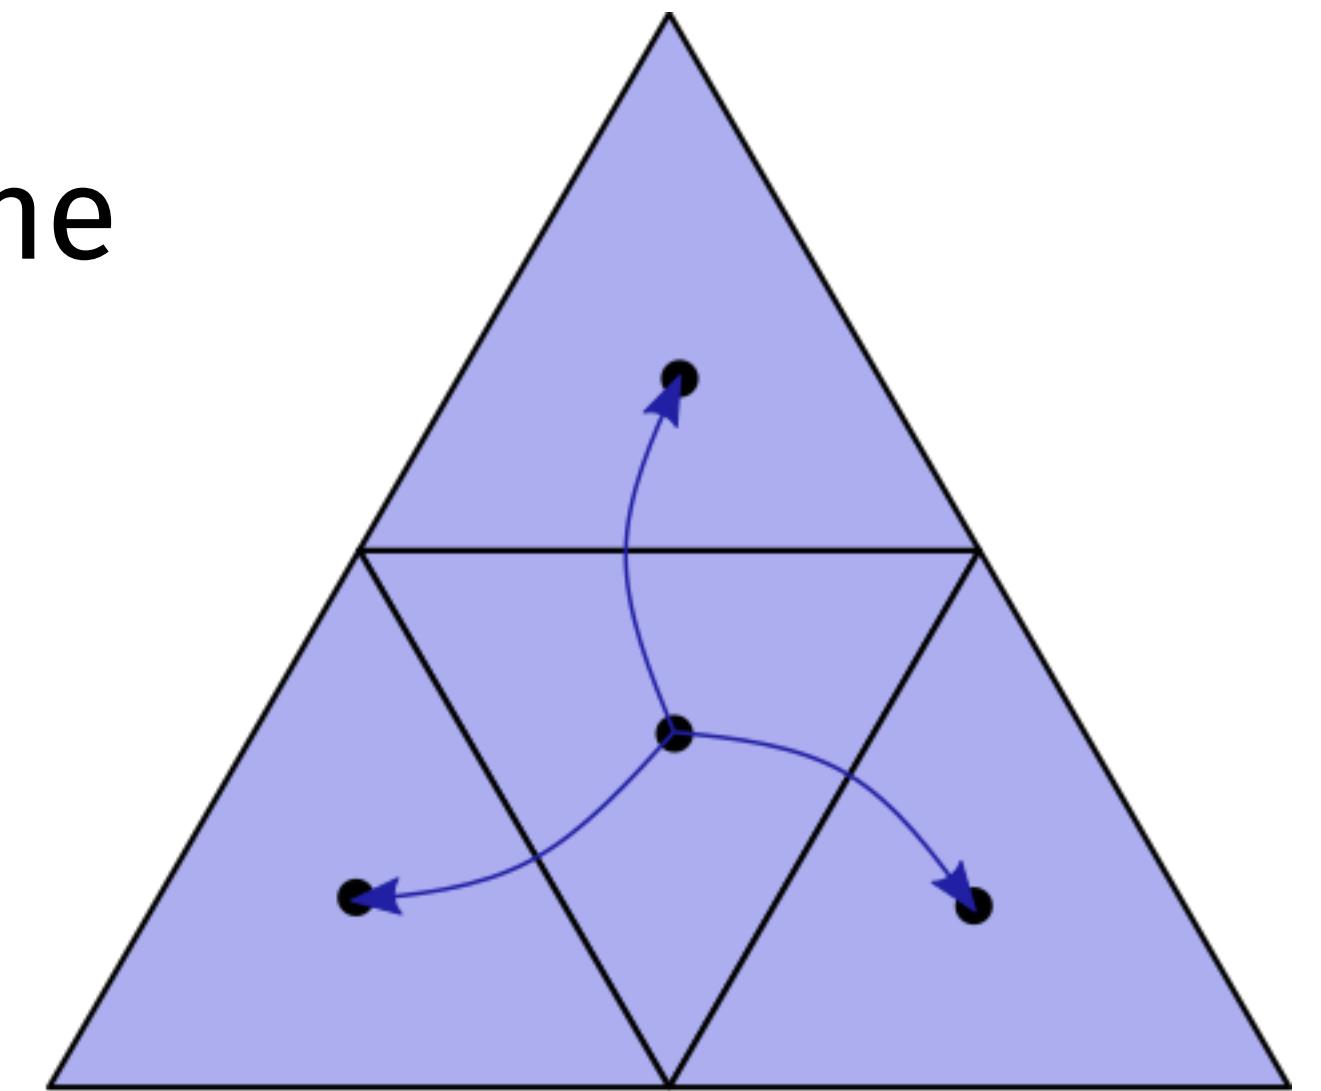
Which graph to use?

- Graph of edges (primal graph):
 - V vertices of the mesh
 - E edges of the mesh
 - Pro: easy and cheap, does not require any pre-processing
 - Con: wiggly paths, highly inaccurate



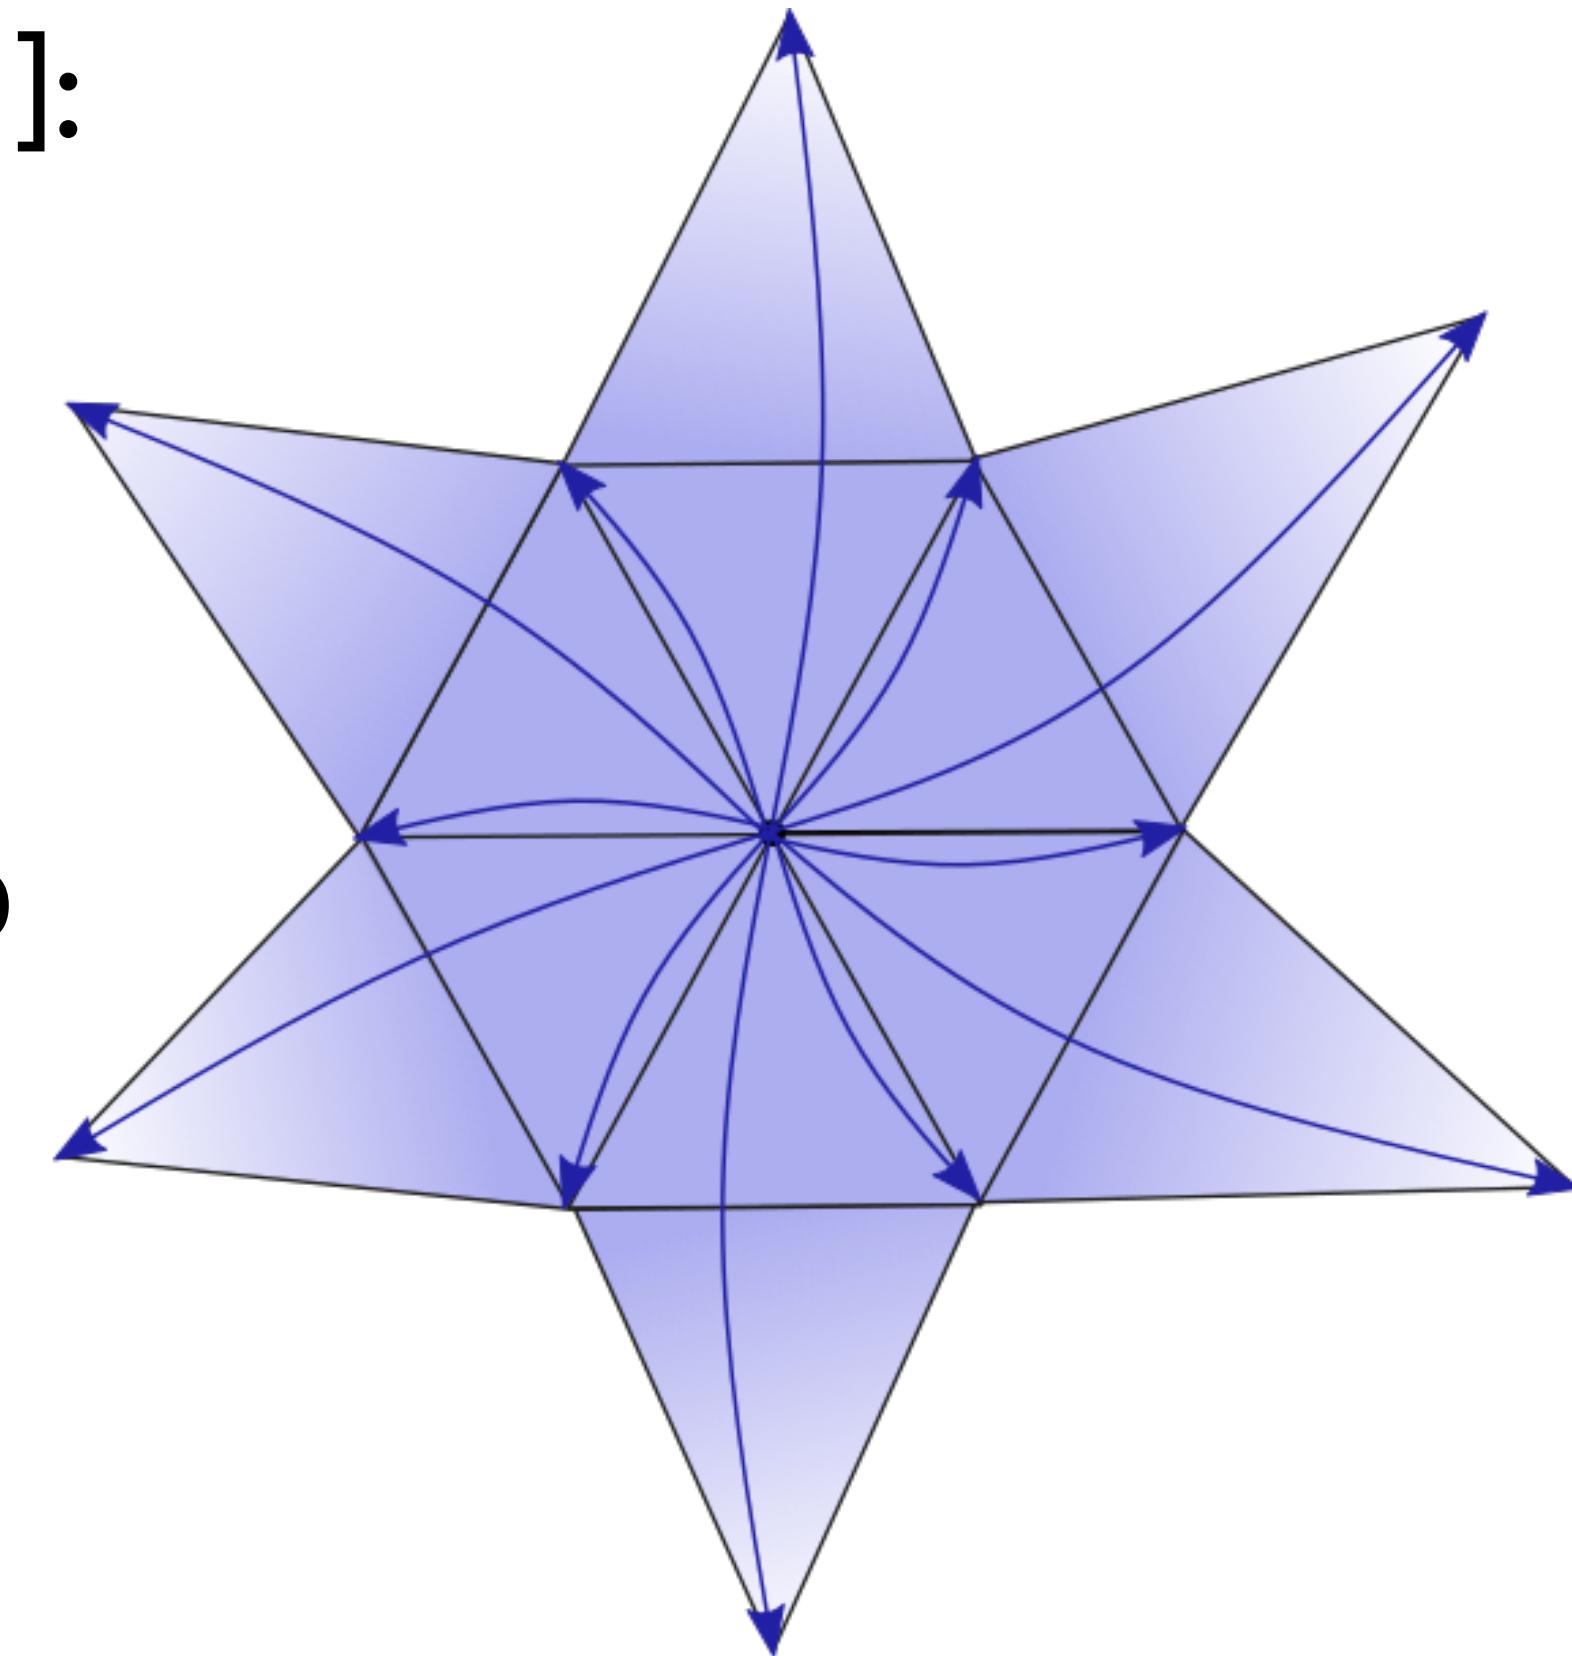
Which graph to use?

- Graph of adjacencies (dual graph):
 - V (centroids of) triangles of the mesh
 - E dual edges = adjacencies between triangles of the mesh
 - Pro: easy and cheap, does not require any pre-processing, valence three at all nodes
 - Con: wiggly paths, highly inaccurate
 - Often used to find initial guess to be refined further



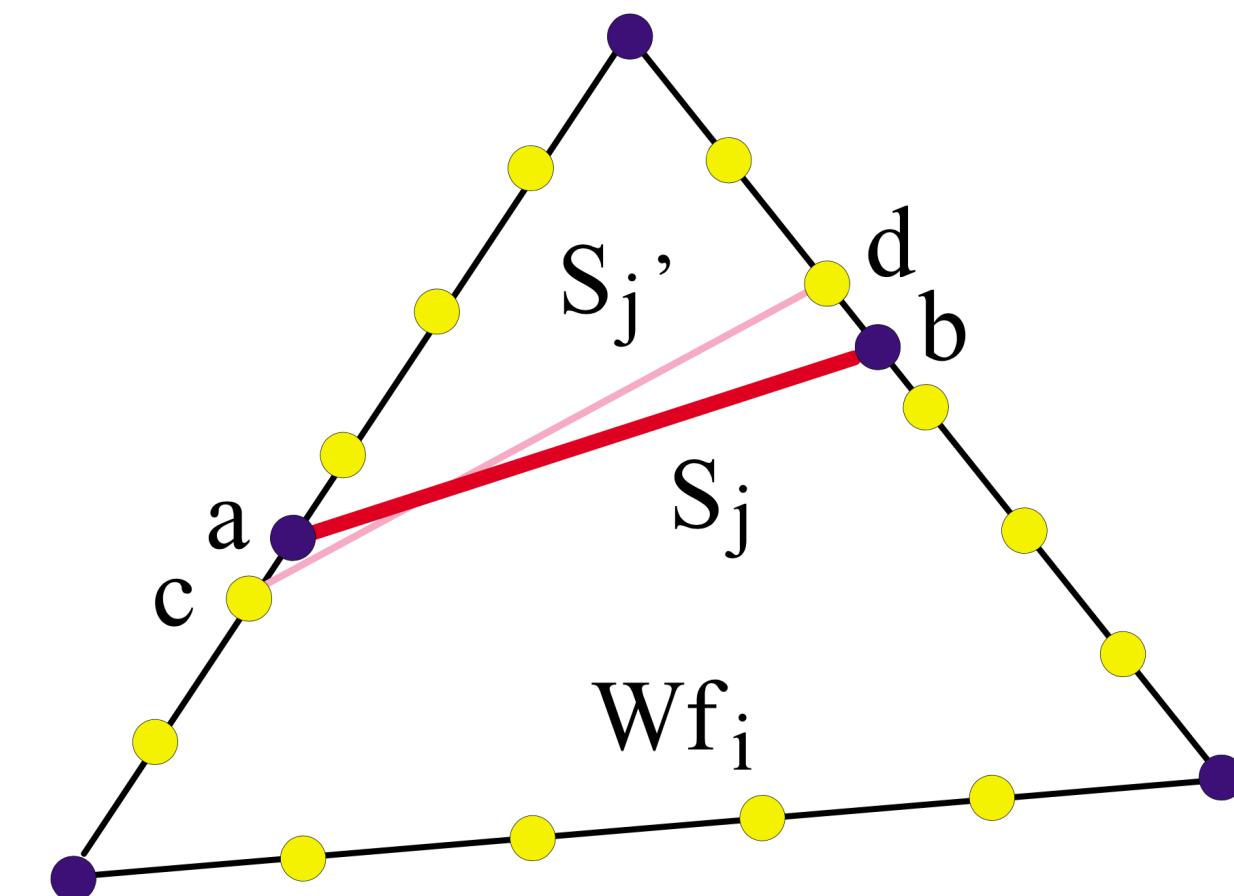
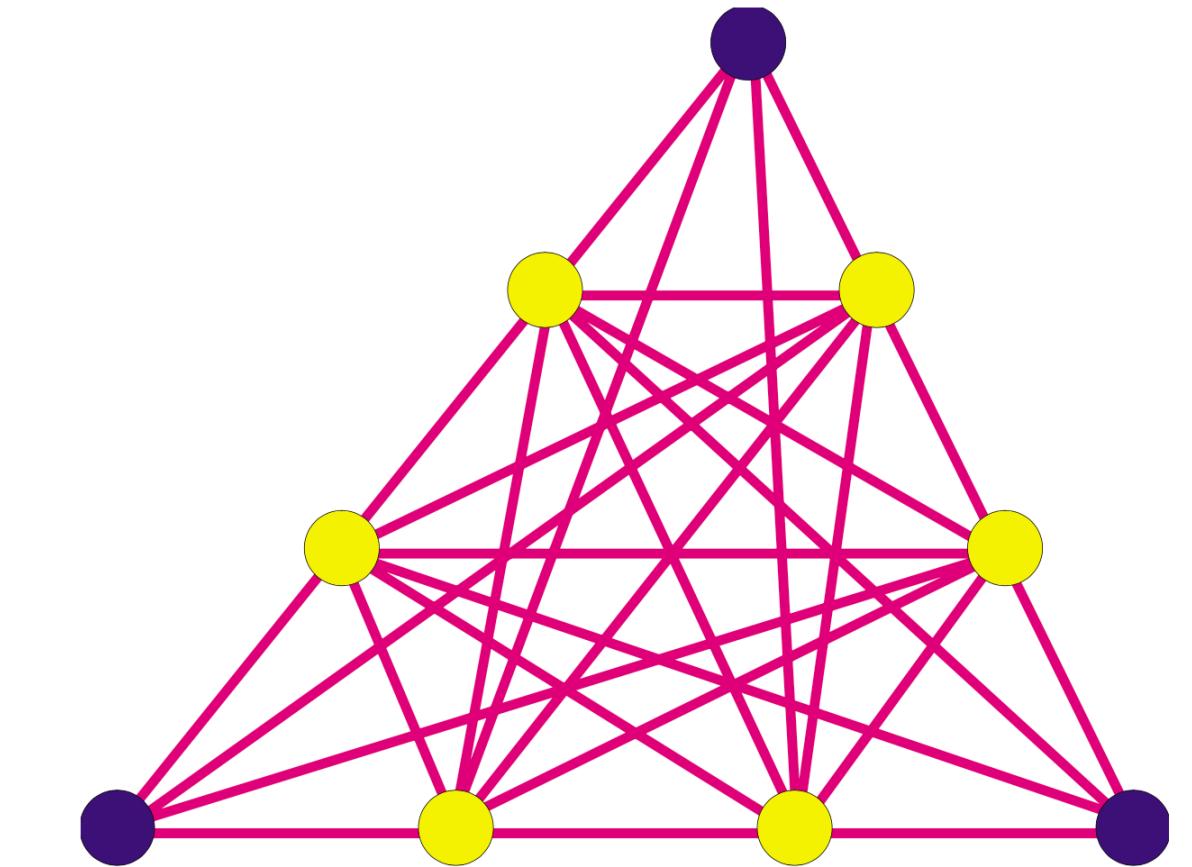
Which graph to use?

- Graph of primal and dual edges [Nazzaro et al., 2021]:
 - V vertices of the mesh
 - E edges of the mesh and dual edges: each vertex is connected with its neighbors and with the vertices opposite to it in the triangles adjacent to its 1-ring
 - easy and cheap, fast pre-processing
 - much better paths, moderately accurate (error 1-2%)



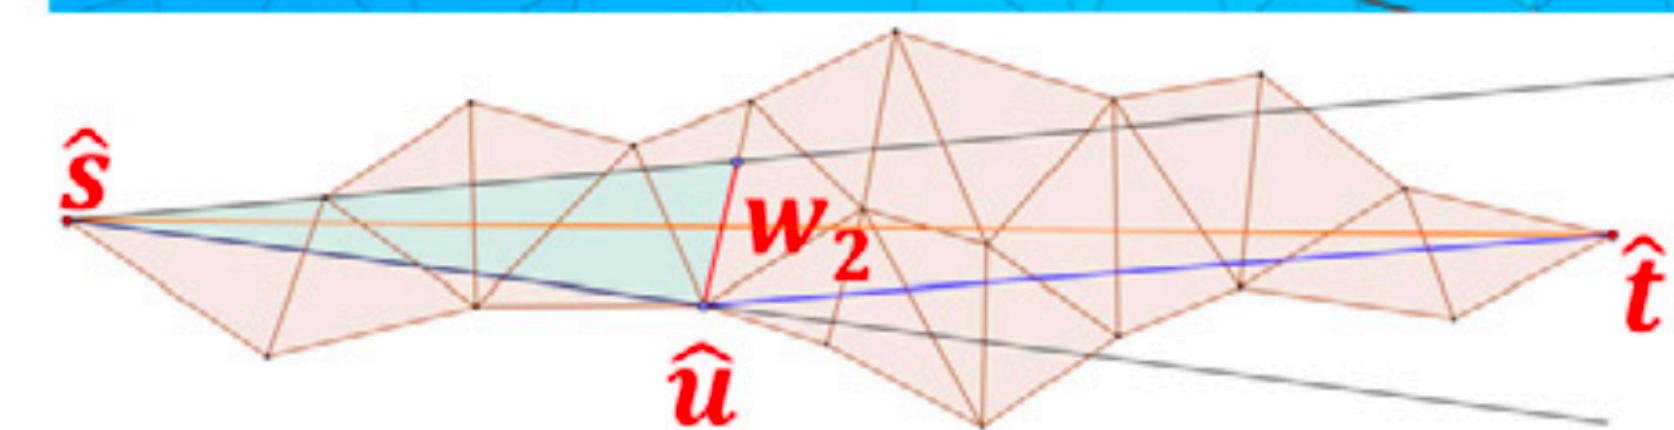
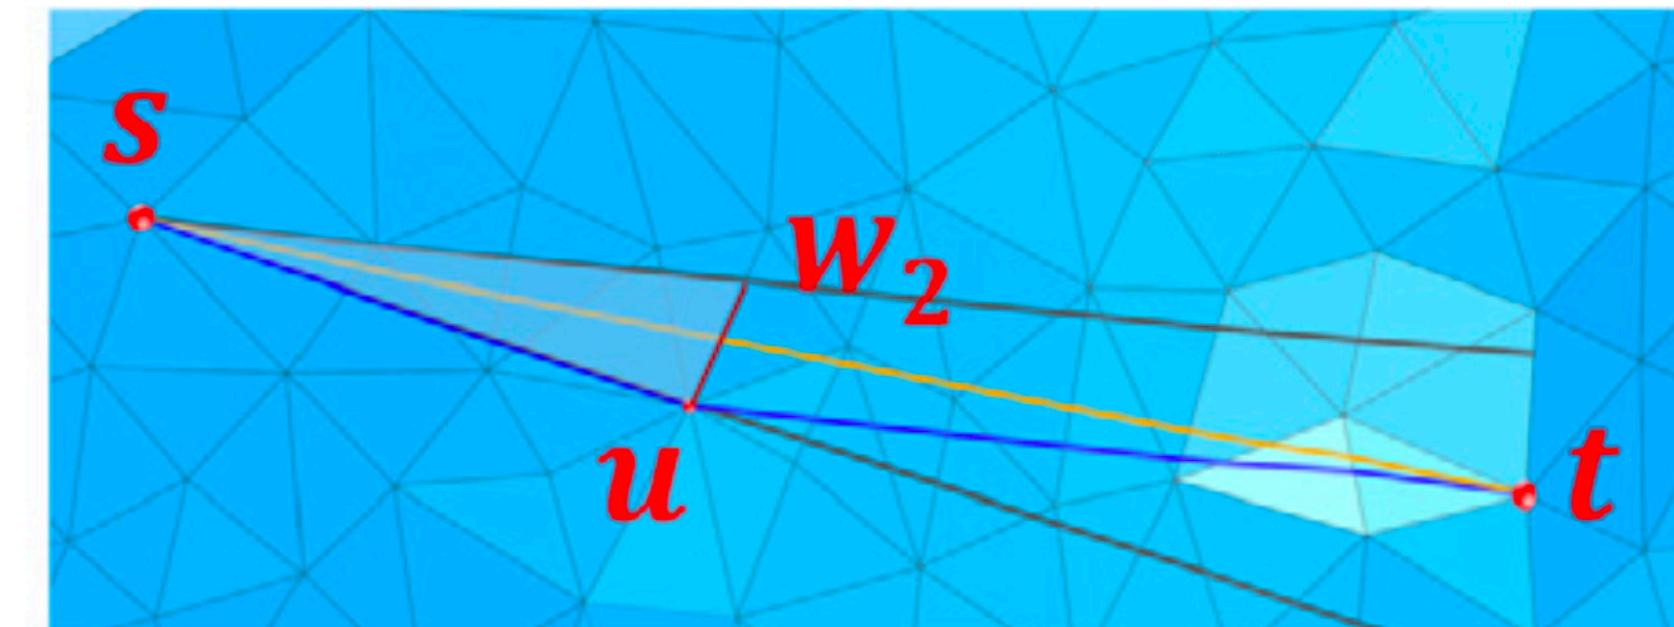
Which graph to use?

- Graph with Steiner points on edges [Lanthier 1999]:
 - edges of the mesh are split by distributing Steiner points along them
 - within each triangle, arcs joining all mutually visible points on its boundary are added
 - Pro: relatively easy to build, improves accuracy
 - Con: adds extra nodes; complexity explodes with the number of Steiner points per edge; query time slows down

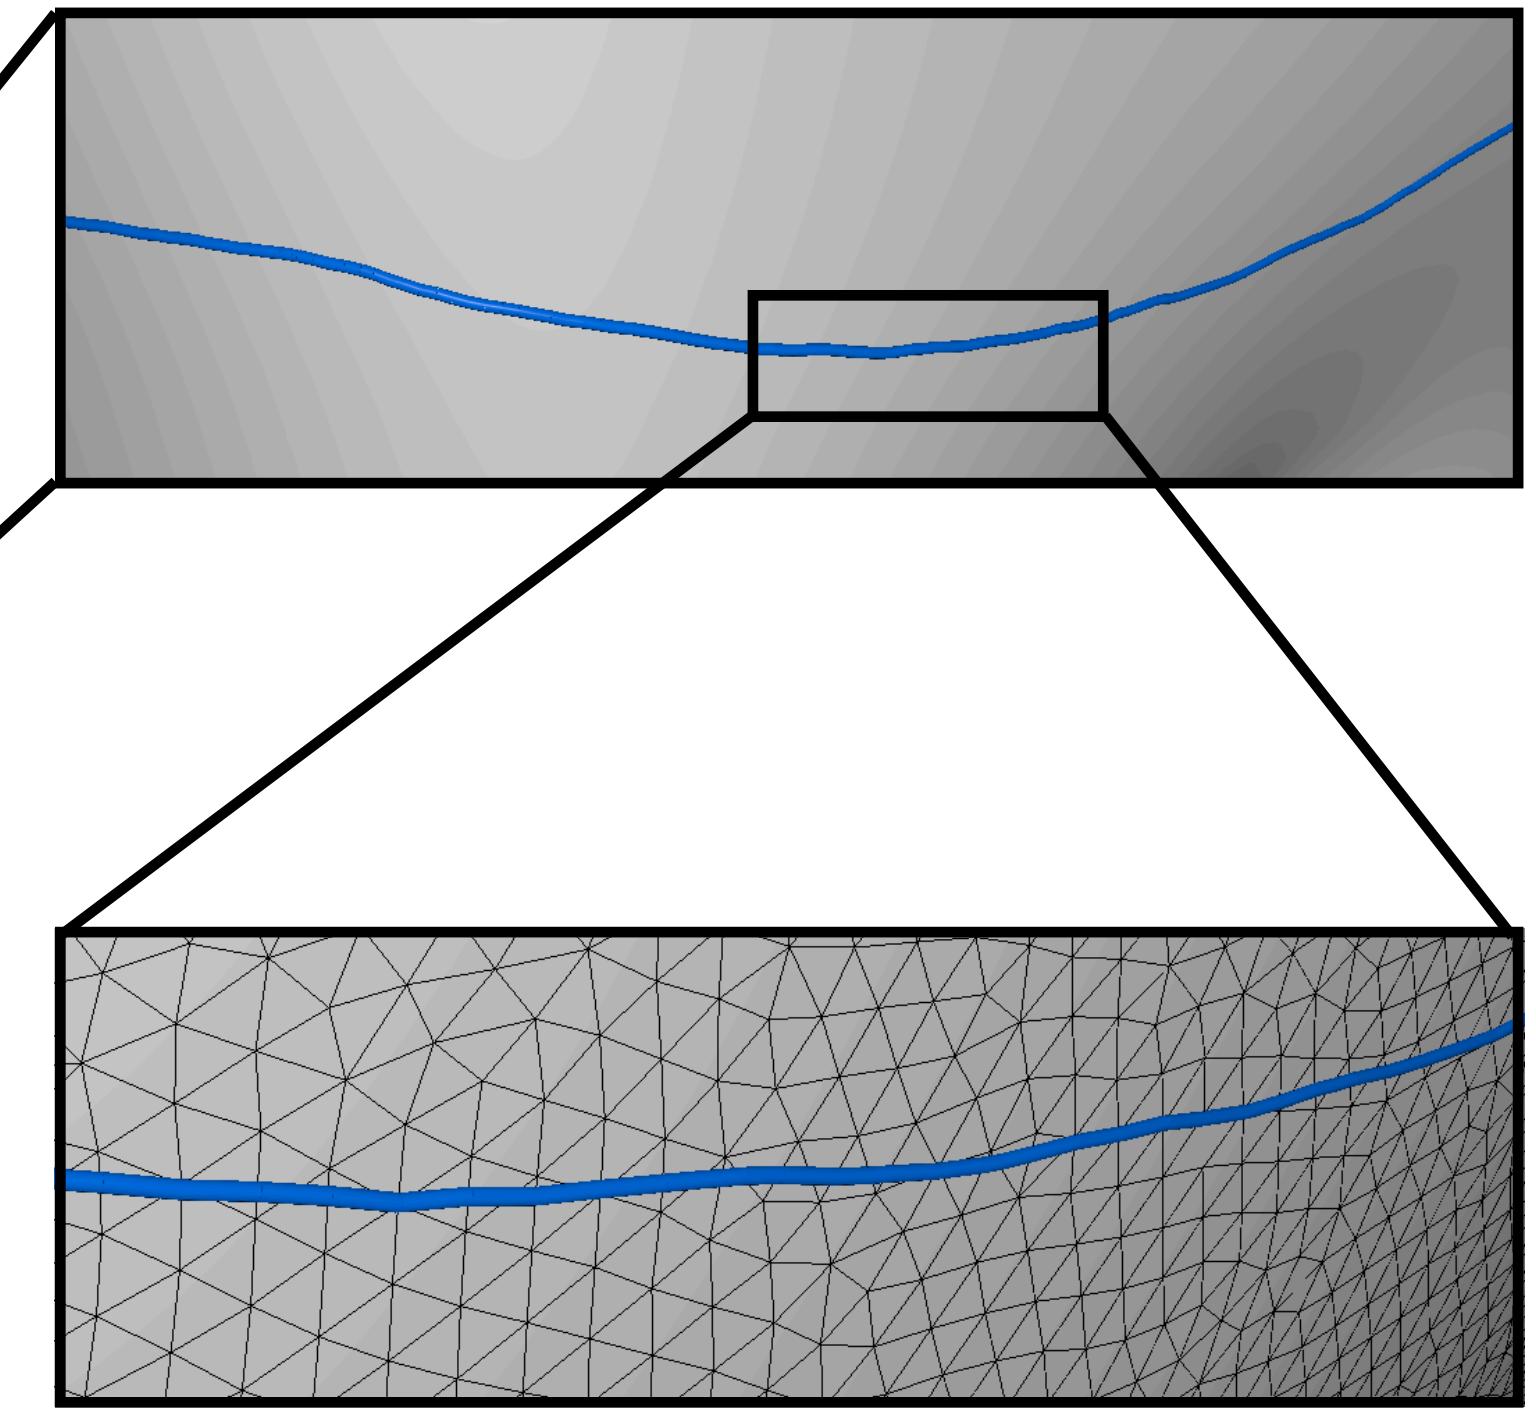
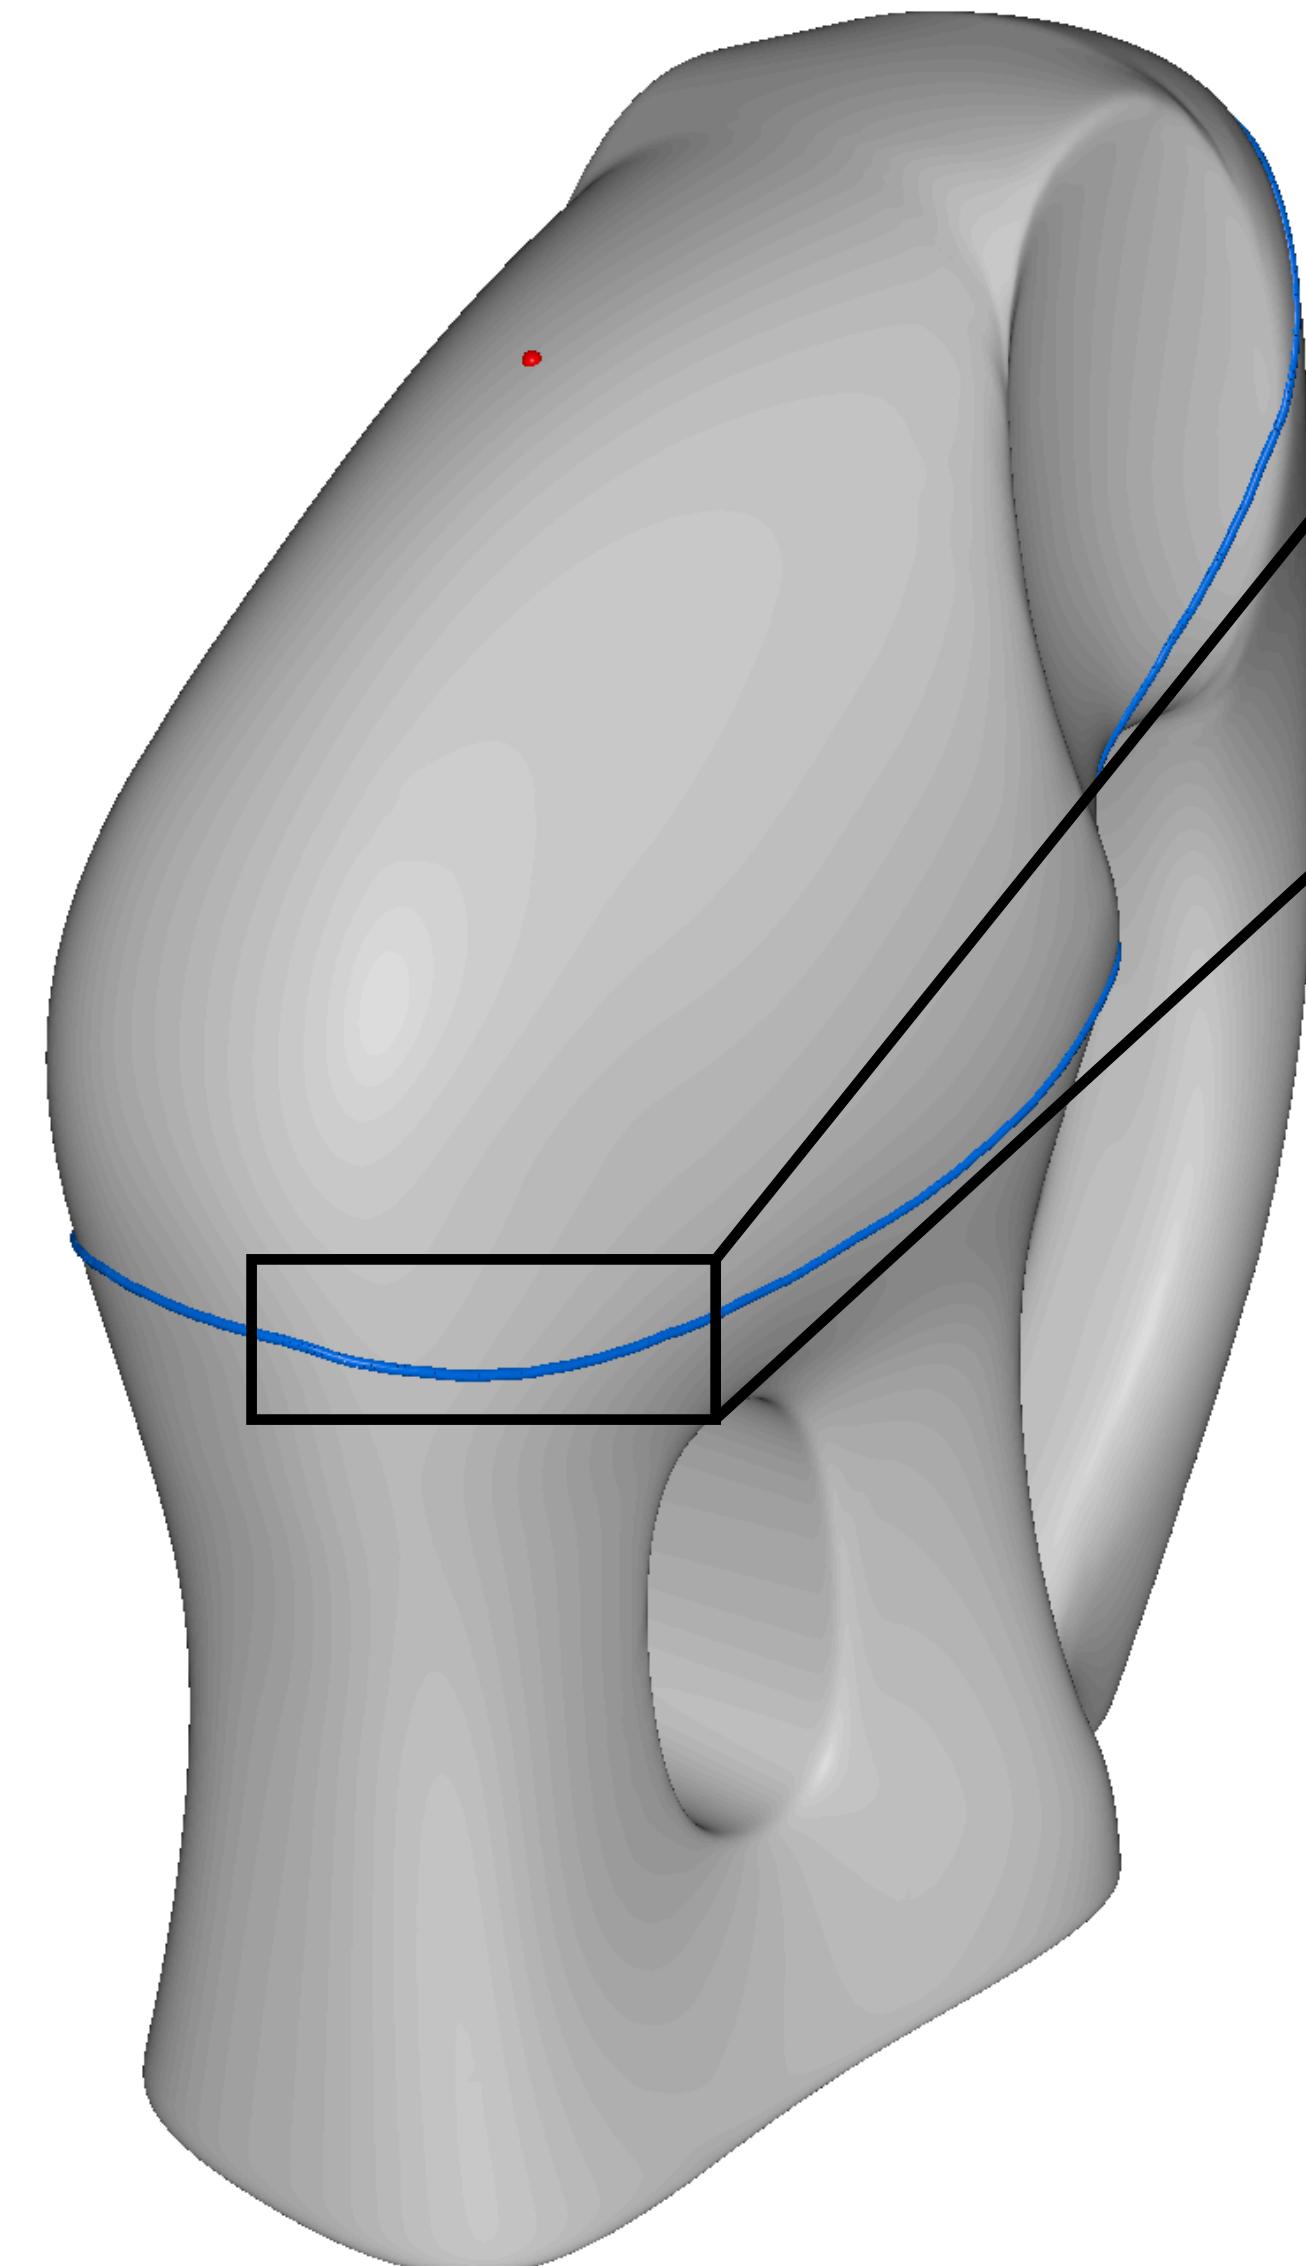
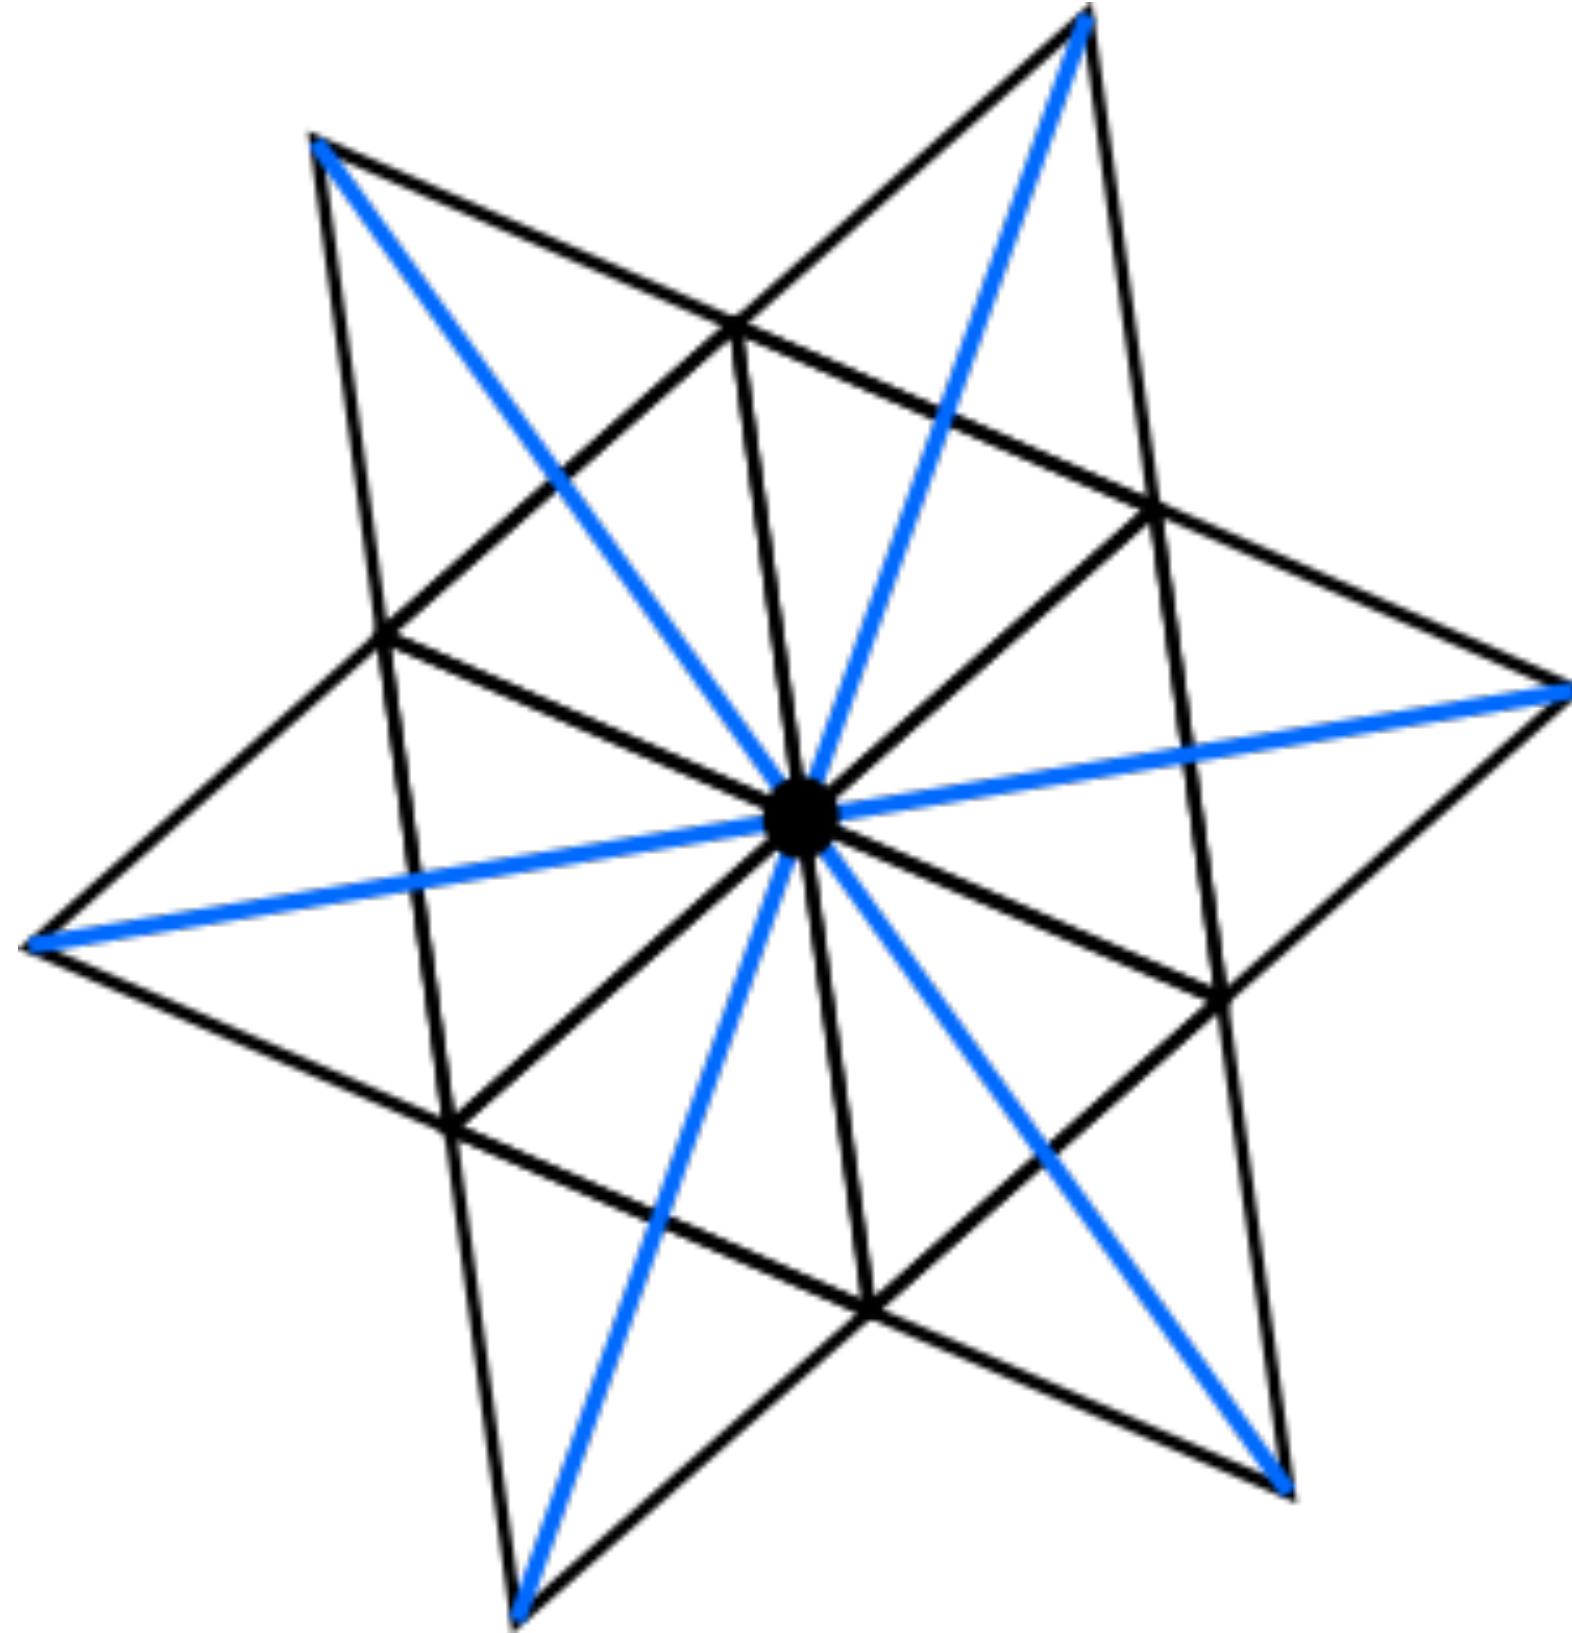


Which graph to use?

- Discrete Geodesic Graph [Wang et al. 2017, Adikusuma et al. 2020]:
 - V vertices of the mesh
 - E subset of arcs of the total graph
 - each arc is an exact shortest path between its endpoints
 - any exact path can be approximated with a path in the graph within a given tolerance ε
 - Pro: controlled accuracy, nodes are just vertices
 - Con: hard and slow to build, heavy storage (vertex valence can be $10^2\text{-}10^3$), query time is slower

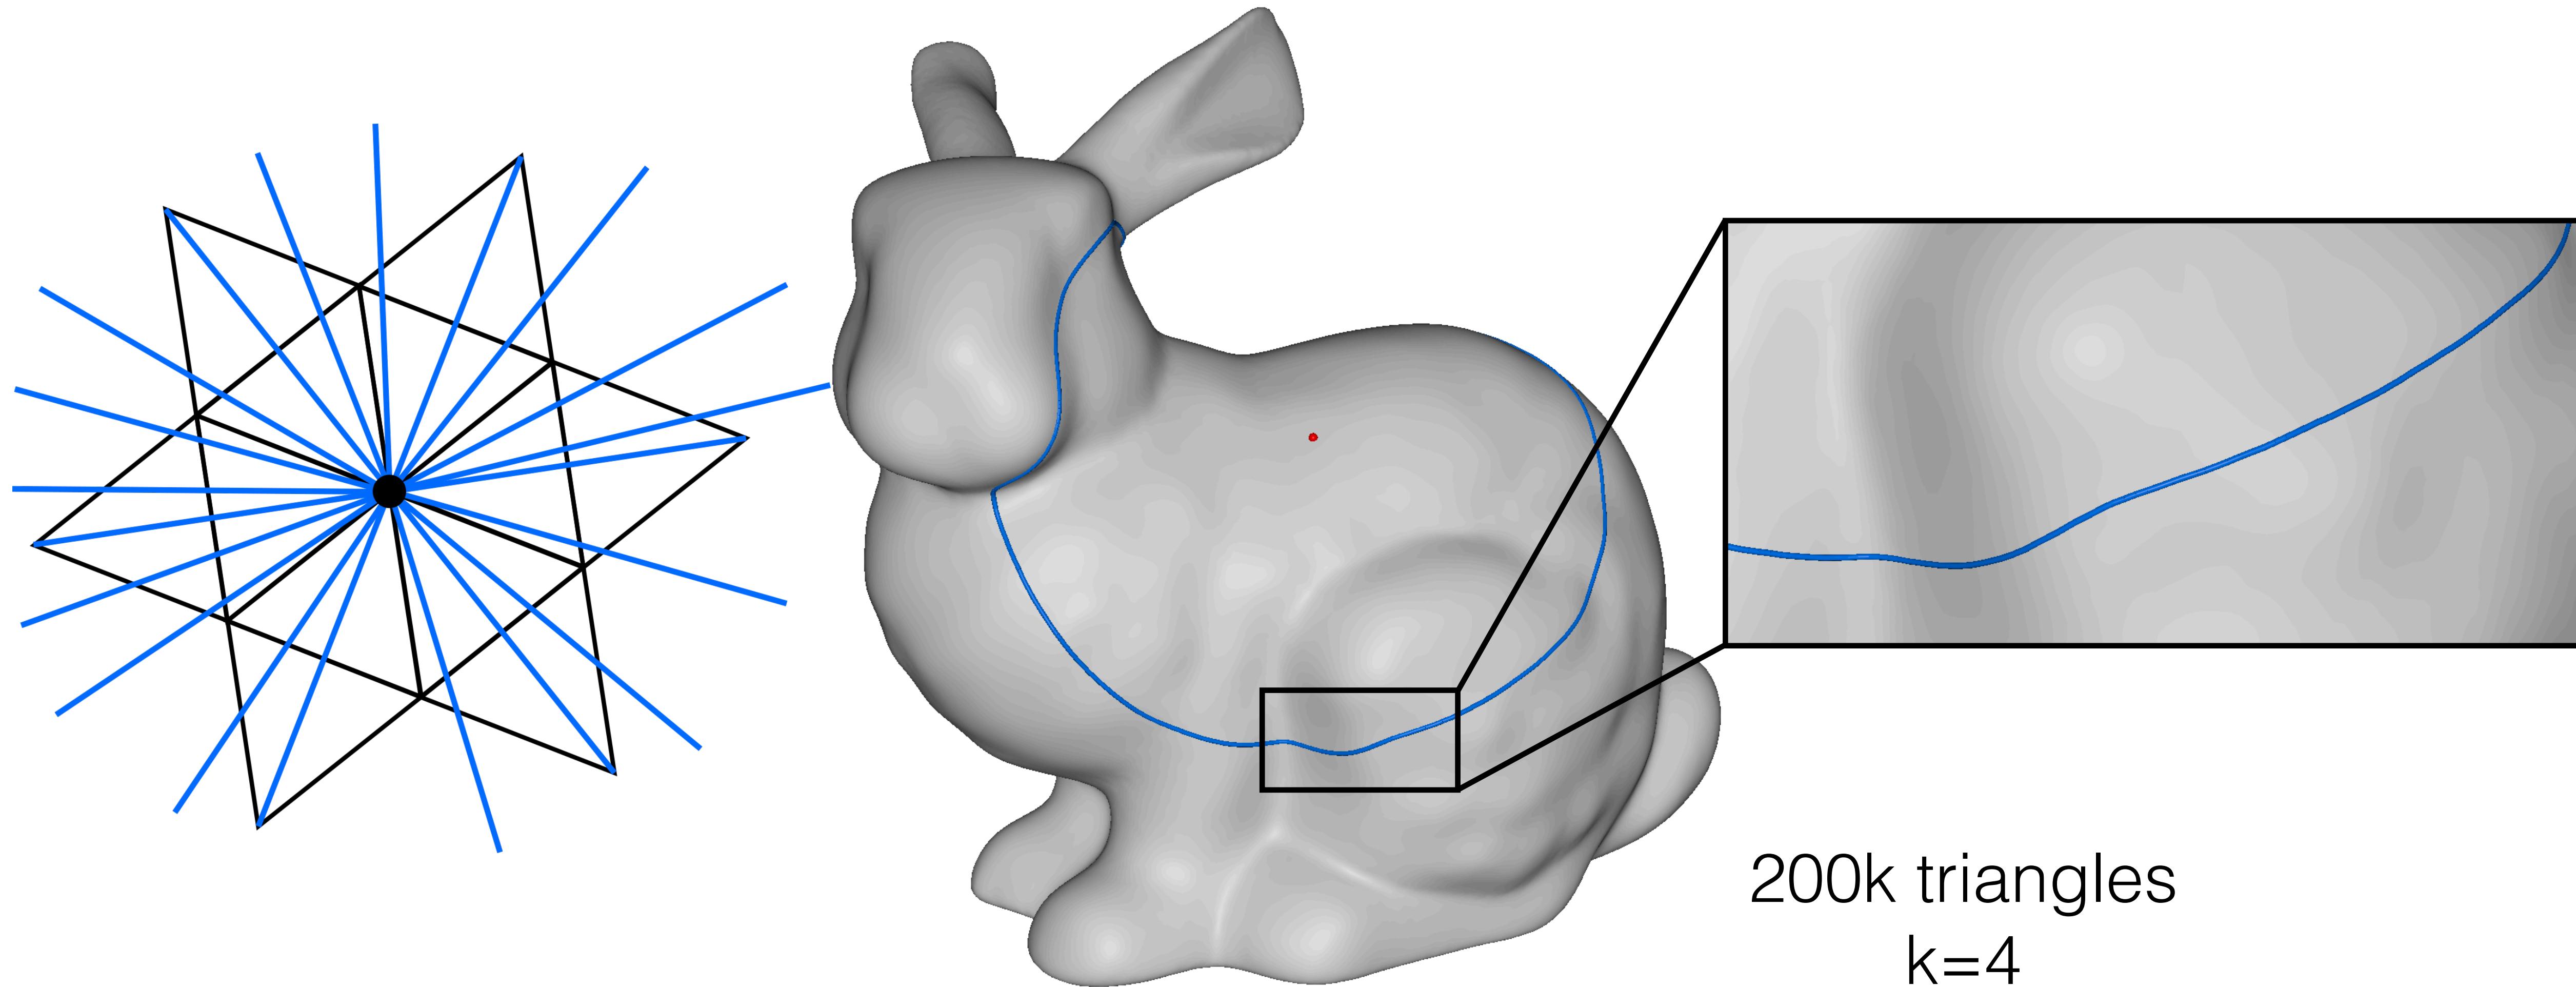


Which graph to use?

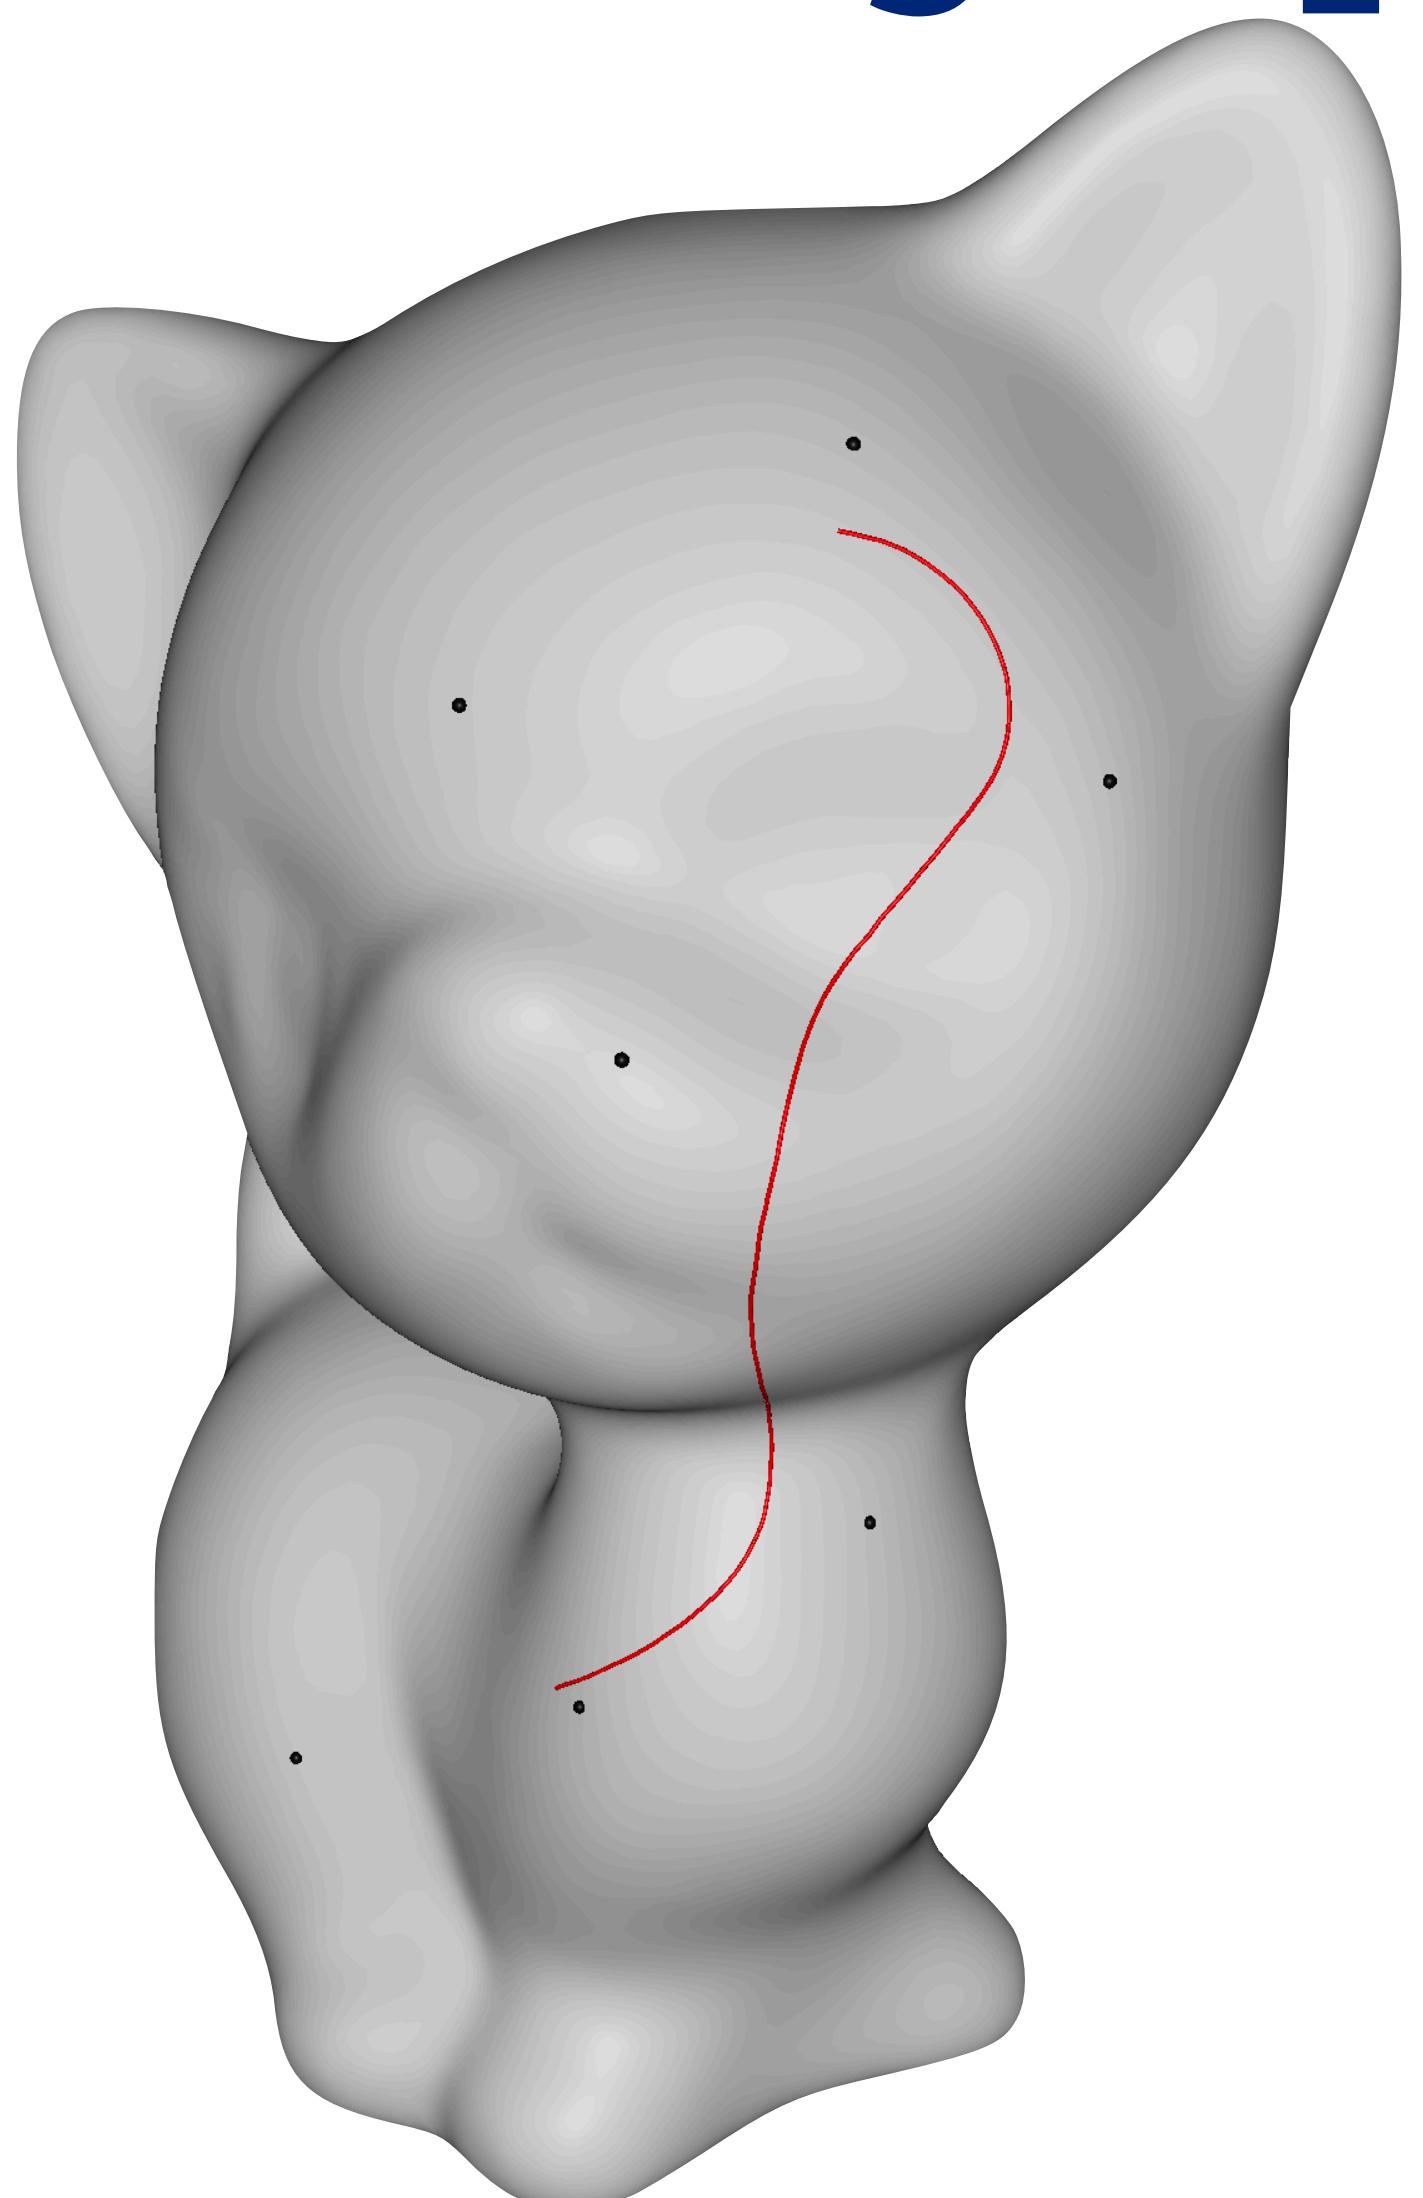


1M triangles

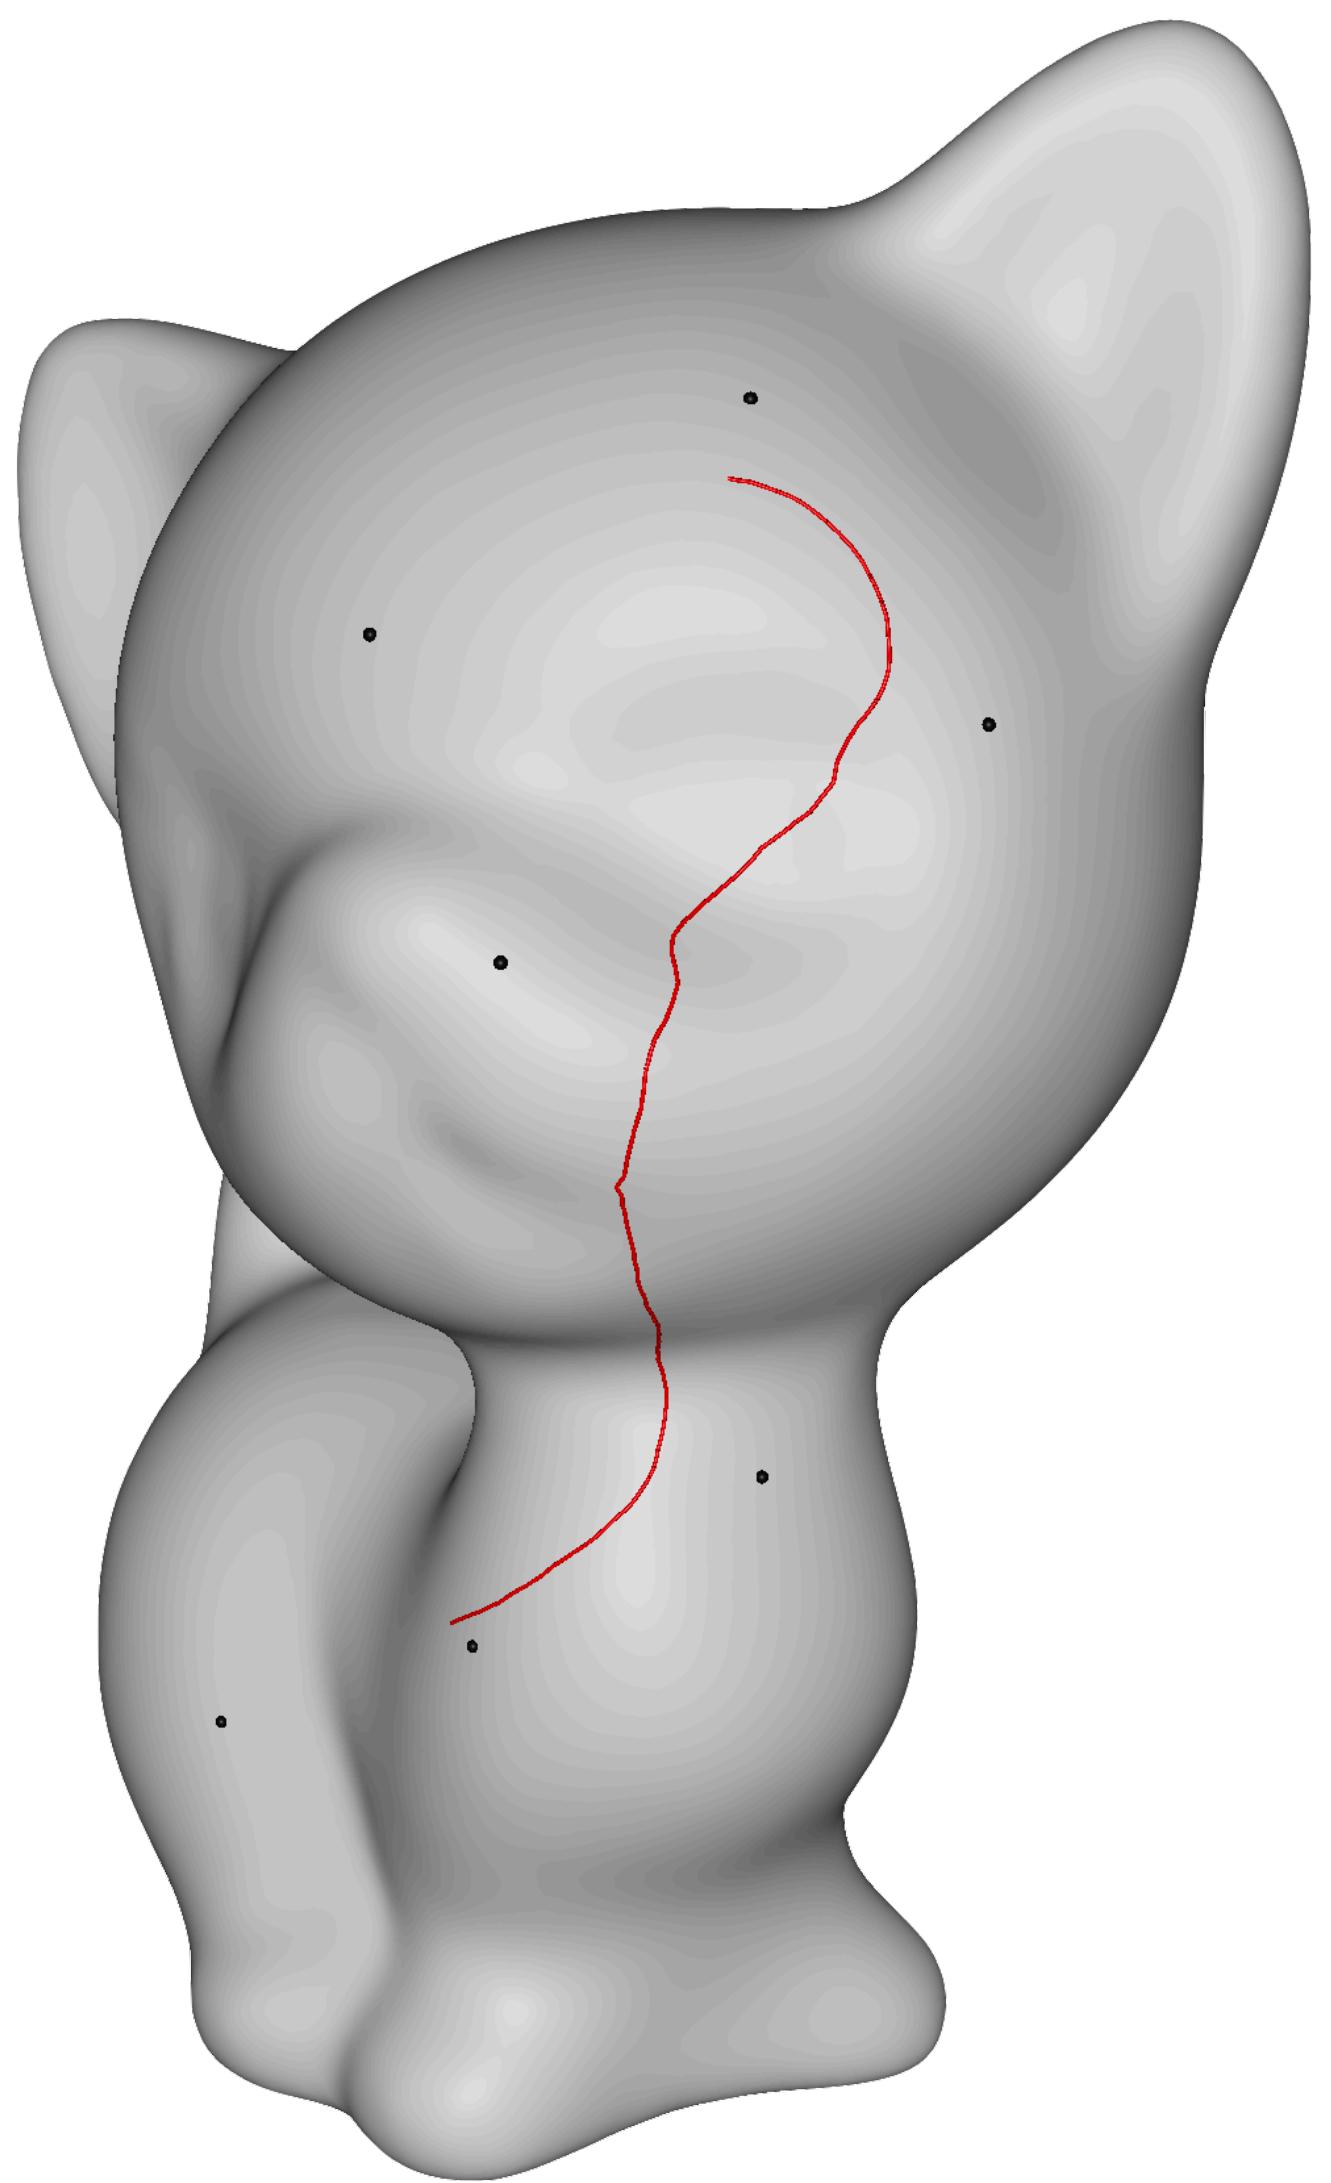
Which graph to use?



Which graph to use?



$k=6$



$k=2$

Exact polyhedral methods

Continuous Single Source Geodesic Distance/Path

- Input:
 - Mesh \mathbf{M} and source s
- Output:
 - An explicit representation of the geodesic distance function

$$d_s : \mathbf{M} \longrightarrow \mathbb{R}$$

computable at any point p on \mathbf{M} (not just a vertex)

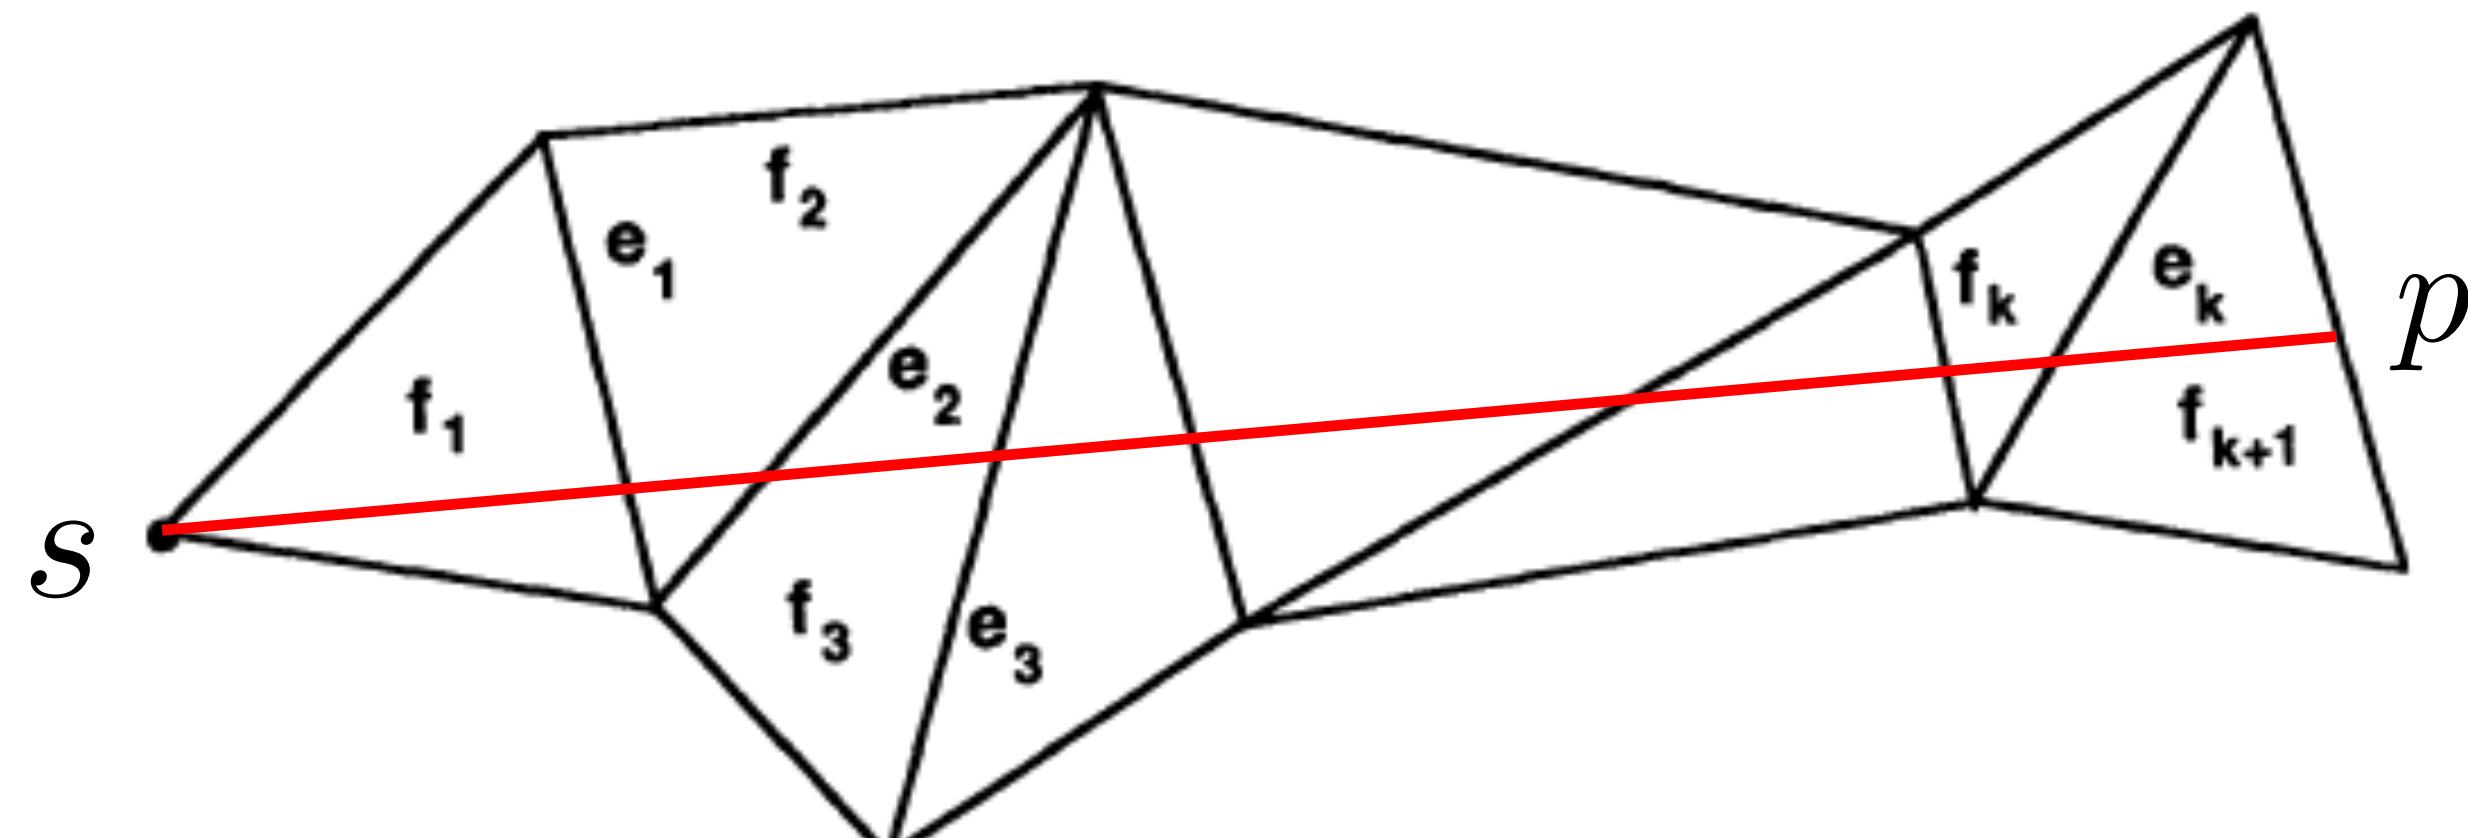
- An explicit representation of geodesic shortest paths $\delta(s,p)$

Polygonal waveform expansion

- Algorithm MMP [Mitchell, Mount, Papadimitriou, 1987] and many subsequent improvements
- Basic idea:
 - expand a front of edges starting at the neighborhood of source s until covering the whole surface
 - encode distance function d_s for points on each edge e as a set of windows splitting e

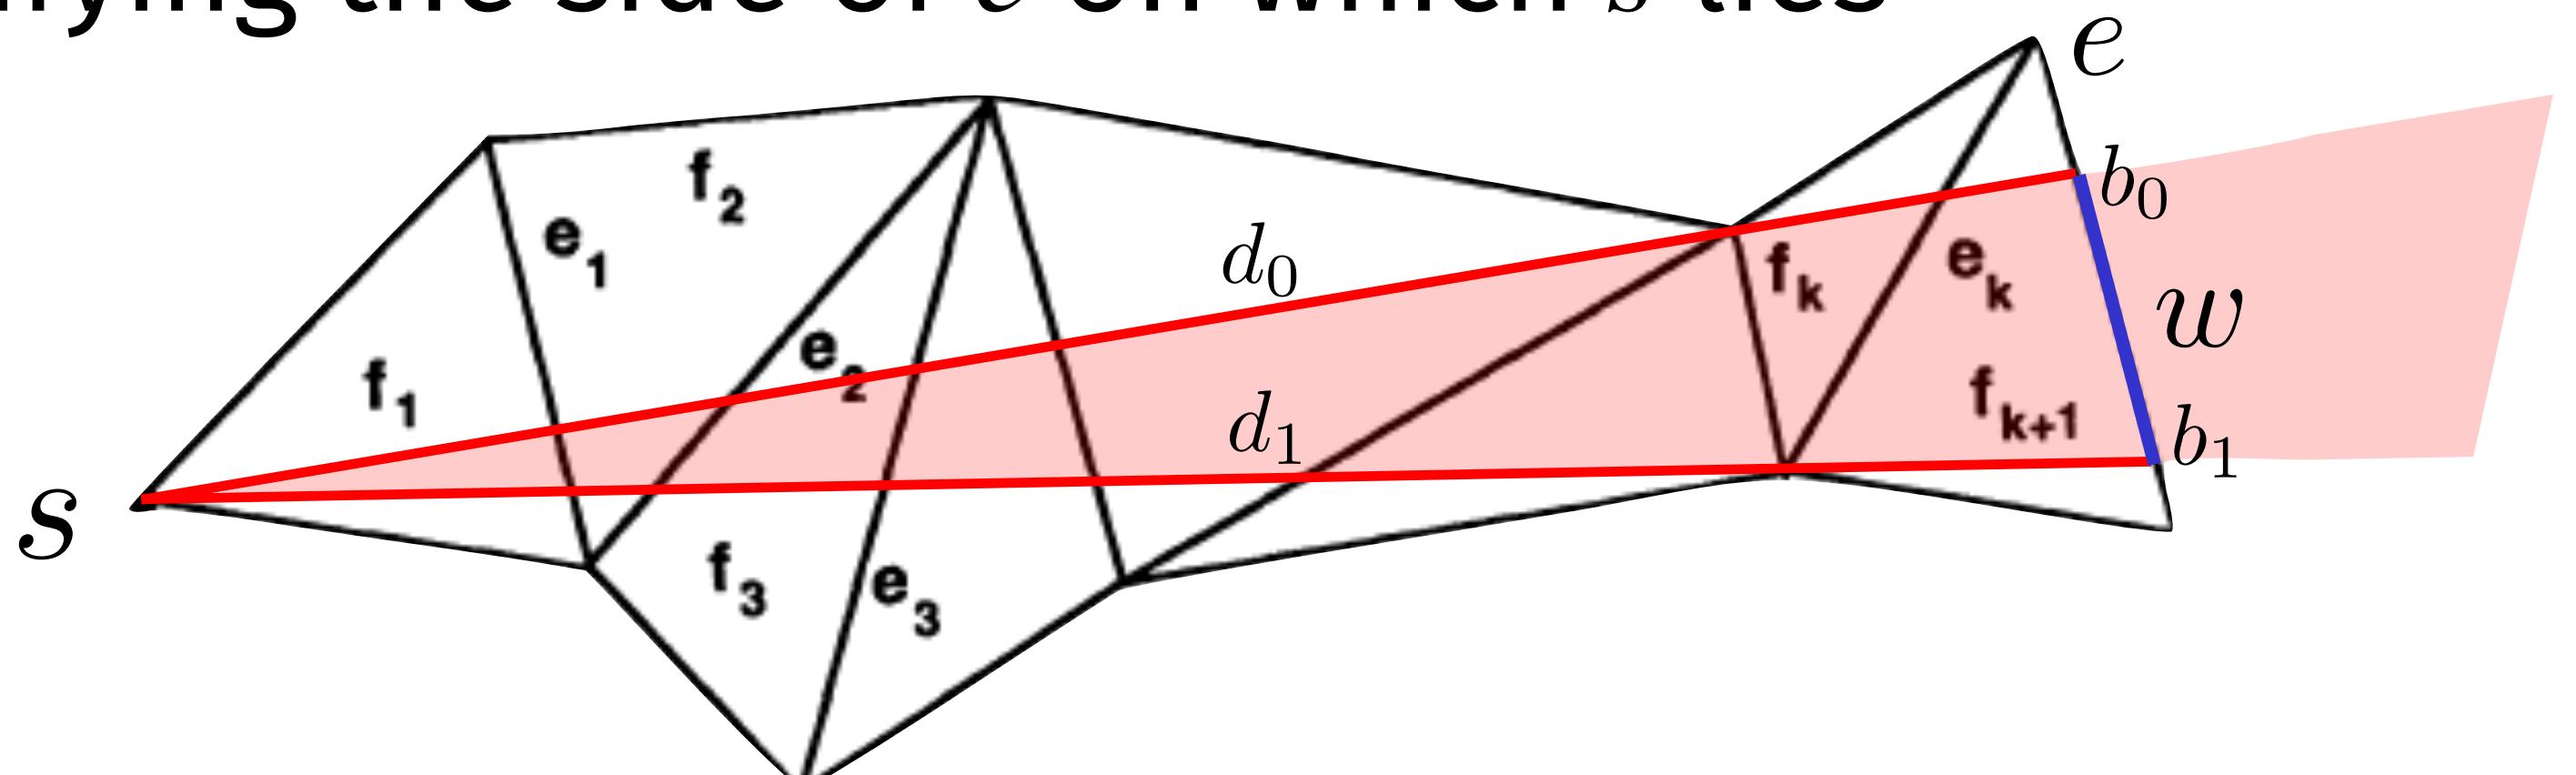
Geodesics through a triangle strip

- Assume s is a vertex (if not, add it as a vertex by splitting the triangle containing it on-the-fly)
- Consider a triangle strip starting at s and flatten it to the plane
- Assume there exists at least one straight line from s contained inside the strip and crossing all its triangles (if not, we are not interested in that strip)
- The folding of such line to \mathbf{M} is a shortest geodesic connecting s to its other endpoint p



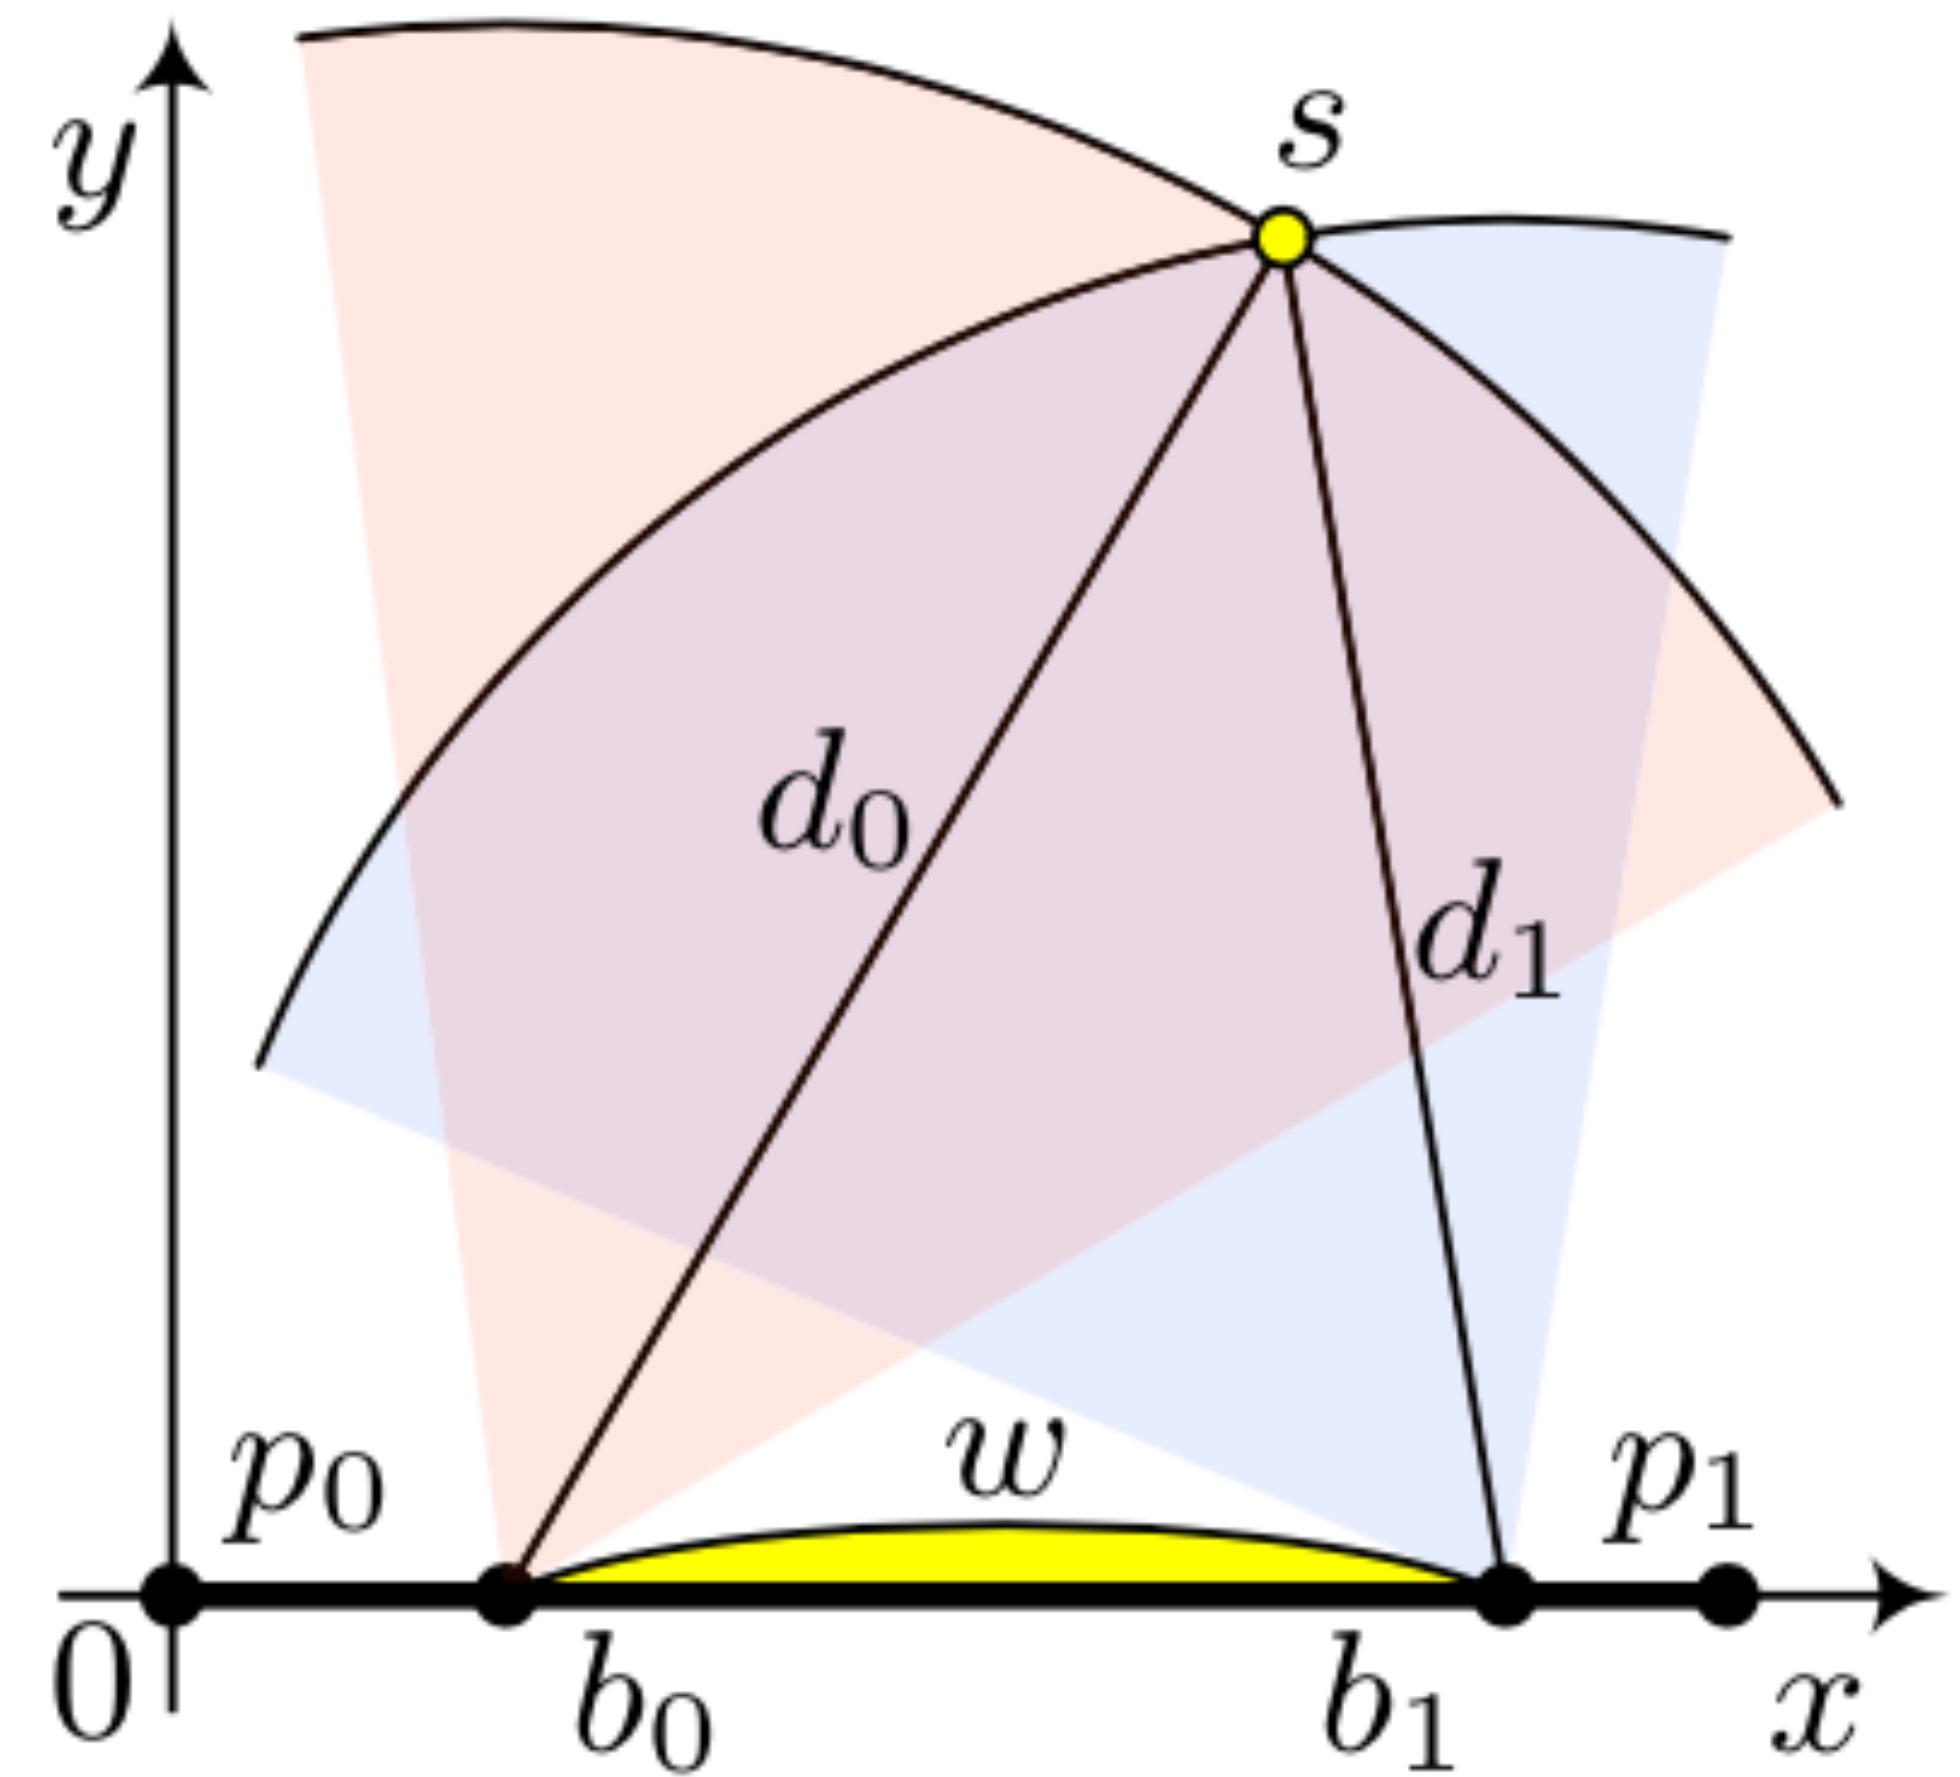
Geodesics through a triangle strip

- The set of all lines from s through the strip, and piercing the same edge e on the opposite end of it, form a wedge, encoded as a *window* w on e :
 - parametric coordinates b_0 and b_1 of endpoints of w on e
 - distances d_0 and d_1 of endpoints of w from s
 - binary direction τ specifying the side of e on which s lies



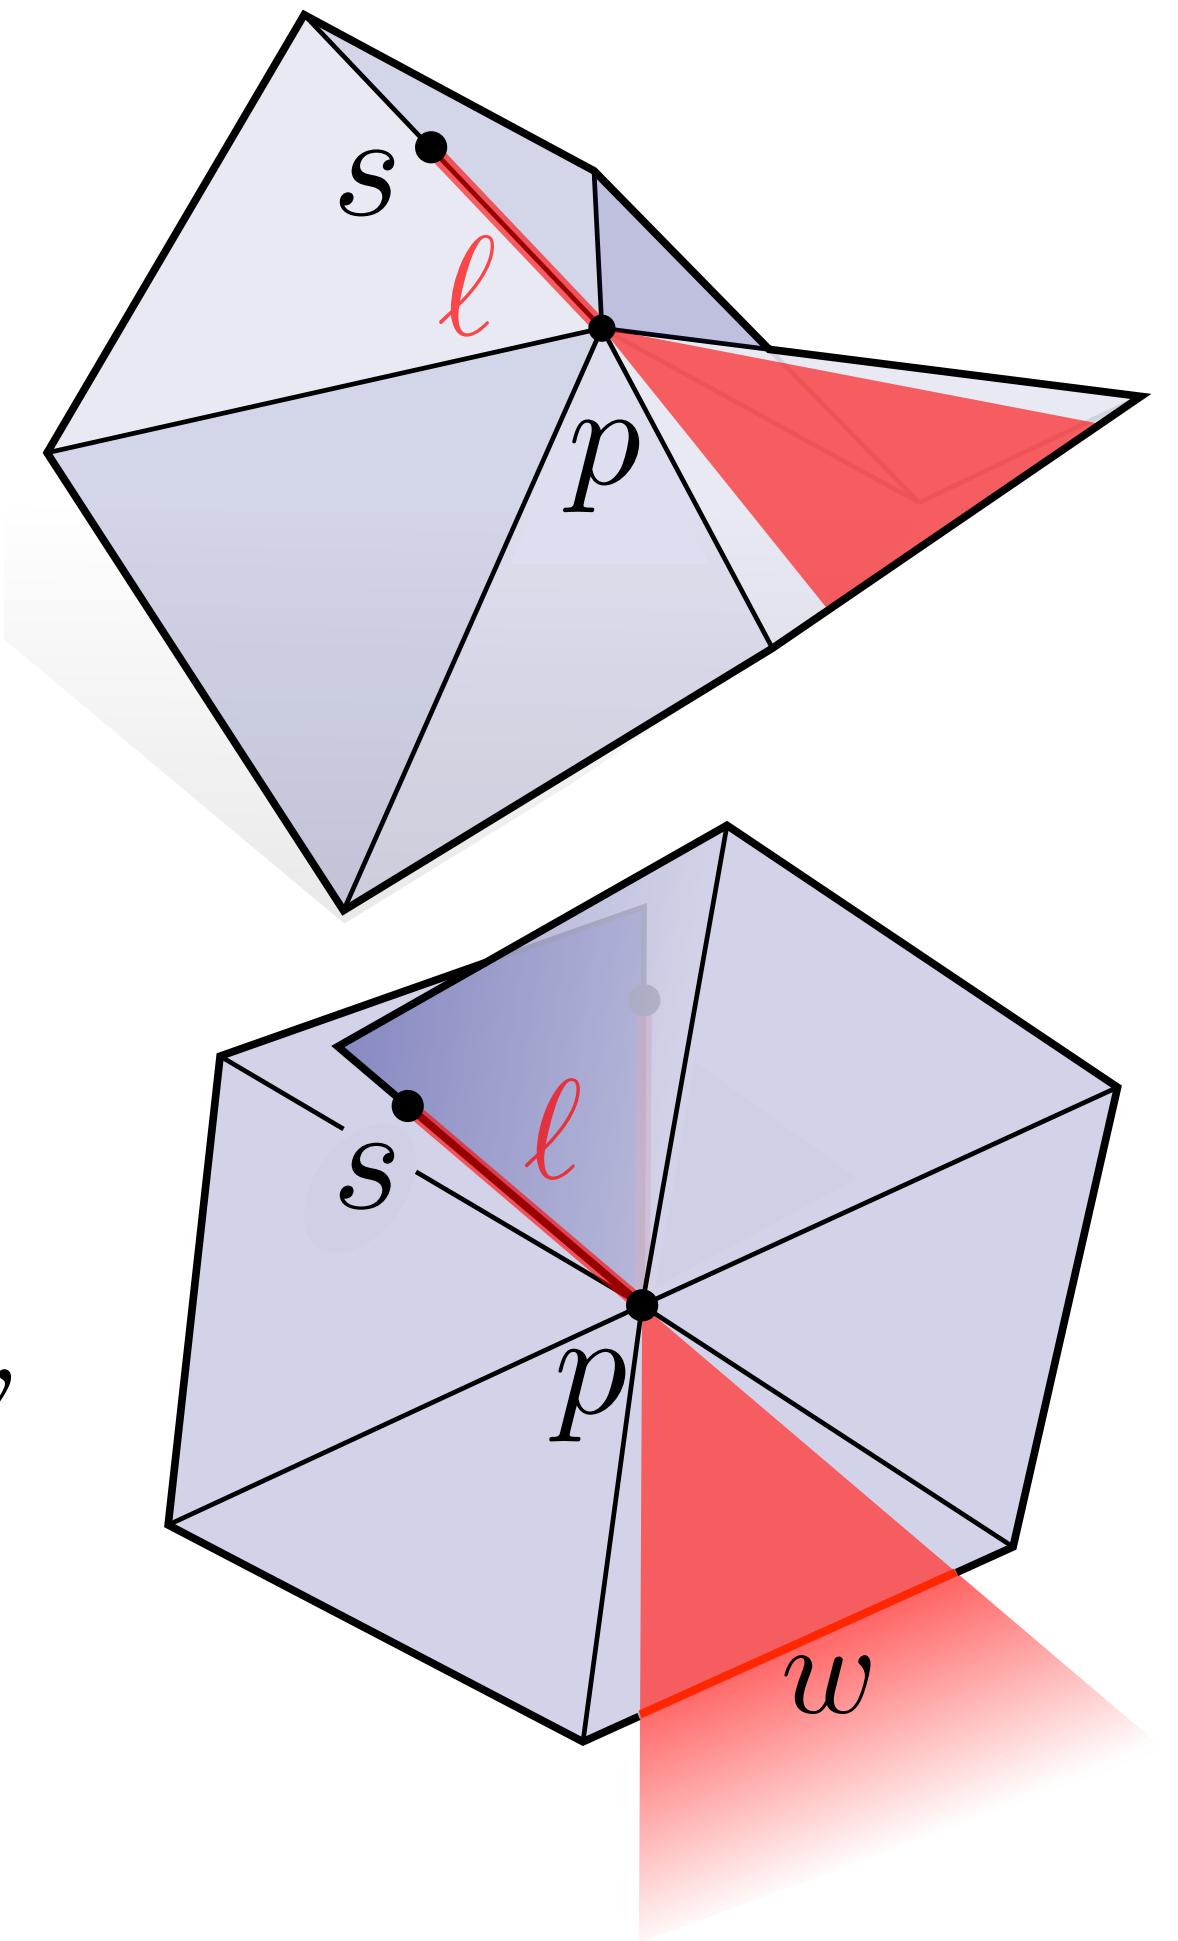
Distance function in a window

- Coordinates (x_s, y_s) of s are computed by intersecting the two circles centered at endpoints b_0 and b_1 :
 - $(x_s - b_0)^2 + y_s^2 = d_0^2$
 - $(x_s - b_1)^2 + y_s^2 = d_1^2$
- $d_s(x) = \sqrt{(x_s - x)^2 + y_s^2}$ for $x \in [b_0, b_1]$



General case

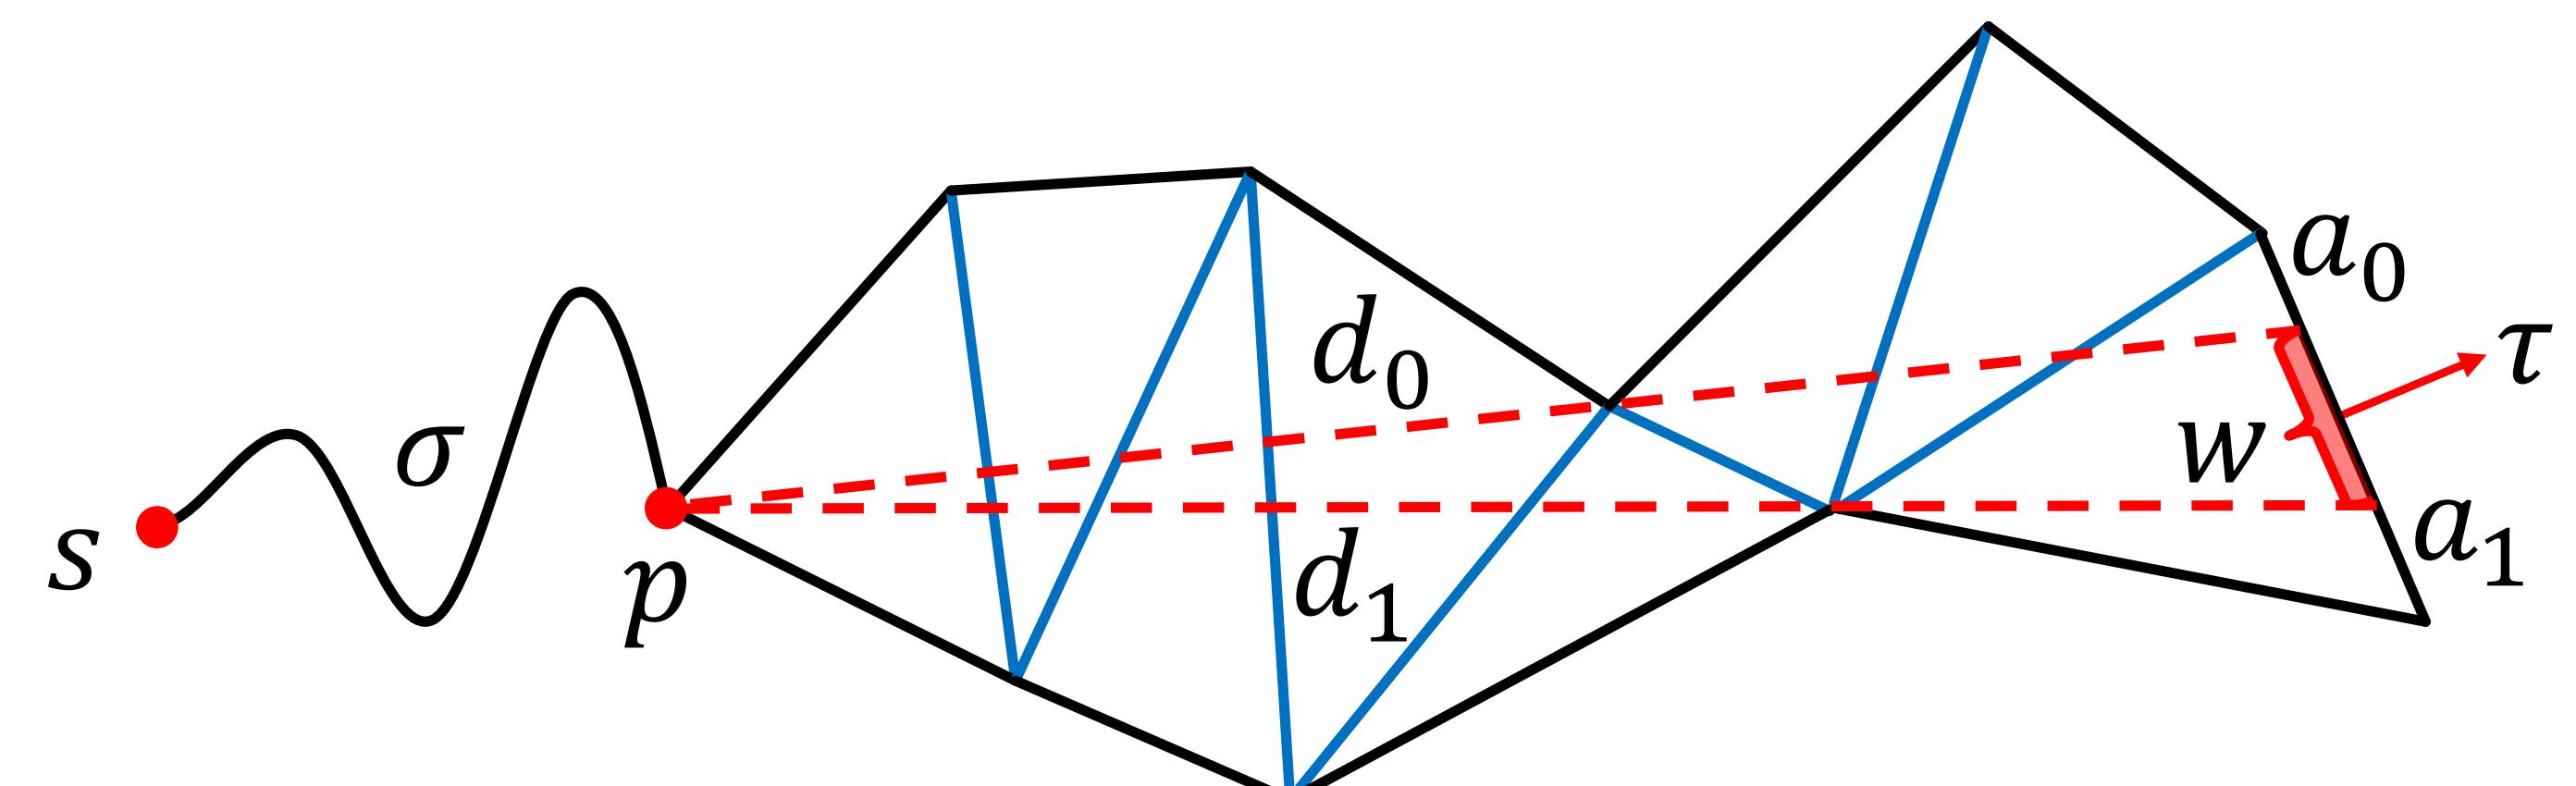
- Not all points can be reached from straight lines through strips
- What happens at saddle vertices?
 - the unfolding of the 1-ring overlaps
 - points in the red wedge cannot be reached by straight lines from s
 - we need to turn about p to reach the red window w
 - line ℓ is a mandatory route to all points of w



General case

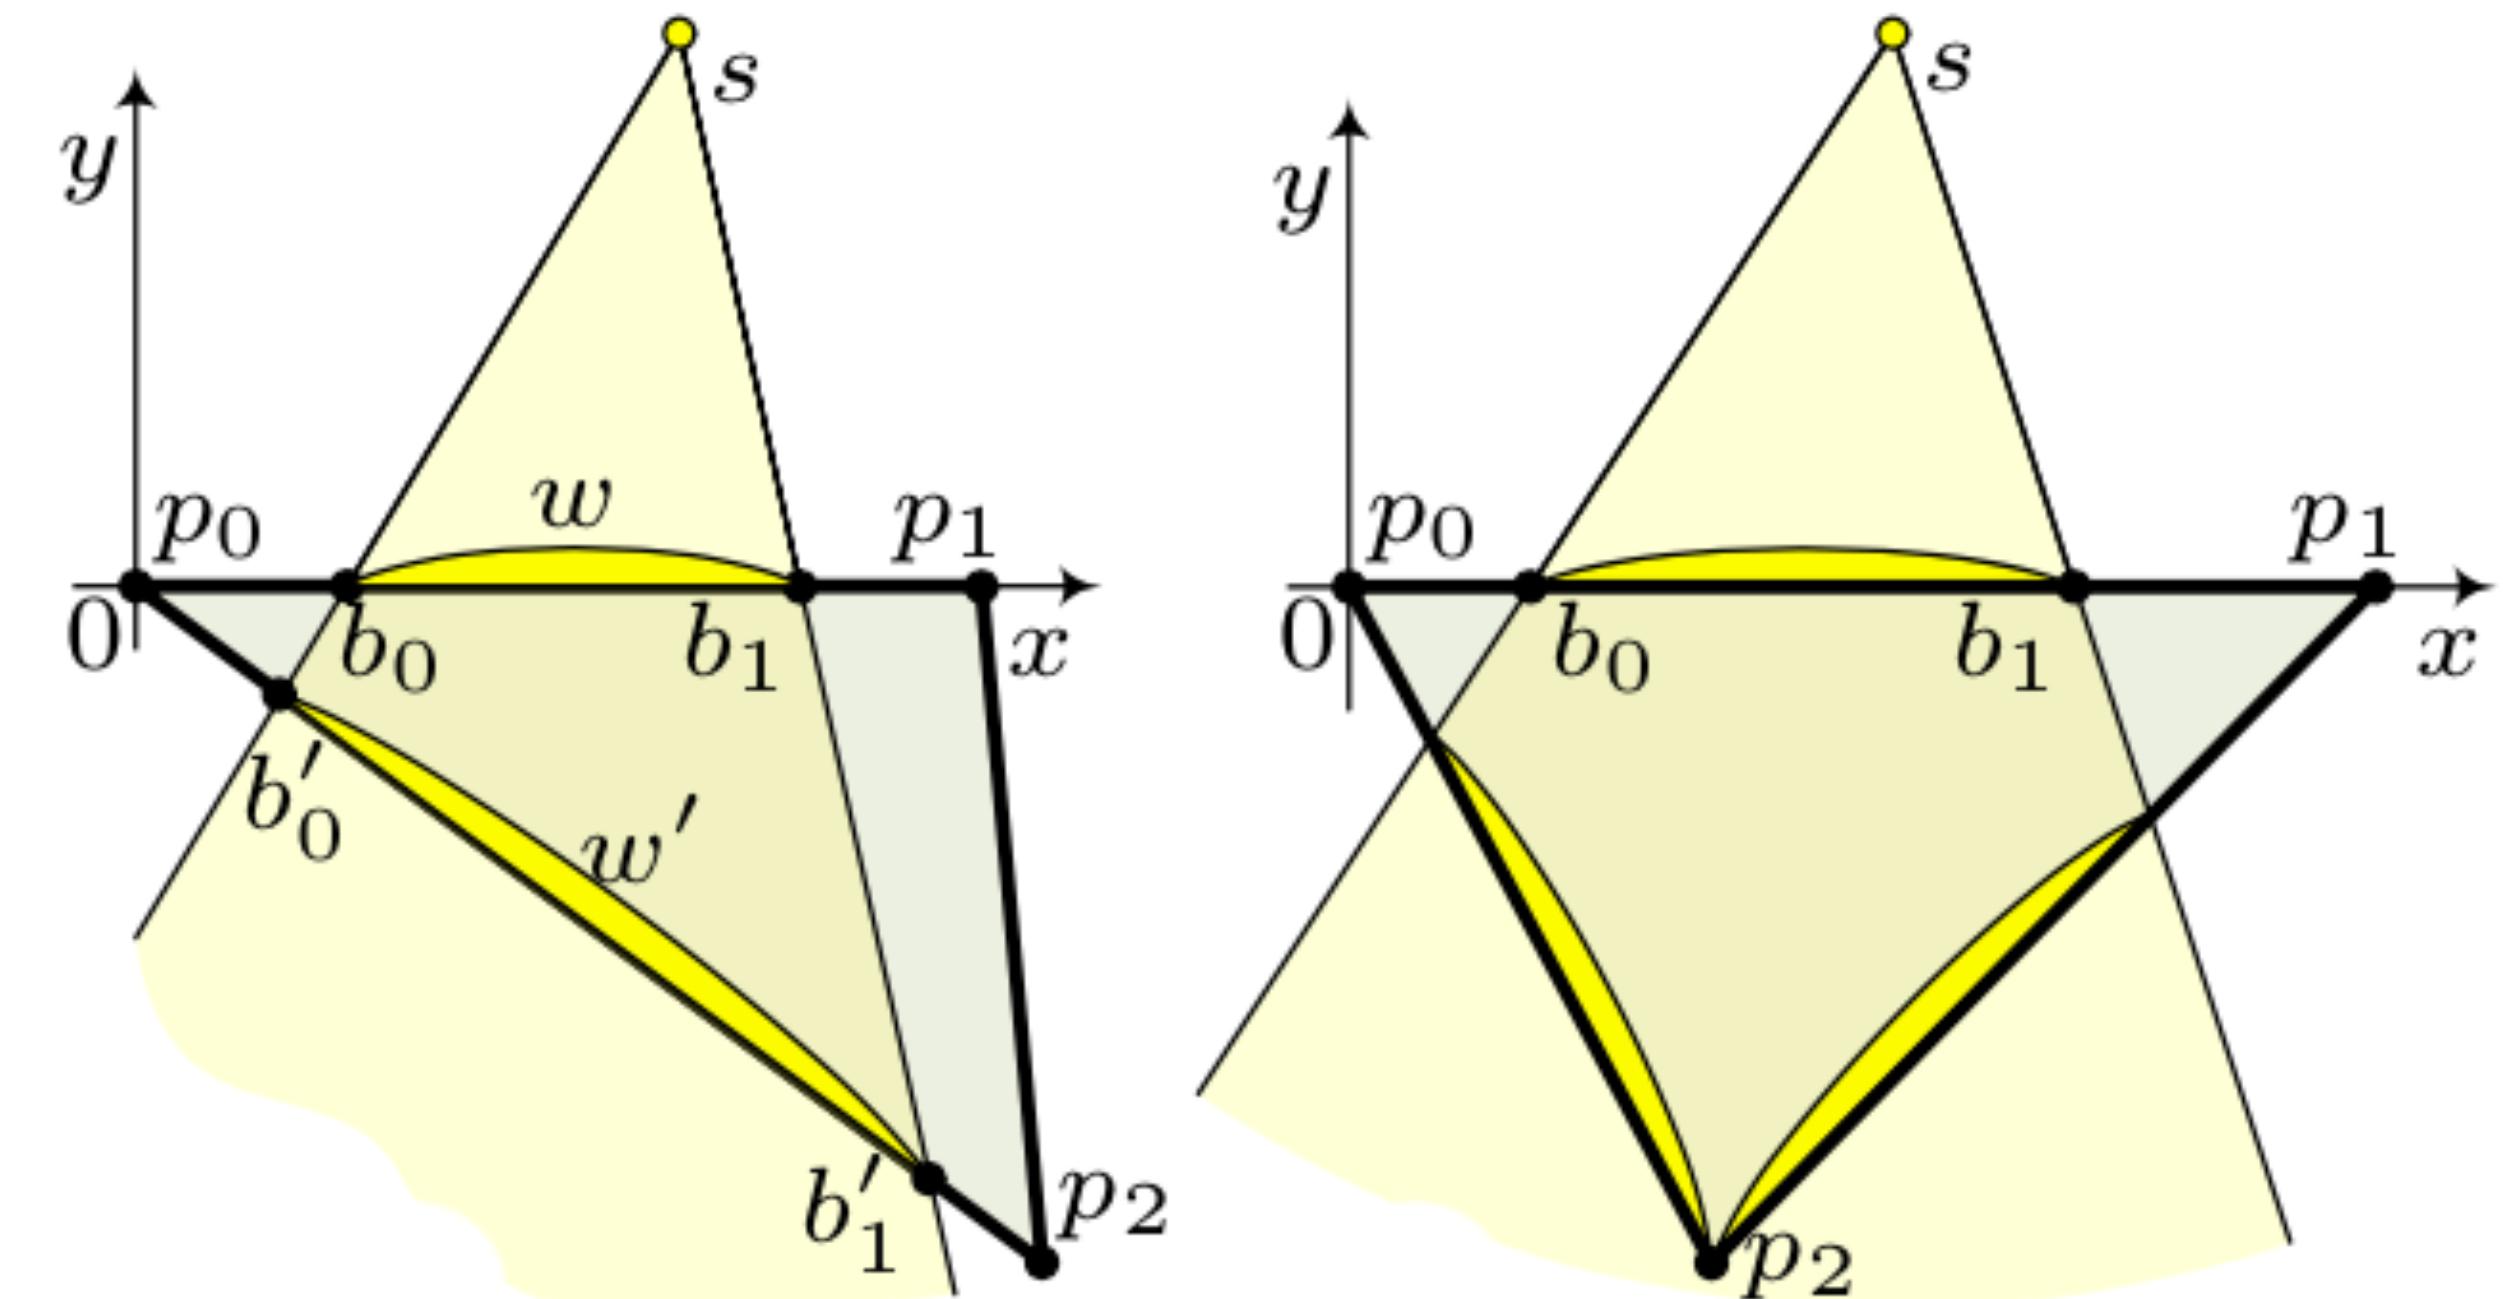
- We decompose the paths to a generic window w into a *funnel* consisting of two parts:
 - a polygonal path of length σ from source s to *pseudo-source* p
 - a wedge from p to w (as before)
- The distance function for all x in w is encoded as:

$$(b_0, b_1, d_0, d_1, \sigma, \tau)$$



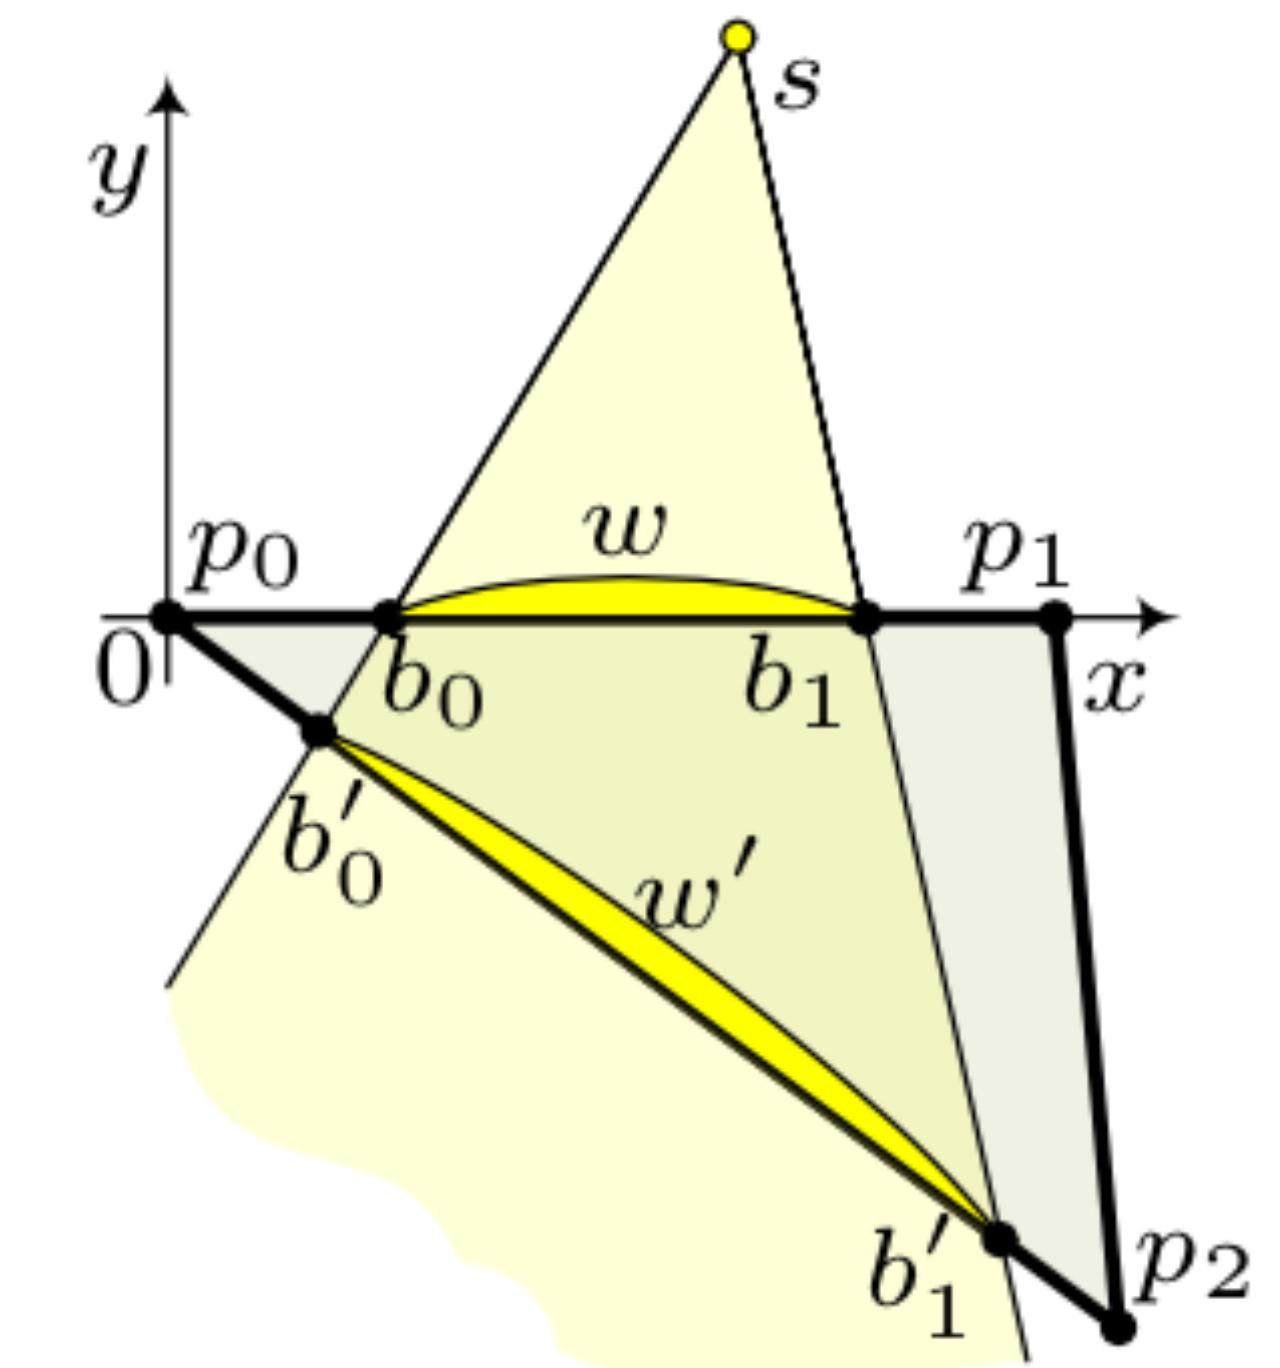
Window propagation

- Given a window w on an edge e_1 , propagate its distance field across an adjacent triangle t to define new potential windows on the two opposing edges e_2, e_3



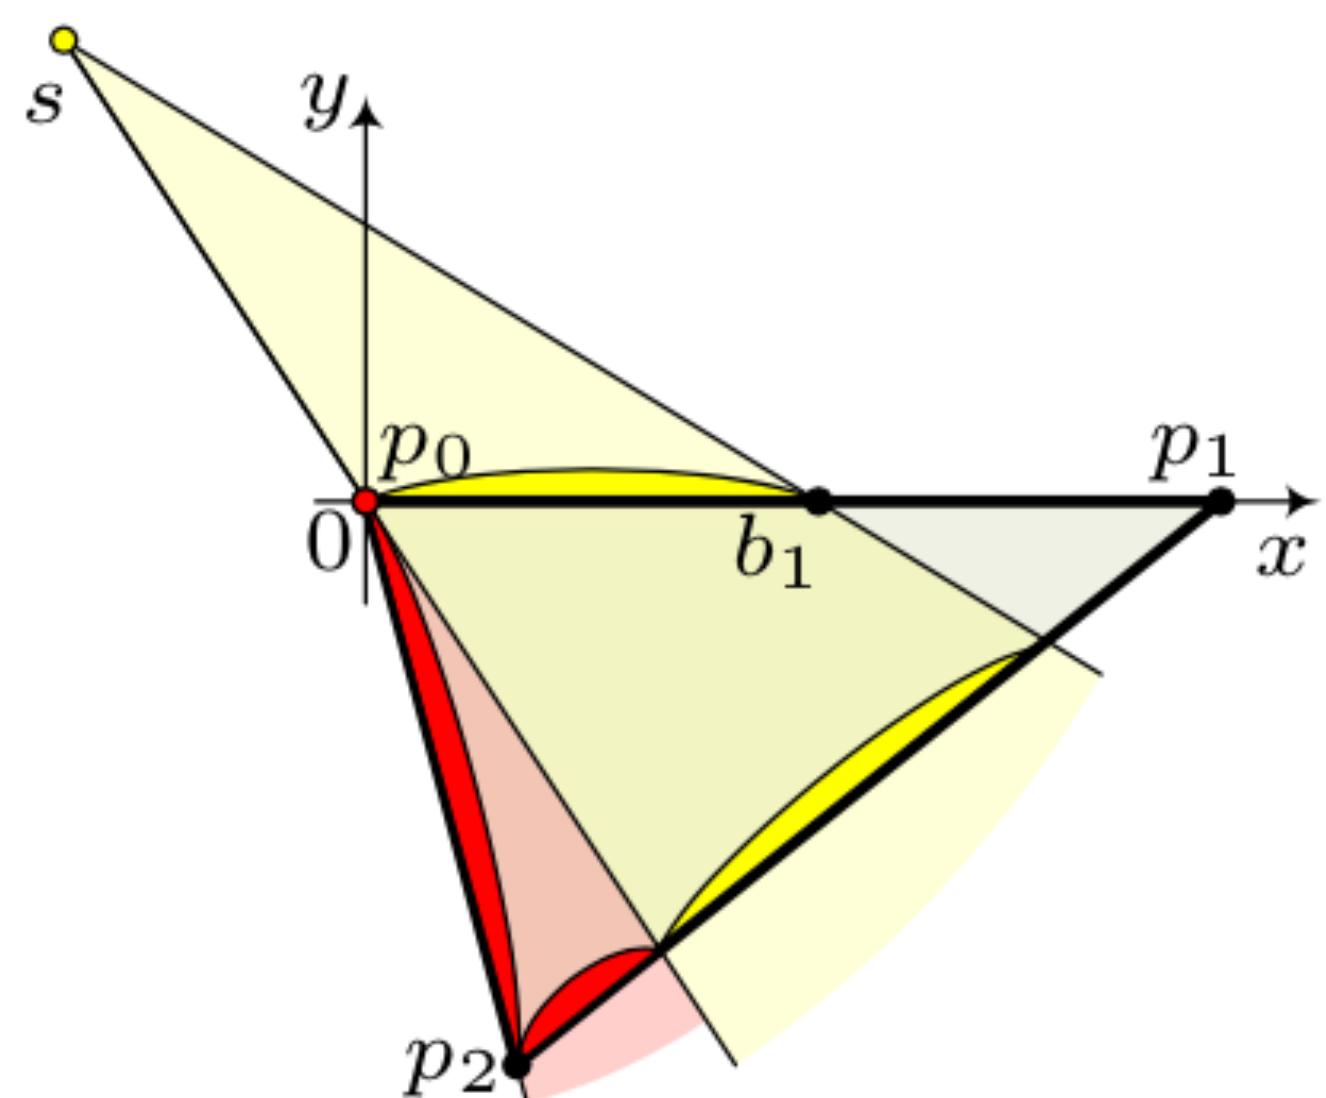
Window propagation

- To define the distance field over w' :
- Extend the rays from the pseudo-source s through the endpoints of w and intersect them with the new edge, to obtain the new interval $[b'_0, b'_1]$
- Compute the new distances d'_0, d'_1 from pseudo-source s
- Pseudo-source distance $\sigma' = \sigma$ is unchanged
- Direction τ is assigned to point inside triangle t



Window propagation

- Special case if one of the endpoints of w is a saddle vertex p_0 :
 - Consider the other edge \bar{e} of face t incident at p_0
 - If \bar{e} lies outside the funnel, then generate the two new red windows in the figure, having p_0 as pseudo-source
 - Distance σ must be updated $\sigma = \sigma + d_0$



Window propagation

- Each new window w_0 may overlap other windows already computed for the same edge
- Let us suppose window w_1 on edge e overlaps w_0 and let

$$\delta = w_0 \cap w_1 = [b_0, b_1]$$

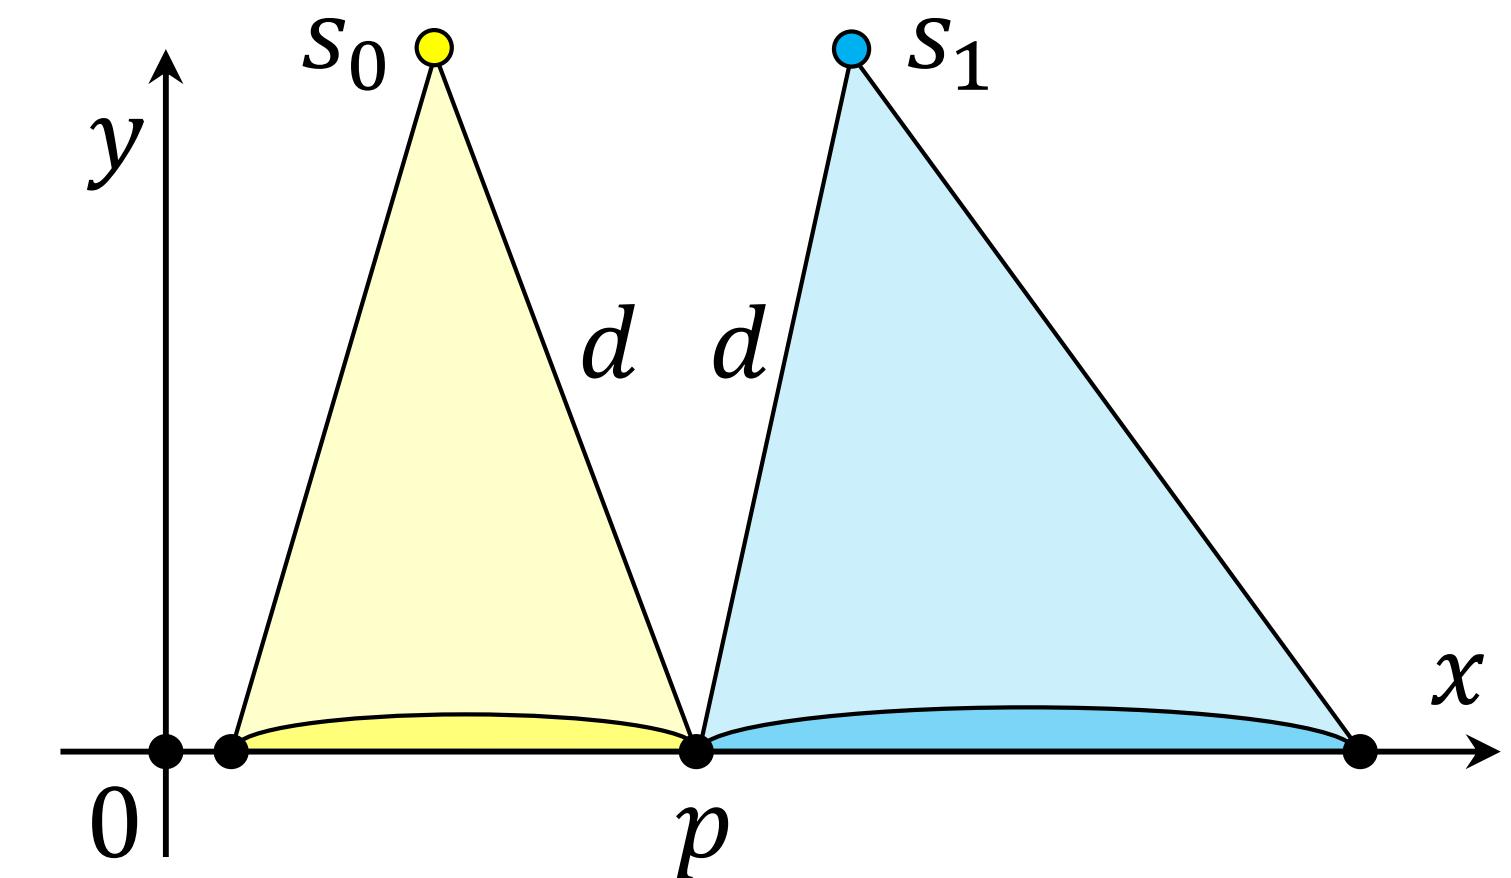
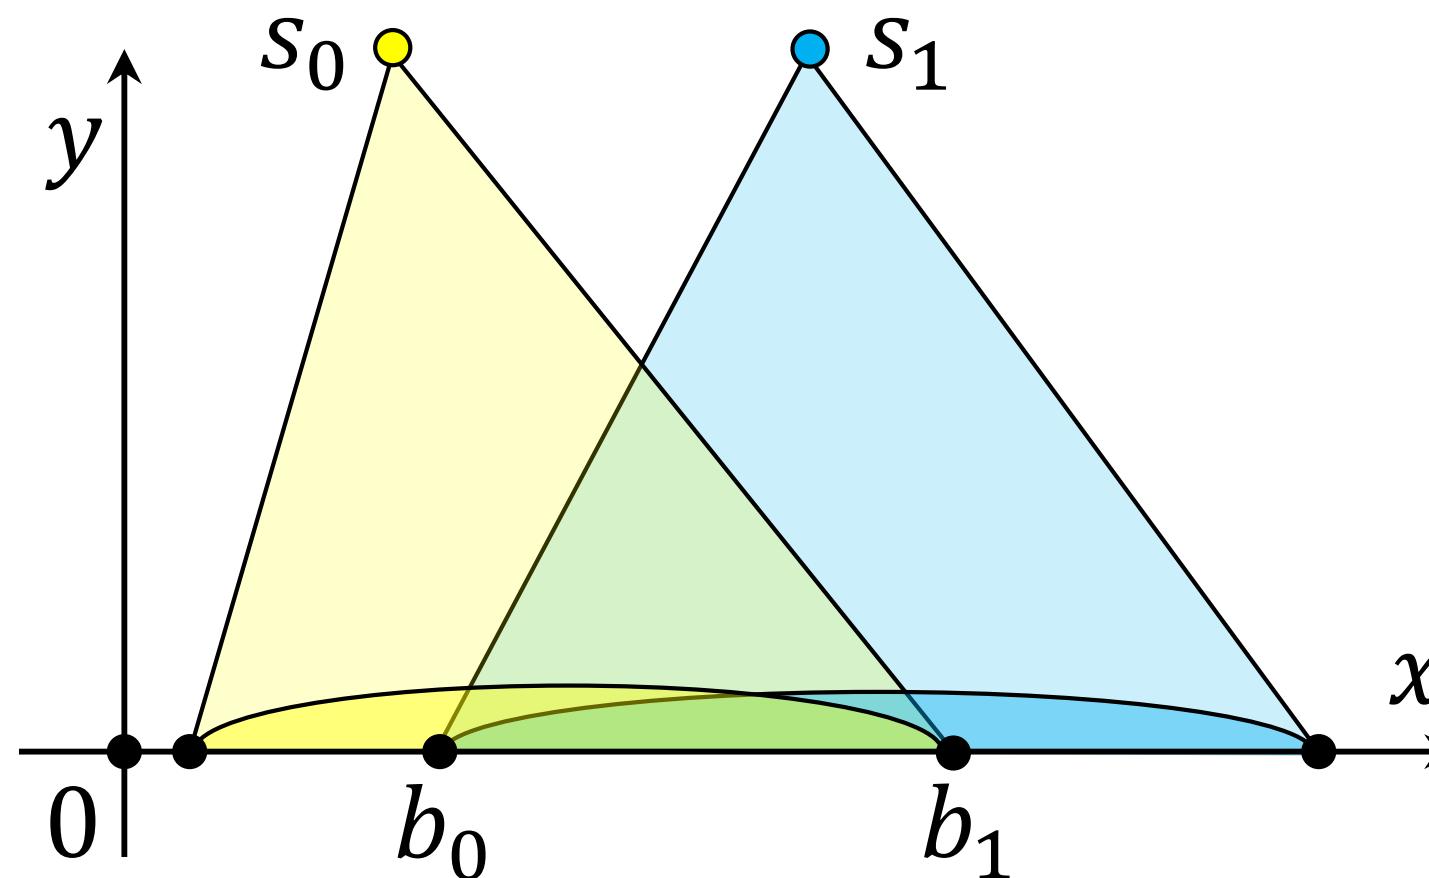
- We must decide which window defines the minimal distance function for each point in δ , and update the windows accordingly

Window propagation

- Find point $p \in [b_0, b_1]$ where the two windows define the same distance

$$\sqrt{(p_x - s_{0x})^2 + s_{0y}^2} + \sigma_0 = \sqrt{(p_x - s_{1x})^2 + s_{1y}^2} + \sigma_1$$

- Solve a quadratic equation to find point p (exercise)



Continuous Dijkstra propagation

- Maintain a priority queue Q of windows
- Priority is minimum distance from source
- Initialization of Q :
 - If s is inside triangle t , then compute one window per edge of t and insert them in the queue
 - If s is on an edge e , then compute one window per edge of the two faces incident at e and insert them in the queue. Also compute two windows on e (no need to insert them in the queue)
 - If s is a vertex, then for every triangle t in the star of s compute one window per edge of t opposite to s and insert them in the queue. Also compute one window for each edge e incident at s (no need to insert them in the queue)

Continuous Dijkstra propagation

- Loop while Q becomes empty:
 - Pop a window w from Q
 - Propagate window w
 - Insert each non-empty window generated by propagation into Q
- Propagation may:
 - Add a new window
 - Modify existing windows
 - Delete existing windows
- In the latter two cases, elements already in Q must be either modified or deleted:
 - Q must support min, member, insert, and delete operations in logarithmic time

Geodesic path construction

- Once all edges are covered by windows representing geodesic distance, it is easy to trace a shortest path from any surface point p back to the source
- For a point p inside a face t :
 - Minimize $|p-p'|+d_s(p')$ for all p' on the edges of t
 - Jump to the window containing the point that minimizes that distance

Geodesic path construction

- For a point p on a window
 - Find the adjacent triangle t according to direction τ
 - Reconstruct the position of the pseudo-source s in the plane of t and intersect the line to s with the other two edges of t
 - Jump to the intersection point, which is on new window, and repeat
- When reaching a pseudo-source s
 - Explore the windows incident at s until a window with a pseudo-source different from s is found
- Repeat until reaching the source

Time complexity

- n = number of vertices of M
- Each edge may have up to $O(n)$ windows [Mitchell et al. 1987]
- In the worst case there are $O(n^2)$ windows
- Each window is processed in $O(\log n)$ time (access to priority queue and binary search to find overlapping windows)
- Total: $O(n^2 \log n)$

Geodesic distance between two points

Same algorithm, but we prune the search when a large enough area has been visited

Input: mesh M , points u and v

- Start computing single source geodesic distance from u
- As soon as v is reached by a funnel, compute its distance $d_u(v)$ with respect to that funnel (this is an upper bound to geodesic distance)
- Stop as soon as the window at the top of the queue has a distance larger than $d_u(v)$

PDE-based methods

The Eikonal equation

- The SSGD problem can be formalized as a PDE:

$$|\nabla \phi|^2 = 1 \quad \text{on } \mathbf{M} \setminus \{s\}$$

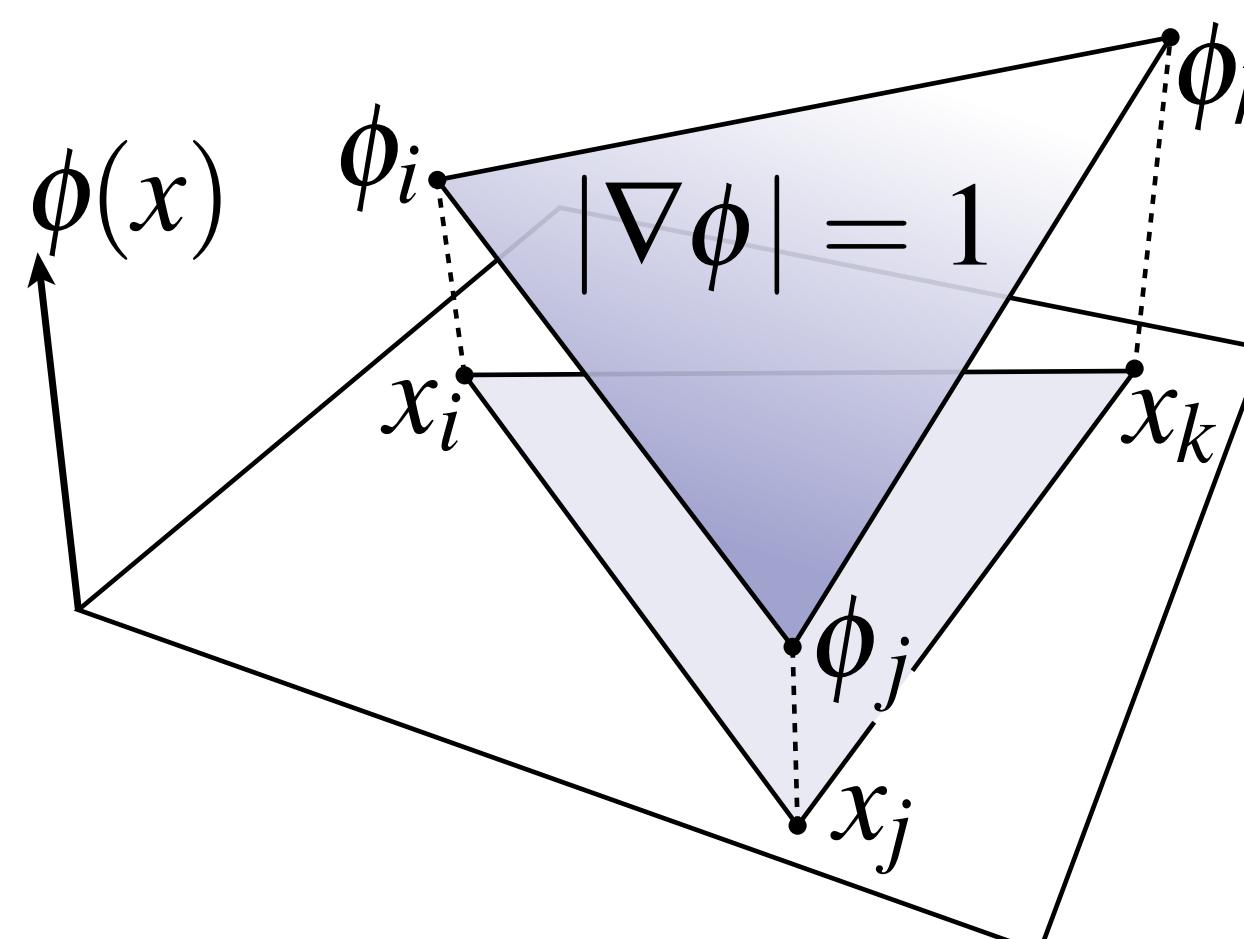
$$\phi(s) = 0$$

- “the *change in distance per unit distance (i.e., the speed)* along the direction of greatest increase should be 1”
- geodesic lines are integral curves of the gradient of the distance field
- the eikonal equation is a non-linear hyperbolic equation, which cannot be approximated using standard linear finite element methods

The fast marching method

[Kimmel and Sethian, 1998]

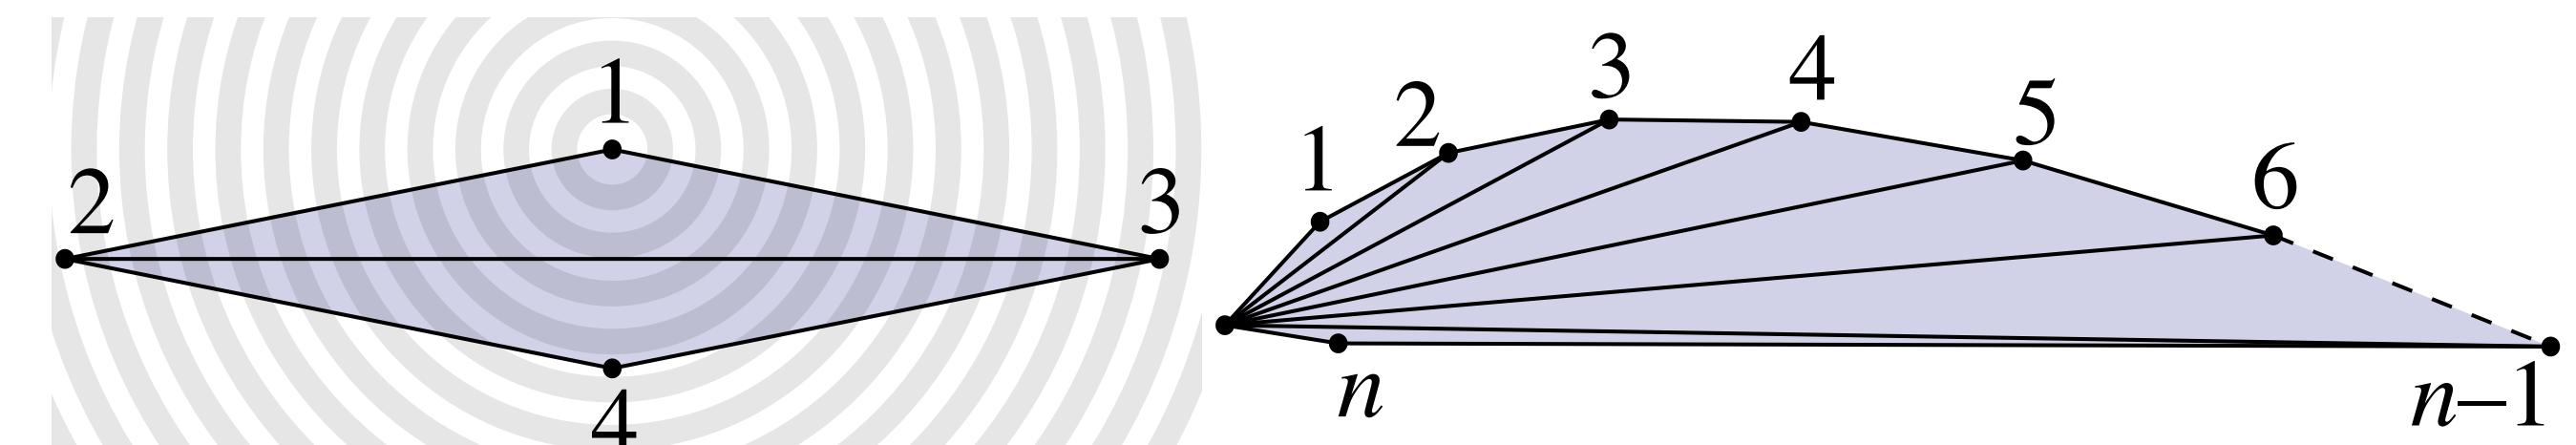
- Follow a Dijkstra-like traversal of the mesh, starting at the source
- Distances are not updated according to paths along edges, but by solving for the linear function that satisfies the eikonal equation
- If the values ϕ_i, ϕ_j at two corners are known, find a value ϕ_k so that the slope $|\nabla\phi|$ of a triangle passing through all three values equals +1



The fast marching method

[Kimmel and Sethian, 1998]

- Linear interpolation of the distance field is inaccurate
 - approximation error increases with the misalignment of the principal direction of the triangle with respect to the (unknown) gradient of ϕ
- The traversal order, given from triangle adjacencies, may fail to analyze closer vertices first, and the problem can be non-local
- Resolving this issue makes the algorithm involved and slow



The Varadhan formula

- The distance field ϕ_x is connected with the *heat kernel* $k_{x,t}$

$$\phi_x(y) = \lim_{t \rightarrow 0} \sqrt{-4t \log k_{x,t}(y)}$$

where $k_{x,t}$ is the solution at time t to the heat equation

$$\frac{d}{dt} k_t = \Delta k_t$$

$$k_0 = \delta_x$$

with δ_x a Dirac delta centered at x

- “Start with heat concentrated at x and let it evolve for a short time”

The Varadhan formula

- The heat equation is linear and can be resolved easily
- However, heat decays exponentially with distance, hence the norm of the result becomes soon unreliable:
 - the smaller t the faster the decay
 - the bigger t the worse the approximation of Varadhan formula

The heat method

[Crane, Weischedel and Wardetzky, 2013]

- Compute the heat kernel by resolving

$$\frac{d}{dt} k_t = \Delta k_t, \quad k_0 = \delta_x$$

for fixed t

- We retain just the gradient direction of the heat kernel
- From the eikonal equation, we know that the gradient of the distance function must have norm 1

- Compute the normalized gradient

$$X = -\nabla k / |\nabla k|$$

- Resolve the Poisson equation

$$\Delta \phi = \nabla \cdot X$$

- We integrate the gradient to obtain the distance field

The heat method

[Crane, Weischedel and Wardetzky, 2013]

- To compute the heat kernel, resolve linear system

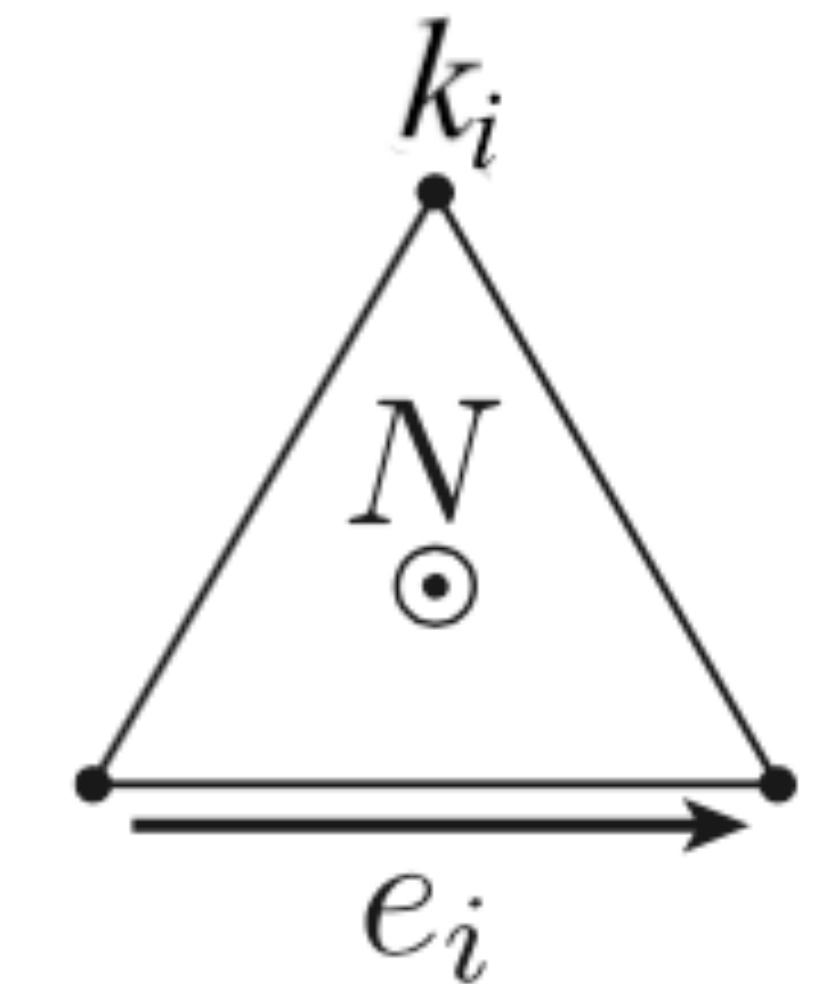
$$(I - tL)k_t = \delta_x$$

for fixed t , where δ_x is a vector with all zero entries and a 1 at x

- The normalized gradient is computed per triangle

$$\nabla k_t = \frac{1}{2A_t} \sum_i k_i (N \times e_i) \quad X = -\nabla k_t / |\nabla k_t|$$

where A_t is the area of the triangle, N is its normal e_i is the edge opposite to vertex i and k_i is the value of the heat kernel at vertex i



The heat method

[Crane, Weischedel and Wardetzky, 2013]

- The divergence associated to vertex i is computed by discrete integration as

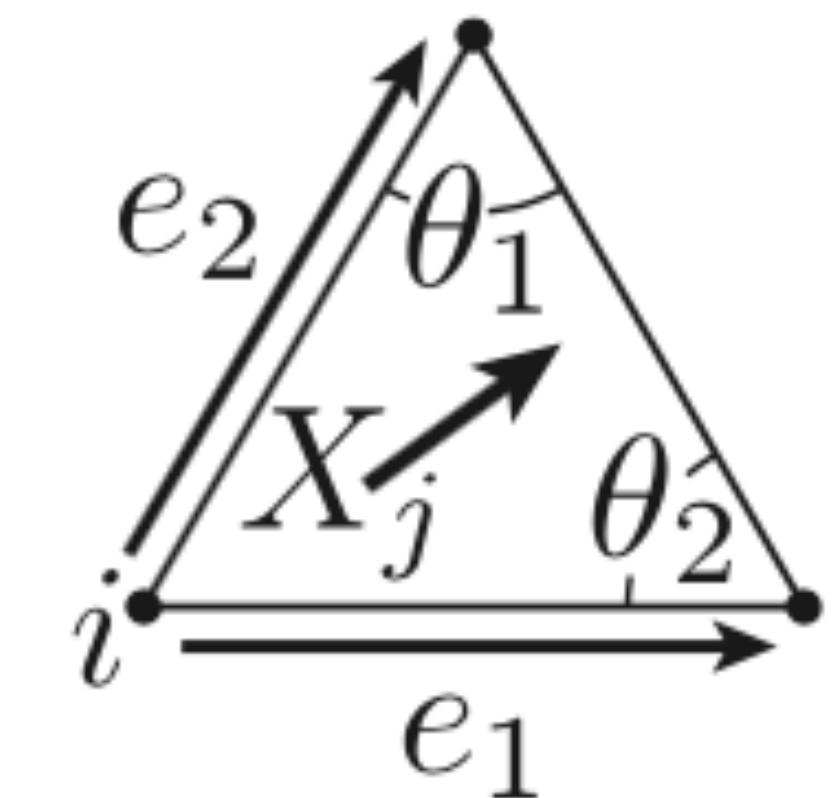
$$\nabla \cdot X_i = \frac{1}{2} \sum_j \cot \theta_1 (e_1 \cdot X_j) + \cot \theta_2 (e_2 \cdot X_j)$$

where summation is on triangles t_j incident at vertex i

- The Poisson problem is resolved with the following linear system

$$L_w \phi = b$$

where L_w is the Laplacian stiffness matrix and b is a vector collecting the divergences $\nabla \cdot X_i$ for all i



The heat method

[Crane, Weischedel and Wardetzky, 2013]

Summary:

- We need to solve two linear systems

$$(I - tL)k_t = \delta_x$$

$$L_w\phi = b$$

that are sparse and have the same size of the input mesh (large!)

- They can be pre-factorized once, so that the solution is fast upon changing the source:
 1. solve a $n \times n$ triangular system
 2. compute gradients and divergences to build vector b
 3. solve another $n \times n$ triangular system

The heat method

[Crane, Weischedel and Wardetzky, 2013]

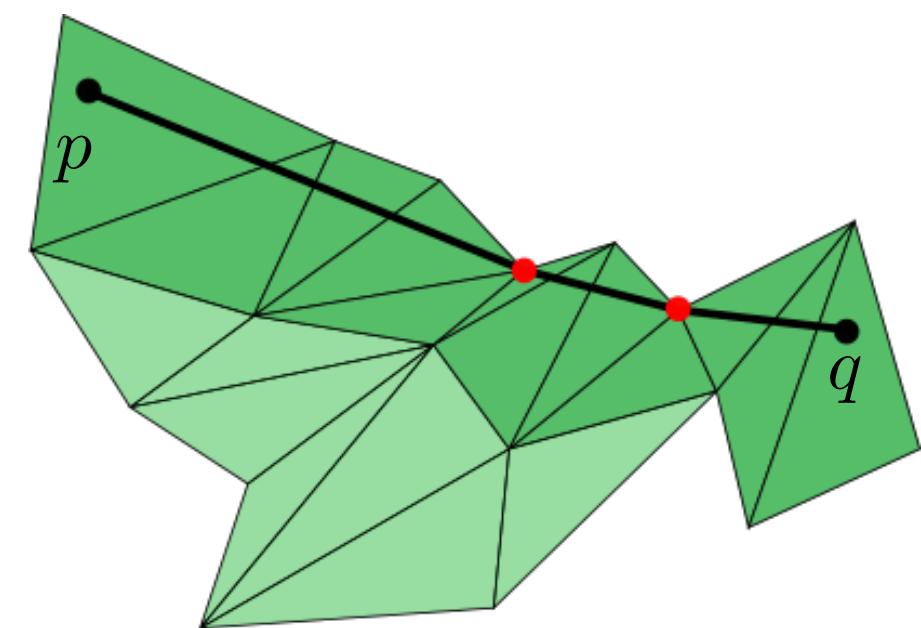
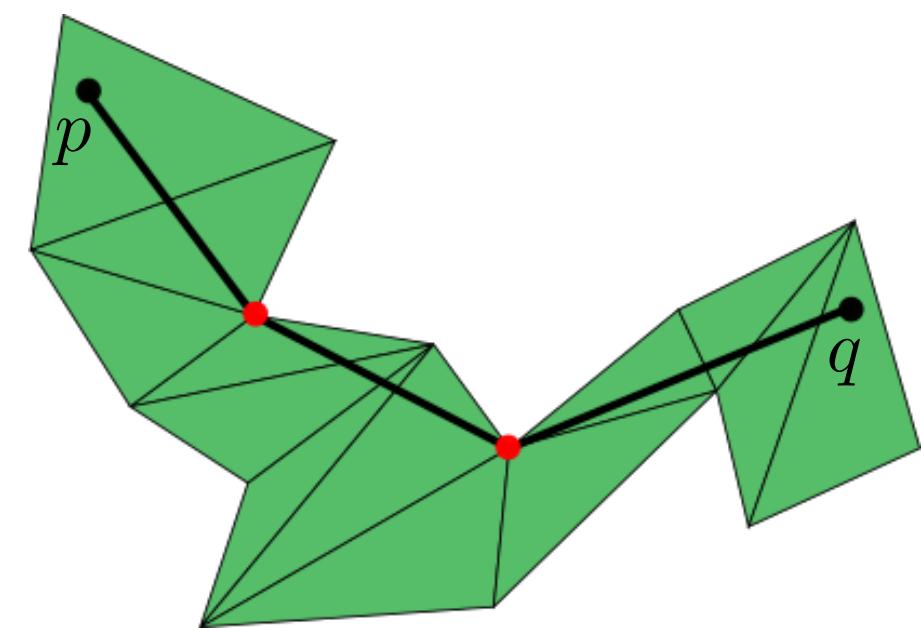
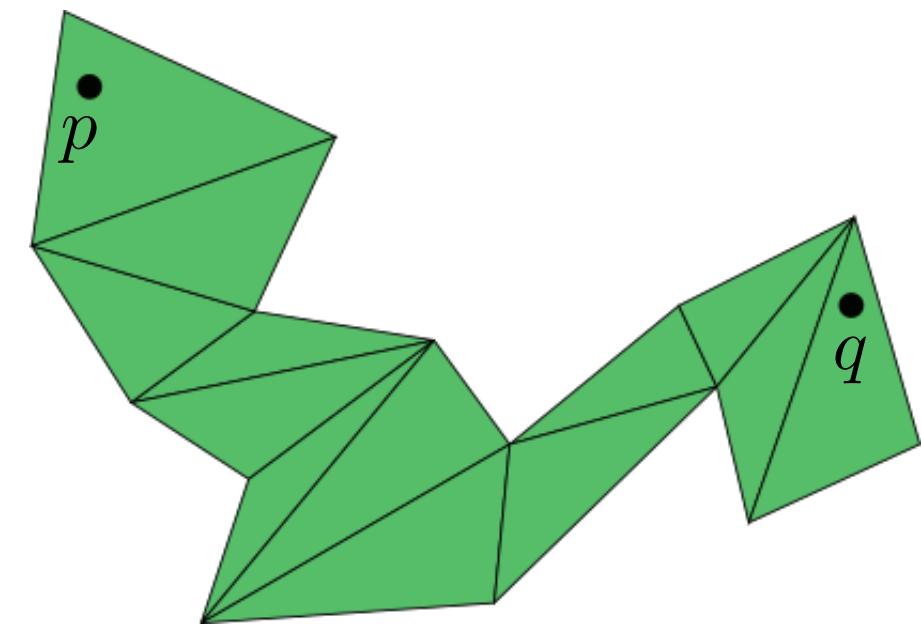
- Pros and cons:
 - Approximated: error ~0.5-1.5% wrt exact polyhedral, but may become large on large meshes
 - Fast after pre-computation: 0.1-0.2 sec on 1M tris
 - Over-smoothed near the cut locus
 - Does not scale well to large meshes
 - Pre-computation becomes very expensive
 - Gradients from heat kernel become too tiny, thus noisy and unreliable even to estimate the gradient direction



Local methods for PPGP

Point to Point Geodesic Path

1. Start with a triangle strip joining p to q (e.g., from shortest path on a graph)
2. Unfold strip to the plane
3. Find shortest path within the strip
4. Straighten the strip to remove turns
5. Repeat 3 and 4 until convergence
 - Pro: fast and scalable
 - Con: finds a *locally shortest* path; step 1 is the bottleneck



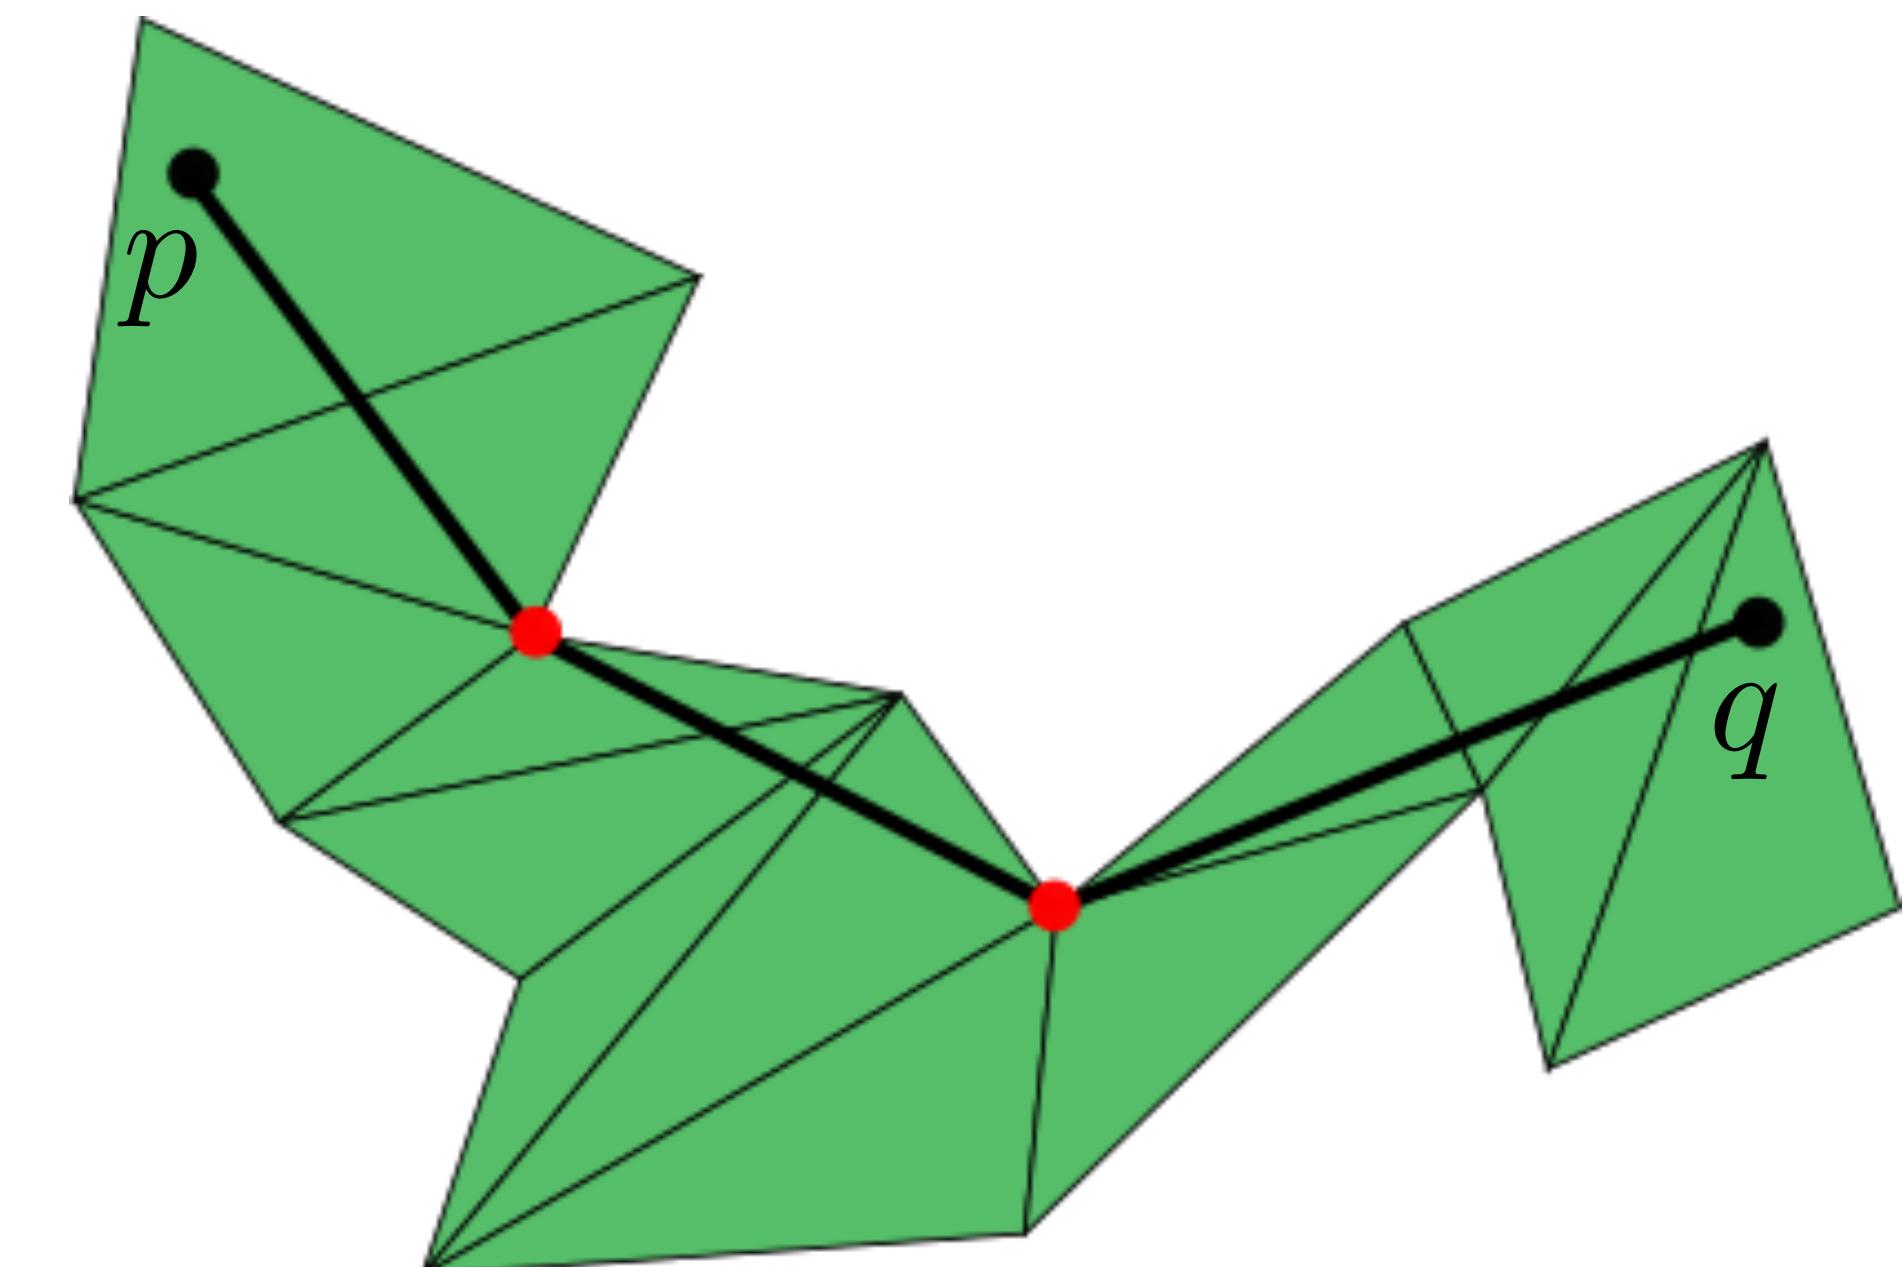
Point to Point Geodesic Path

1. Initial triangle strip:
 - usually found with a graph-based method (e.g., from dual graph)
2. Unfolding:
 - one triangle at a time, by using intersecting circles (similar to locating the source in the MMP algorithm)
3. Shortest path in a strip
 - *Funnel algorithm* [Lee and Preparata, 1984]
4. Straightening the strip
 - *Check angle at turns* [Xia and Wang, 2007]

Shortest path in a triangle strip

[Lee and Preparata, 1984]

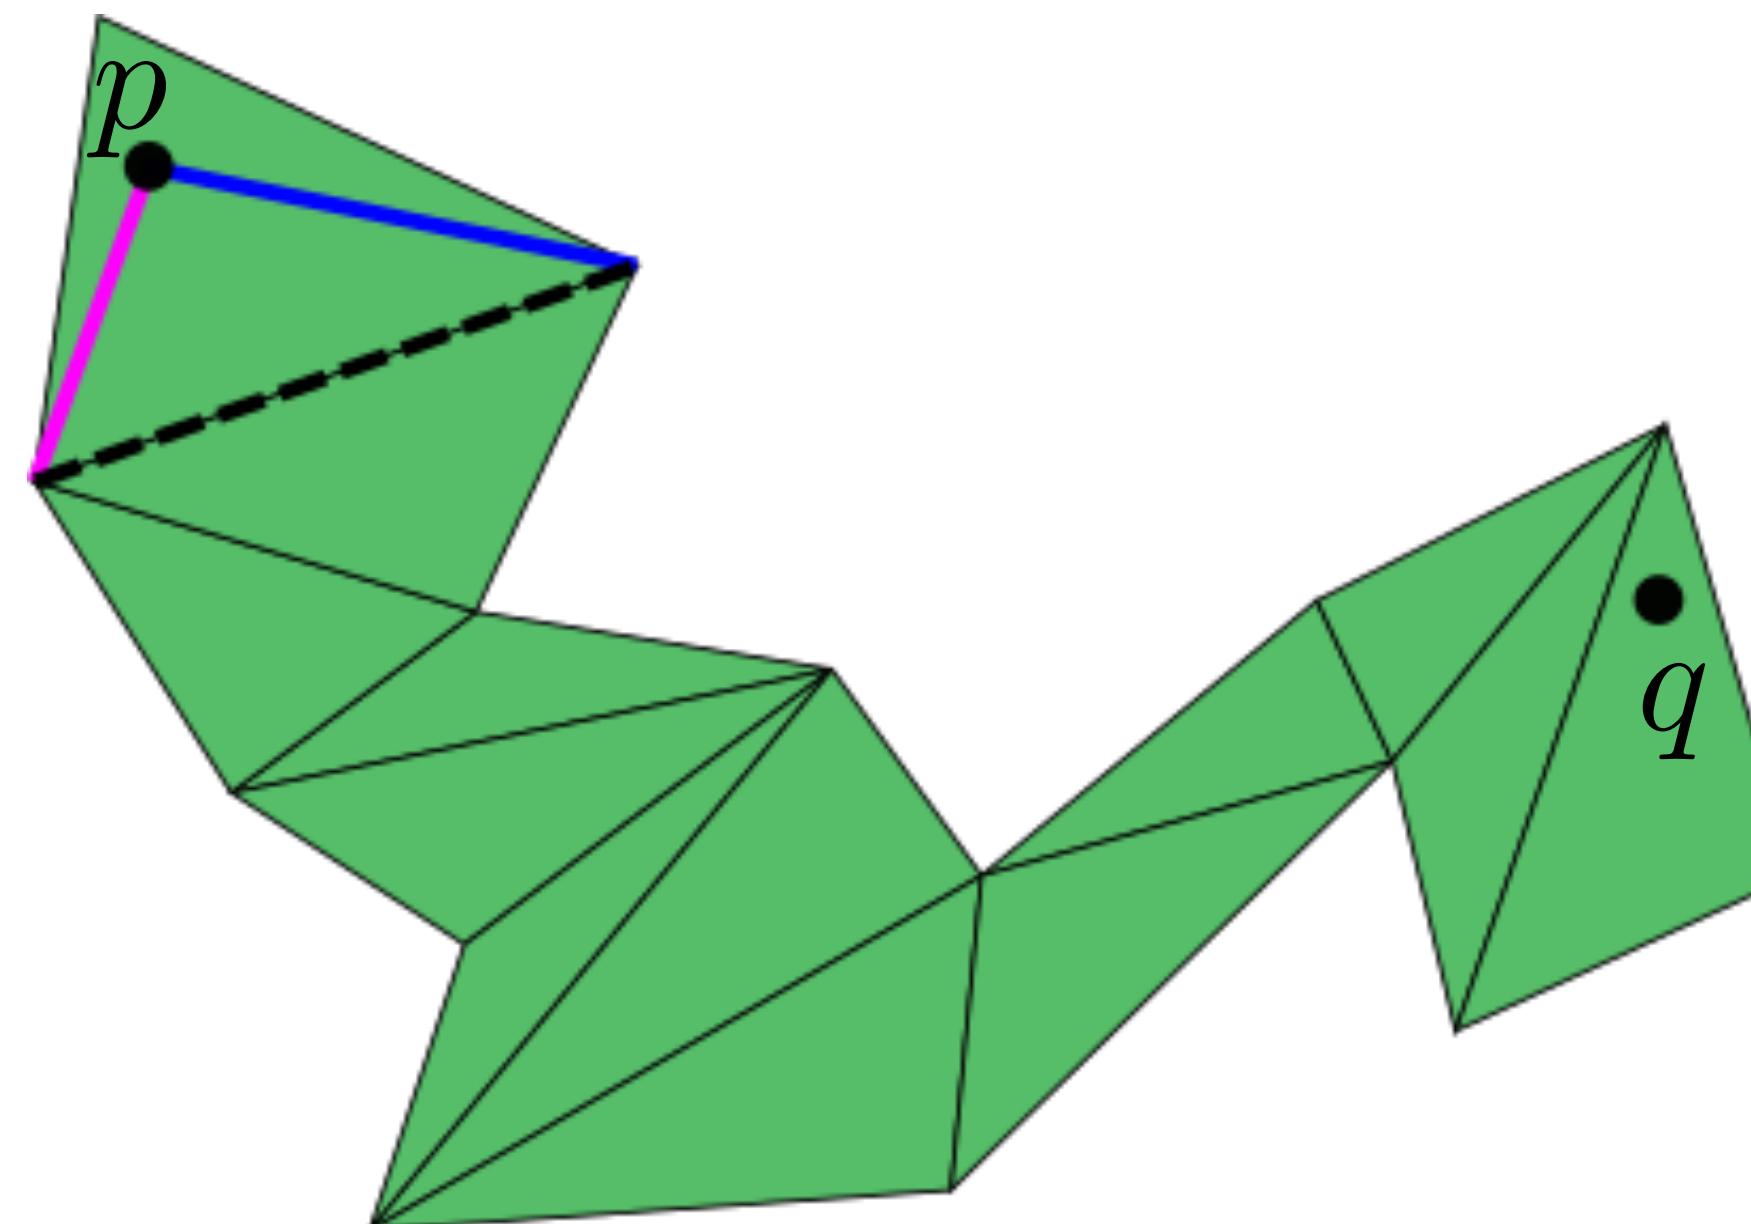
- Input: triangle strip in the plane, containing p and q in its end triangles, respectively
- Output: shortest polyline connecting p to q inside the triangle strip
- The path can turn only at reflex vertices



Shortest path in a triangle strip

[Lee and Preparata, 1984]

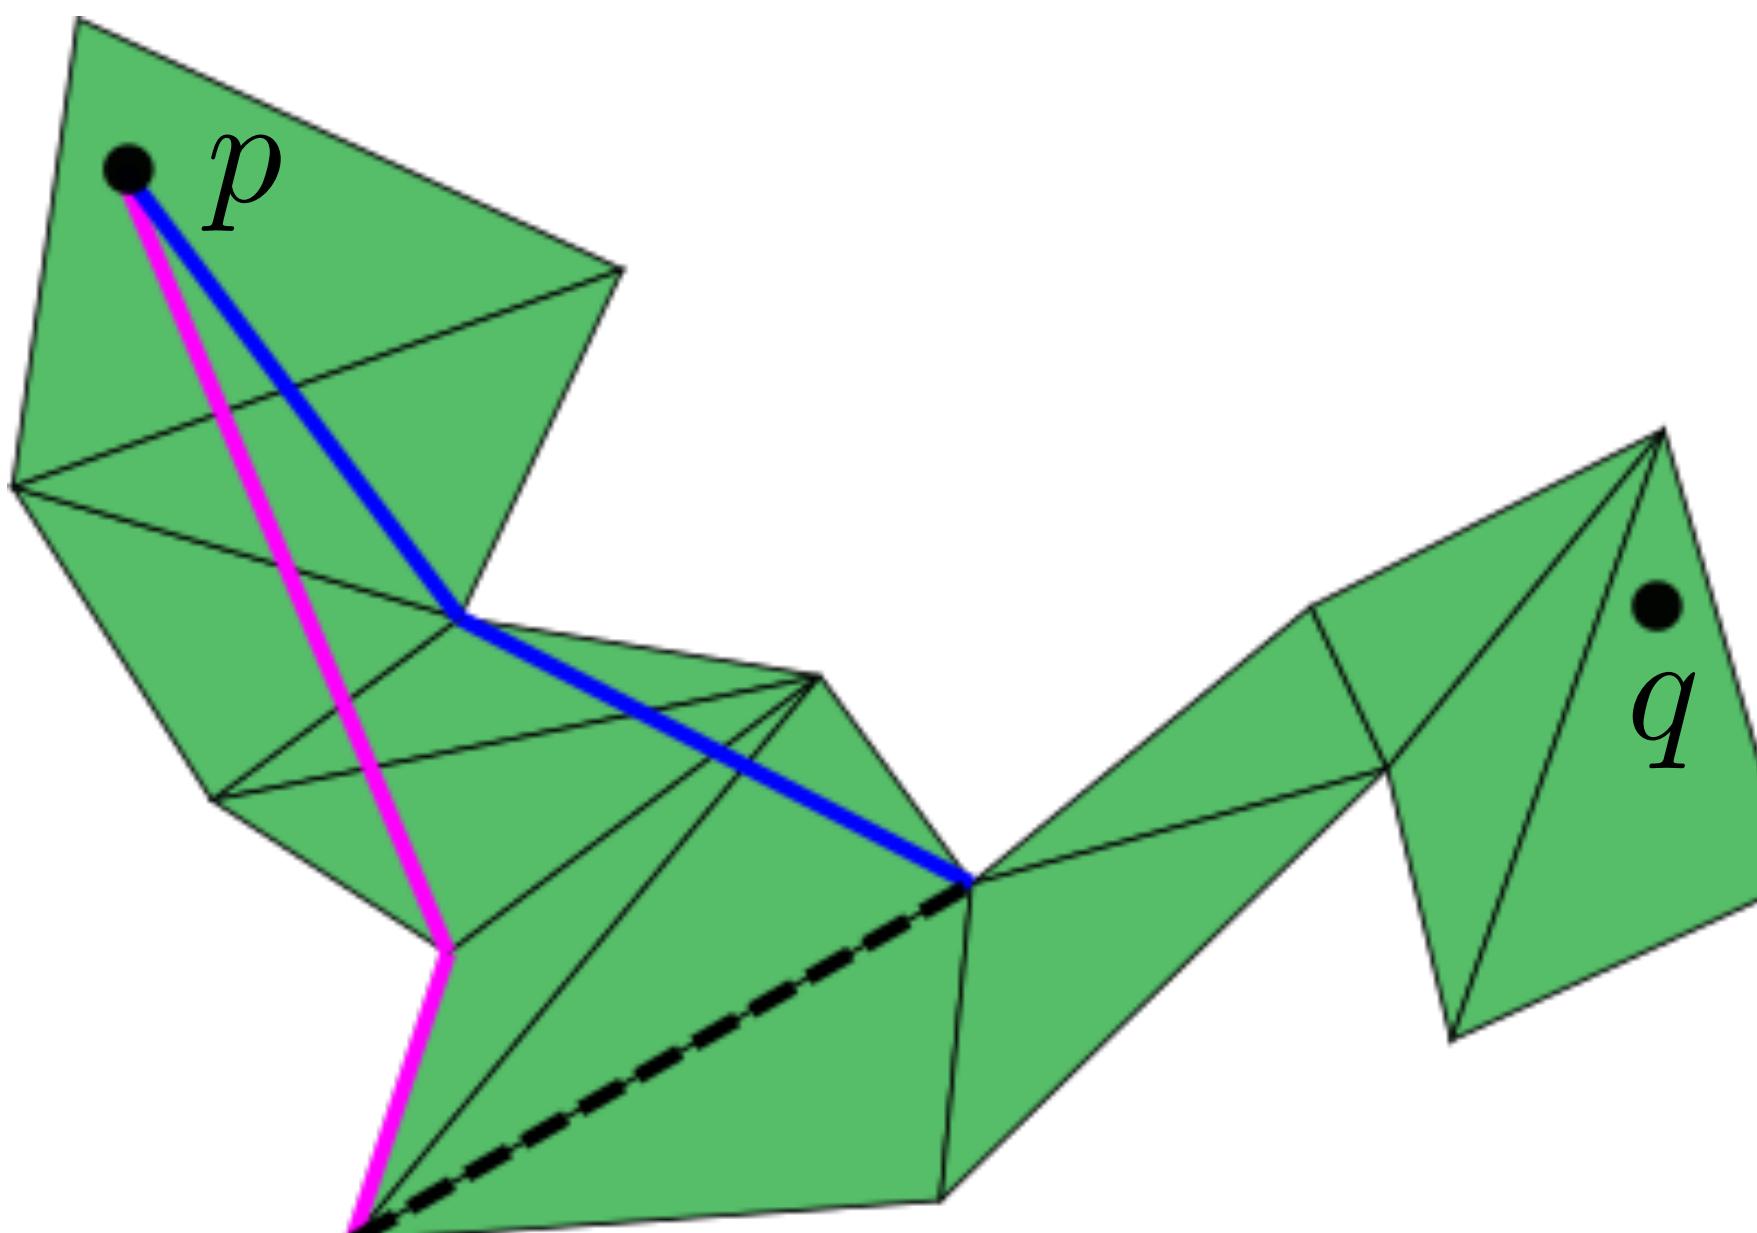
- Initialize funnel: connect p to the endpoints of first transversal edge of the strip; the mouth of the funnel spans all points that can be reached from p



Shortest path in a triangle strip

[Lee and Preparata, 1984]

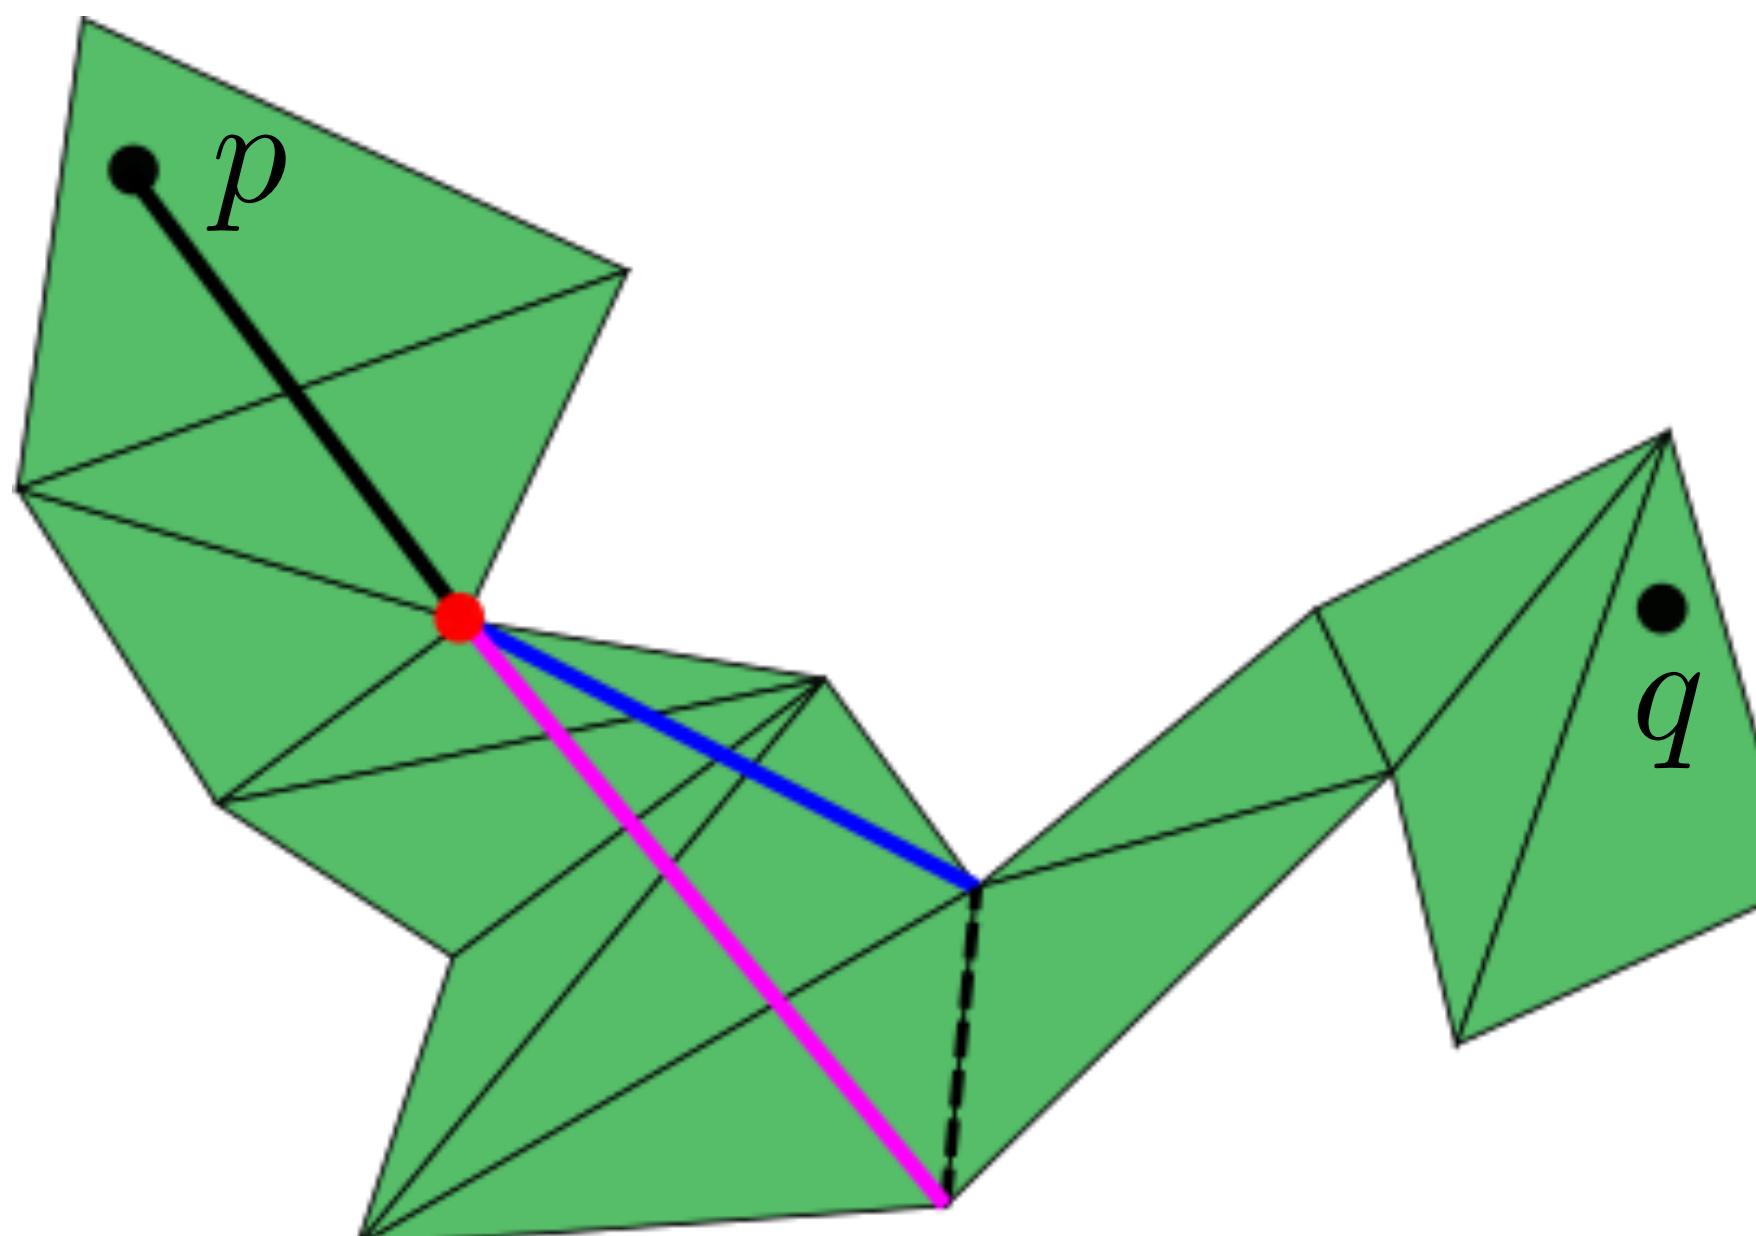
- Expand funnel: move one side of the funnel to the next transversal edge; the sides of the funnel may become concave chains when they overcome reflex edges



Shortest path in a triangle strip

[Lee and Preparata, 1984]

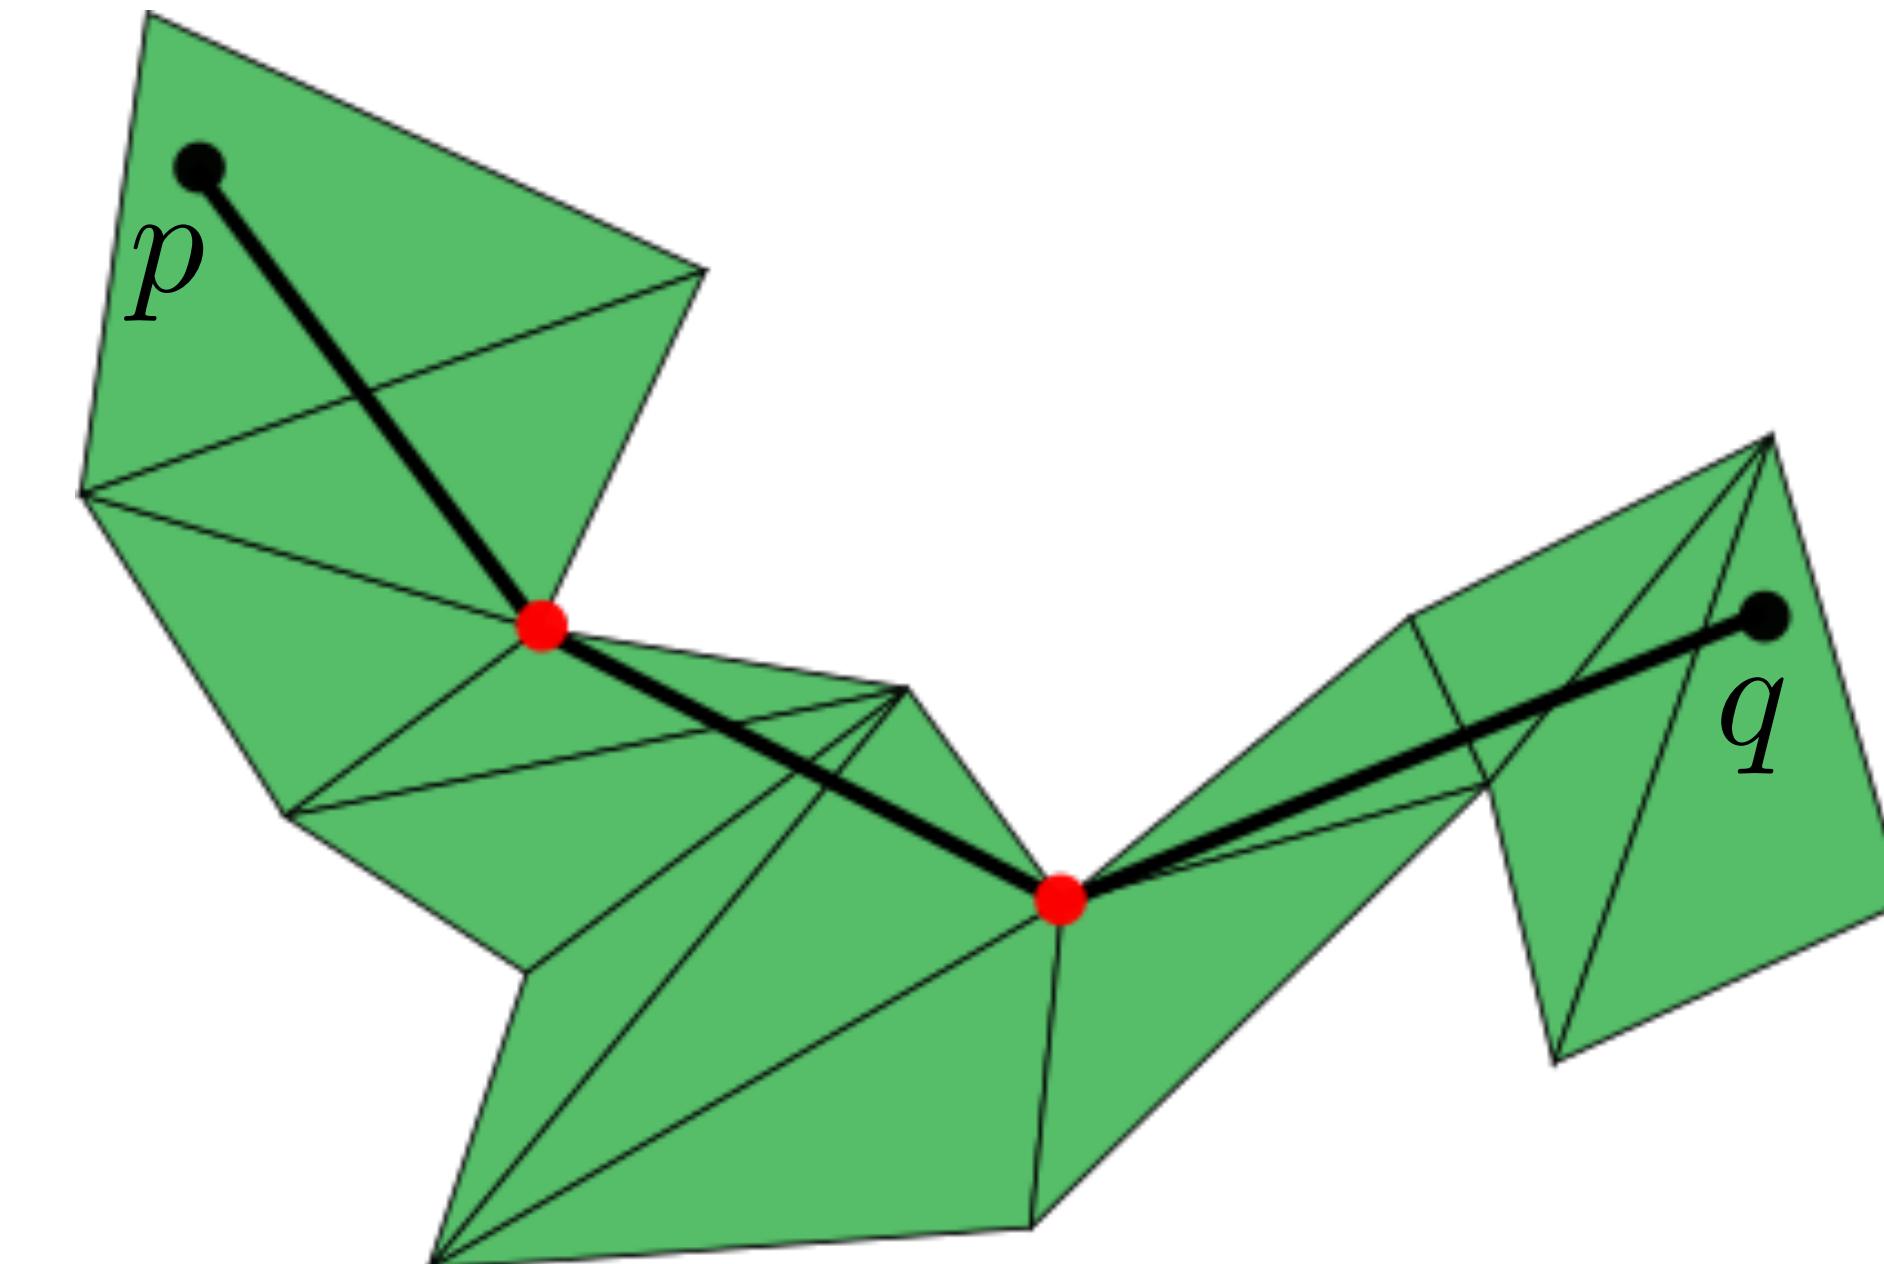
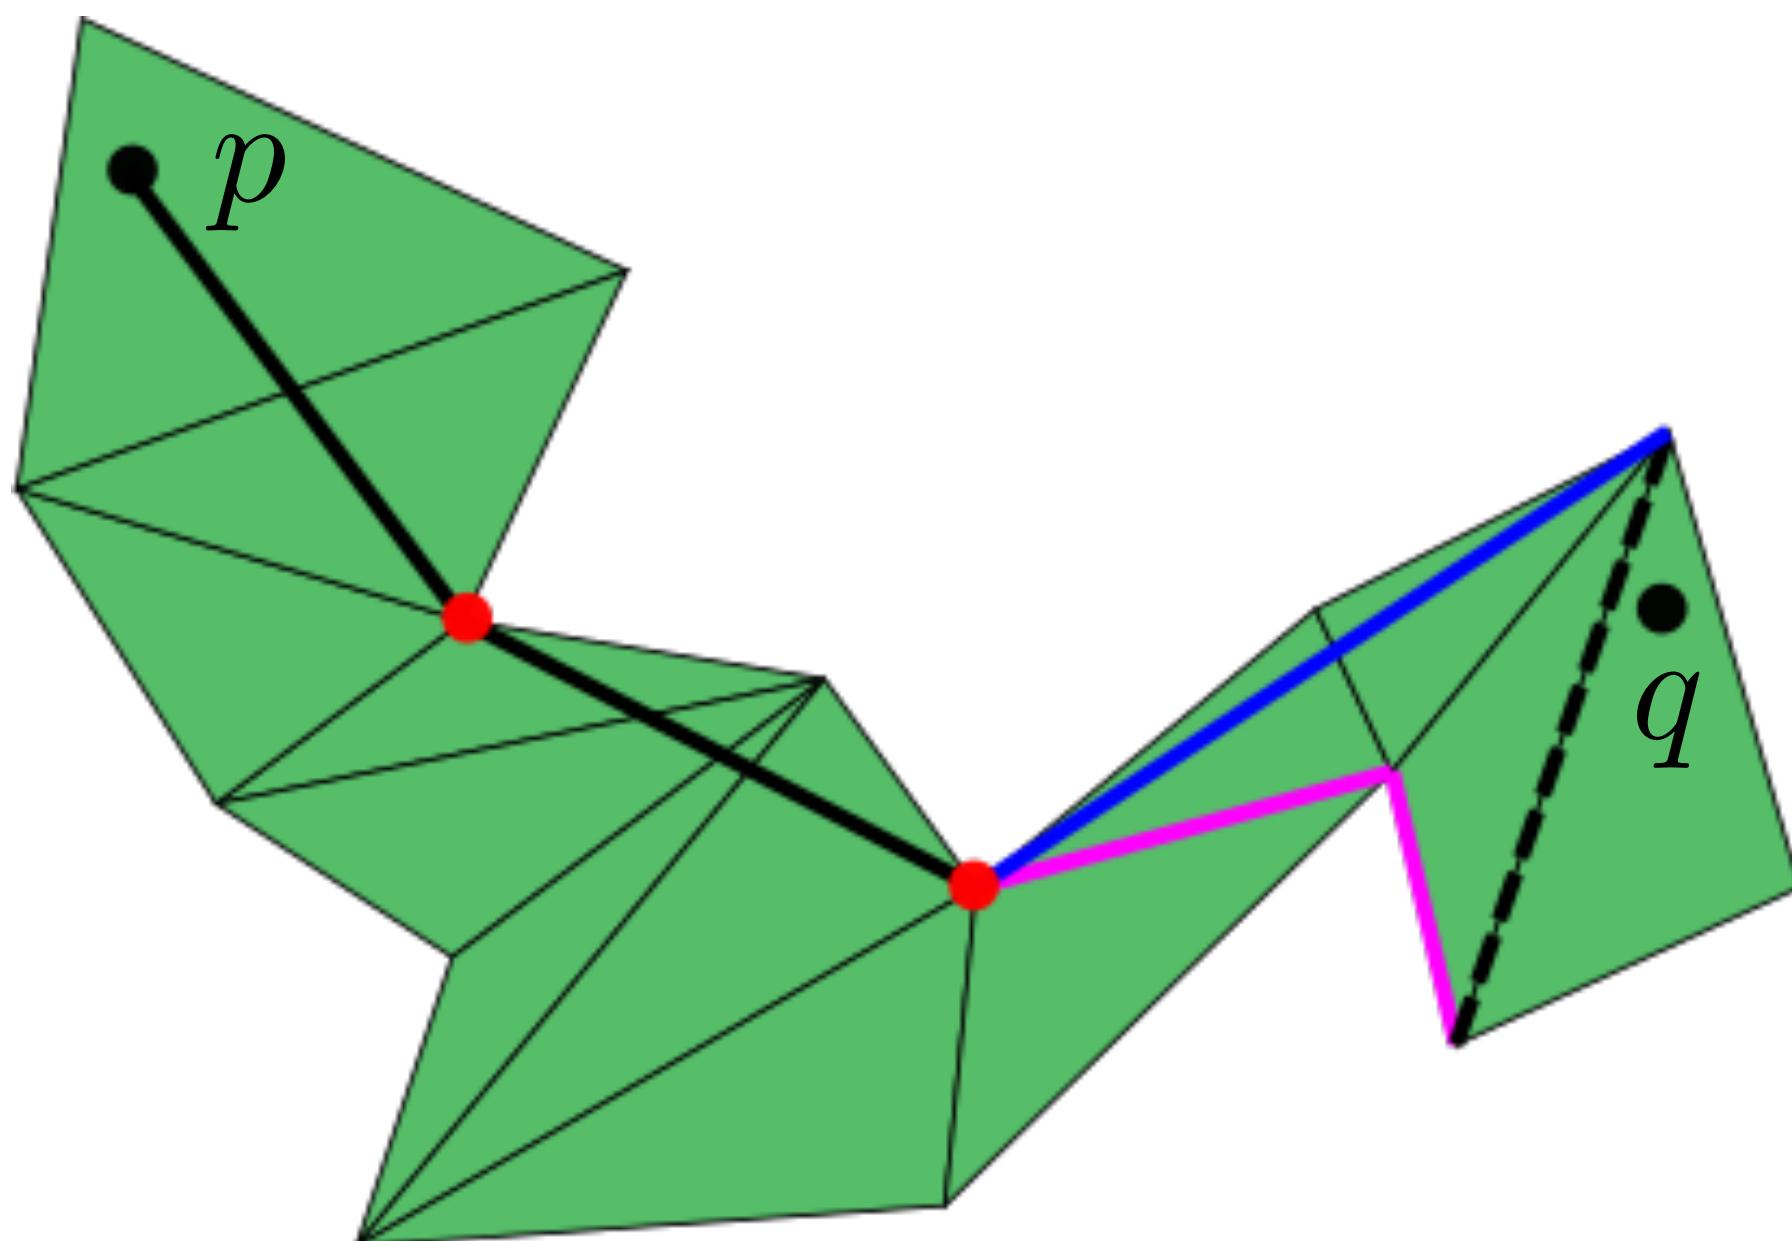
- While advancing, the sides of funnel the may partially collapse to form the beak
- Collapses happen when updating one side overcomes the other side



Shortest path in a triangle strip

[Lee and Preparata, 1984]

- The funnel stops advancing when it reaches the last triangle
- The last pseudo-source is connected to q with a shortest polyline running inside the body of the funnel

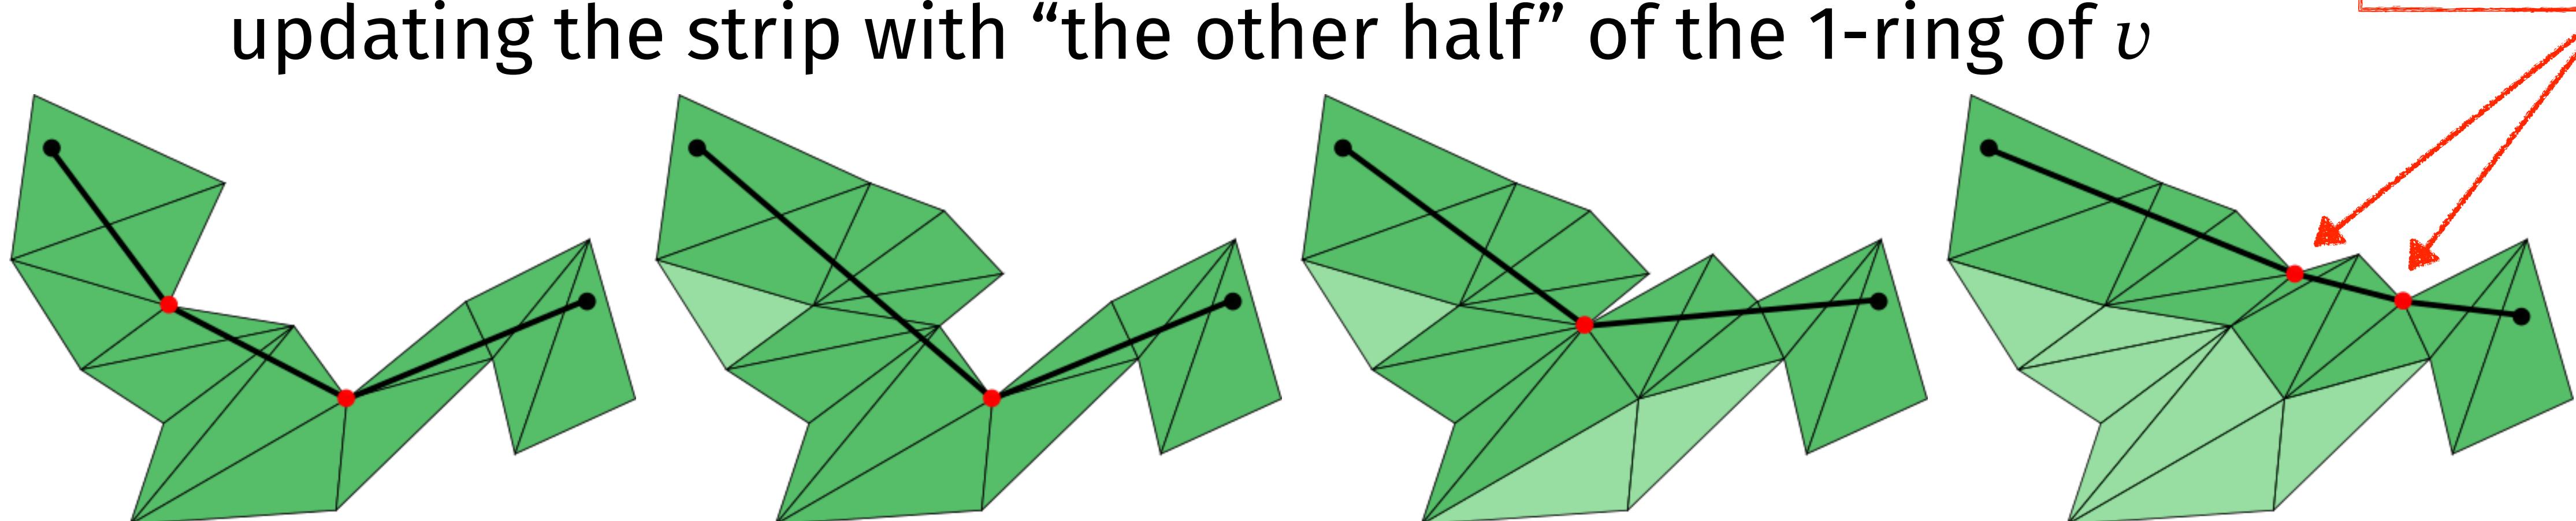


Straightening the strip

[Xia and Wang, 2007]

- A shortest path can bend only at saddle vertices of \mathbf{M}
- A turn at a saddle vertex is necessary only if the total angle of the portion of its 1-ring outside the strip is larger than π
- If a bend occurs at v , which does not fulfill the two conditions above, then the path can be shortened by updating the strip with “the other half” of the 1-ring of v

These two angles are
larger than π in the
unfolded mesh



References

A Survey of Algorithms for Geodesic Paths and Distances
Keenan Crane, Marco Livesu, Enrico Puppo, Yipeng Qin

<https://arxiv.org/abs/2007.10430>

Thank you